

HelixTerminator Security & Zero Trust Architecture Specification

Document Version: 1.0.0

Classification: Internal — Engineering Confidential

Effective Date: 2026-06-28

Owner: Platform Security Team

Review Cycle: Quarterly

Trust Domain: helixterminator.io

Table of Contents

1. [Zero Trust Architecture](#)
 2. [Authentication Architecture](#)
 3. [SSH Certificate Authority](#)
 4. [Vault Encryption Architecture](#)
 5. [SSH Key Management Security](#)
 6. [RBAC Model](#)
 7. [Audit Logging](#)
 8. [Network Security](#)
 9. [Compliance](#)
 10. [Incident Response](#)
-

Preface

HelixTerminator is an enterprise-grade SSH client platform that processes some of the most sensitive assets in any organization: SSH private keys, server credentials, encrypted vaults containing infrastructure secrets, and live terminal sessions into production systems. A breach of HelixTerminator is not merely a data leak — it is a complete compromise of the infrastructure the product is meant to secure.

This specification establishes the non-negotiable security architecture for every component of the platform. It is designed to be implementation-ready: every architectural decision maps to code, every policy maps to enforceable configuration, every threat maps to a countermeasure.

The overarching security principle is **Zero Trust**: no actor — user, device, workload, or network — is trusted by default. Every access decision is made fresh, based on verified identity, device health, network context, and policy. Trust is never implicit, always explicitly granted, and continuously re-evaluated.

1. Zero Trust Architecture

1.1 Principles and Philosophy

Zero Trust is not a product or a network zone — it is a security model predicated on the assumption that the network perimeter has already been breached. Every request, regardless of origin, is treated as potentially hostile until proven otherwise. The HelixTerminator Zero Trust implementation embodies three foundational axioms:

1. **Never Trust, Always Verify.** Every service-to-service call, every user request, every device connection must present verifiable proof of identity. Network location — being “inside” the VPC — confers no implicit trust.
2. **Assume Breach.** Design all systems assuming an adversary has already established a foothold somewhere in the environment. Lateral movement must be impossible even from a compromised internal service. Each workload sees only what it needs to function.
3. **Least-Privilege Access.** Every identity — user, device, or workload — operates with the minimum permissions required. Permissions are time-bound, scoped to specific resources, and revocable without restart.

1.2 Trust Boundaries and Threat Model

HelixTerminator operates across four trust zones, each with strictly enforced boundaries:

Zone	Description	Trust Level
External	Public internet, end-user clients	Zero trust — all traffic authenticated
Edge	API Gateway, CDN, load balancers	Untrusted — terminate TLS, enforce rate limits
Service Mesh	Internal microservices	mTLS-enforced trust via SPIFFE SVIDs
Data Tier	Databases, Kafka, Redis	Authenticated connections only, encrypted at rest

Threat actors modeled in this specification:

- **External attacker:** Internet-based adversary with no credentials
- **Compromised credential:** Legitimate user credentials stolen (phishing, credential stuffing)
- **Compromised service:** An internal microservice running malicious code (supply chain attack)
- **Malicious insider:** A legitimate employee with valid credentials acting outside their authorization
- **Nation-state adversary:** Advanced persistent threat with long time horizons and sophisticated techniques

1.3 SPIFFE/SPIRE Workload Identity

1.3.1 Overview

SPIFFE (Secure Production Identity Framework for Everyone) provides a universal workload identity standard. Every HelixTerminator workload receives a SPIFFE Verifiable Identity Document (SVID) cryptographically signed by the trust domain root. This eliminates the need for static secrets in service-to-service communication entirely.

Trust Domain: `helixterminator.io`

All SPIFFE IDs follow the format:

```
spiffe://helixterminator.io/<service-path>
```

Service identity examples:

```
spiffe://helixterminator.io/ns/default/sa/auth-service
spiffe://helixterminator.io/ns/default/sa/pki-service
spiffe://helixterminator.io/ns/default/sa/vault-service
spiffe://helixterminator.io/ns/default/sa/session-manager
```

```
spiffe://helixterminator.io/ns/default/sa/audit-service
spiffe://helixterminator.io/ns/vault/sa/vault-api
spiffe://helixterminator.io/ns/pki/sa/ca-signer
```

1.3.2 SPIRE Server High-Availability Configuration

The SPIRE server is the root of workload identity issuance. It must be highly available, as a SPIRE server outage prevents new SVIDs from being issued and prevents existing SVIDs from rotating, eventually causing service authentication failures.

Deployment topology: 3-node SPIRE server cluster with external datastore (PostgreSQL) and distributed KMS (AWS KMS / HashiCorp Vault).

```
# spire-server-config.hcl
server {
  bind_address      = "0.0.0.0"
  bind_port         = "8081"
  trust_domain      = "helixterminator.io"
  data_dir          = "/run/spire/data"
  log_level         = "INFO"
  log_format        = "json"

  # CA TTL – SVIDs issued with 24h lifetime
  # Root CA rotates every 365 days, Intermediate every 30 days
  ca_ttl            = "720h"      # 30 days – intermediate CA lifetime
  default_svid_ttl  = "1h"       # default SVID TTL (overridable per
                                # entry)

  # Key management – AWS KMS in production
  key_manager "aws_kms" {
    plugin_data {
      region          = "us-east-1"
      key_policy_file = "/run/spire/key-policy.json"
    }
  }

  # External datastore for HA clustering
  data_store "sql" {
    plugin_data {
      database_type      = "postgres"
      connection_string  = "dbname=spire user=spire password=$
                           {SPIRE_DB_PASSWORD} host=spire-db.internal
                           sslmode=require"
      max_open_conns     = 20
    }
  }
}
```

```

        max_idle_conns      = 5
        conn_max_lifetime   = "5m"
    }
}

# Node attestation – Kubernetes PSAT
NodeAttestor "k8s_psat" {
    plugin_data {
        clusters = {
            "helixterm-prod" = {
                service_account_allow_list = ["default:spire-agent"]
                audience                 = ["spire-server"]
                allowed_node_label_keys   = []
                allowed_pod_label_keys    = []
                kube_config_file          = ""
                token_path                 = ""
            }
        }
    }
}

# Workload attestation – Kubernetes
WorkloadAttestor "k8s" {
    plugin_data {
        skip_kubelet_verification = false
        node_name_env             = "MY_NODE_NAME"
        max_poll_attempts          = 5
        poll_retry_interval        = "500ms"
    }
}

# Upstream authority – offline root CA via AWS KMS
UpstreamAuthority "aws_pca" {
    plugin_data {
        region = "us-east-1"
        certificate_authority_arn = "arn:aws:acm-pca:us-east-1:123456789:certificate-authority/helixterm-root"
        signing_algorithm = "SHA256WITHRSA"
        ca_signing_template_arn = "arn:aws:acm-pca:::template/SubordinateCACertificate_PathLen1/V1"
        assume_role_arn = ""
        endpoint = ""
        supplemental_bundle_path = "/run/spire/upstream-bundle.pem"
    }
}

```

```

# Health checks
health_checks {
  listener_enabled = true
  bind_address     = "0.0.0.0"
  bind_port        = "8080"
  live_path        = "/live"
  ready_path       = "/ready"
}

# Federation – allows cross-domain trust (future: partner IdP)
federation {
  bundle_endpoint {
    address = "0.0.0.0"
    port    = 8443
  }
}
}

```

SPIRE Server Kubernetes Deployment:

```

apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: spire-server
  namespace: spire
spec:
  replicas: 3
  serviceName: spire-server
  selector:
    matchLabels:
      app: spire-server
  template:
    metadata:
      labels:
        app: spire-server
    spec:
      serviceAccountName: spire-server
      # Anti-affinity: spread across availability zones
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            - labelSelector:
                matchExpressions:
                  - key: app

```

```
        operator: In
        values: ["spire-server"]
        topologyKey: topology.kubernetes.io/zone
containers:
- name: spire-server
  image: ghcr.io/spiffe/spire-server:1.9.0
  args: ["-config", "/run/spire/config/server.conf"]
  ports:
  - containerPort: 8081
    name: grpc
  - containerPort: 8080
    name: health
  volumeMounts:
  - name: spire-config
    mountPath: /run/spire/config
    readOnly: true
  - name: spire-data
    mountPath: /run/spire/data
  - name: spire-bundle
    mountPath: /run/spire/bundle
  resources:
    requests:
      cpu: 200m
      memory: 256Mi
    limits:
      cpu: 1000m
      memory: 512Mi
  livenessProbe:
    httpGet:
      path: /live
      port: 8080
    initialDelaySeconds: 15
    periodSeconds: 10
  readinessProbe:
    httpGet:
      path: /ready
      port: 8080
    initialDelaySeconds: 5
    periodSeconds: 5
volumes:
- name: spire-config
  configMap:
    name: spire-server
- name: spire-bundle
```

```

        emptyDir: {}
    volumeClaimTemplates:
    - metadata:
        name: spire-data
      spec:
        accessModes: ["ReadWriteOnce"]
        storageClassName: gp3-encrypted
        resources:
          requests:
            storage: 1Gi

```

1.3.3 SPIRE Agent DaemonSet

A SPIRE agent runs on every node in the cluster, exposing the Workload API (Unix domain socket) to workloads on that node.

```

apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: spire-agent
  namespace: spire
spec:
  selector:
    matchLabels:
      app: spire-agent
  updateStrategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
  template:
    metadata:
      labels:
        app: spire-agent
    spec:
      hostPID: true
      hostNetwork: true
      dnsPolicy: ClusterFirstWithHostNet
      serviceAccountName: spire-agent
      initContainers:
      - name: init-wait-for-spire-server
        image: busybox:1.36
        command: ['sh', '-c', 'until nc -z spire-
          server.spire.svc.cluster.local 8081; do sleep 2; done']
      containers:
      - name: spire-agent

```

```
image: ghcr.io/spiffe/spire-agent:1.9.0
args: ["-config", "/run/spire/config/agent.conf"]
volumeMounts:
- name: spire-config
  mountPath: /run/spire/config
  readOnly: true
- name: spire-bundle
  mountPath: /run/spire/bundle
  readOnly: true
- name: spire-agent-socket
  mountPath: /run/spire/sockets
  readOnly: false
- name: spire-agent-data
  mountPath: /run/spire/data
- name: hostpid
  mountPath: /host/proc
  readOnly: true
env:
- name: MY_NODE_NAME
  valueFrom:
    fieldRef:
      fieldPath: spec.nodeName
resources:
  requests:
    cpu: 50m
    memory: 64Mi
  limits:
    cpu: 200m
    memory: 128Mi
volumes:
- name: spire-config
  configMap:
    name: spire-agent
- name: spire-bundle
  configMap:
    name: spire-bundle
- name: spire-agent-socket
  hostPath:
    path: /run/spire/sockets
    type: DirectoryOrCreate
- name: spire-agent-data
  hostPath:
    path: /run/spire/agent-data
    type: DirectoryOrCreate
```

```
- name: hostpid
  hostPath:
    path: /proc
```

SPIRE Agent Configuration:

```
# spire-agent.hcl
agent {
  data_dir      = "/run/spire/data"
  log_level     = "INFO"
  log_format    = "json"
  trust_domain  = "helixterminator.io"
  server_address = "spire-server.spire.svc.cluster.local"
  server_port   = "8081"
  socket_path   = "/run/spire/sockets/agent.sock"

  # Trust bundle from SPIRE server bootstrap
  trust_bundle_path = "/run/spire/bundle/bundle.crt"

  insecure_bootstrap = false

  NodeAttestor "k8s_psat" {
    plugin_data {
      cluster      = "helixterm-prod"
      token_path   = "/var/run/secrets/tokens/spire-agent"
    }
  }

  WorkloadAttestor "k8s" {
    plugin_data {
      node_name_env      = "MY_NODE_NAME"
      skip_kubelet_verification = false
      max_poll_attempts  = 5
      poll_retry_interval = "500ms"
    }
  }

  KeyManager "memory" {
    plugin_data {}
  }

  health_checks {
    listener_enabled = true
    bind_address     = "127.0.0.1"
    bind_port        = "8080"
  }
}
```

```

    live_path      = "/live"
    ready_path     = "/ready"
  }
}

```

1.3.4 SVID Registration Entries

Registration entries map workload attestation selectors to SPIFFE IDs:

```

# Register auth-service
spire-server entry create \
  -spiffeID spiffe://helixterminator.io/ns/default/sa/auth-
    service \
  -parentID spiffe://helixterminator.io/spire/agent/k8s_psat/
    helixterm-prod/NODE_ID \
  -selector k8s:ns:default \
  -selector k8s:sa:auth-service \
  -selector k8s:pod-label:app:auth-service \
  -ttl 3600

# Register vault-service
spire-server entry create \
  -spiffeID spiffe://helixterminator.io/ns/vault/sa/vault-api \
  -parentID spiffe://helixterminator.io/spire/agent/k8s_psat/
    helixterm-prod/NODE_ID \
  -selector k8s:ns:vault \
  -selector k8s:sa:vault-service \
  -selector k8s:pod-label:app:vault-service \
  -ttl 3600

# Register PKI CA service (higher privilege – shorter TTL)
spire-server entry create \
  -spiffeID spiffe://helixterminator.io/ns/pki/sa/ca-signer \
  -parentID spiffe://helixterminator.io/spire/agent/k8s_psat/
    helixterm-prod/NODE_ID \
  -selector k8s:ns:pki \
  -selector k8s:sa:pki-ca-signer \
  -selector k8s:pod-label:app:pki-ca \
  -ttl 1800

```

1.3.5 SVID Rotation

SVIDs are short-lived (1 hour TTL) and rotated automatically by the SPIRE agent. The Go SDK for SVID rotation in HelixTerminator services:

```

// pkg/identity/workload.go
package identity

import (
    "context"
    "crypto/tls"
    "crypto/x509"
    "fmt"
    "sync"

    "github.com/spiffe/go-spiffe/v2/spiffeid"
    "github.com/spiffe/go-spiffe/v2/spiffetls/tlsconfig"
    "github.com/spiffe/go-spiffe/v2/workloadapi"
)

const workloadAPISocket = "unix:///run/spire/sockets/agent.sock"

// WorkloadIdentity manages the SPIFFE workload identity for this
// service.
type WorkloadIdentity struct {
    mu      sync.RWMutex
    source *workloadapi.X509Source
}

// NewWorkloadIdentity creates a new workload identity manager,
// connecting to
// the SPIRE agent Workload API and watching for SVID rotations.
func NewWorkloadIdentity(ctx context.Context) (*WorkloadIdentity,
    error) {
    source, err := workloadapi.NewX509Source(
        ctx,
        workloadapi.WithClientOptions(
            workloadapi.WithAddr(workloadAPISocket),
        ),
    )
    if err != nil {
        return nil, fmt.Errorf("failed to create X.509 source:
            %w", err)
    }
    return &WorkloadIdentity{source: source}, nil
}

// ServerTLSConfig returns a *tls.Config for use as a TLS server,
// enforcing
// mTLS with SPIFFE peer authentication. Only peers from the
// helixterminator.io

```

```

// trust domain are accepted.
func (w *WorkloadIdentity) ServerTLSConfig() *tls.Config {
    trustDomain :=
        spiffeid.RequireTrustDomainFromString("helixterminator.io")
    return tlsconfig.MTLSServerConfig(w.source, w.source,
        tlsconfig.AuthorizeMemberOf(trustDomain))
}

// ClientTLSConfig returns a *tls.Config for connecting to a
// specific peer SPIFFE ID.
func (w *WorkloadIdentity) ClientTLSConfig(peerID spiffeid.ID)
    *tls.Config {
    return tlsconfig.MTLSClientConfig(w.source, w.source,
        tlsconfig.AuthorizeID(peerID))
}

// SVID returns the current SVID certificate and key.
func (w *WorkloadIdentity) SVID() (*tls.Certificate, error) {
    svids, err := w.source.GetX509SVIDs()
    if err != nil {
        return nil, fmt.Errorf("failed to get X.509 SVIDs: %w",
            err)
    }
    if len(svids) == 0 {
        return nil, fmt.Errorf("no SVIDs available")
    }
    return svids[0].Leaf, nil
}

// Close shuts down the SVID watcher.
func (w *WorkloadIdentity) Close() error {
    return w.source.Close()
}

// ValidatePeer validates that an incoming connection's peer
// certificate
// has the expected SPIFFE ID.
func ValidatePeer(cert *x509.Certificate, expectedID spiffeid.ID)
    error {
    ids, err := x509svid.IDsFromCert(cert)
    if err != nil {
        return
            fmt.Errorf("failed to extract SPIFFE IDs from cert: %w",
                err)
    }
    for _, id := range ids {
        if id == expectedID {
            return nil
        }
    }
}

```

```

    }
}
return fmt.Errorf("peer SPIFFE ID not authorized: expected
    %s", expectedID)
}

```

1.4 Mutual TLS (mTLS) via Istio Service Mesh

1.4.1 Istio Installation and Configuration

HelixTerminator uses Istio as the service mesh layer, enforcing mTLS between all services. Istio integrates with SPIRE via the SPIFFE CSI driver for certificate provisioning.

```

# istio-operator-config.yaml
apiVersion: install.istio.io/v1alpha1
kind: IstioOperator
metadata:
  name: helixterm-istio
  namespace: istio-system
spec:
  profile: production
  meshConfig:
    defaultConfig:
      proxyMetadata:
        SPIFFE_ENDPOINT_SOCKET: unix:///run/spire/sockets/
          agent.sock
    # Enforce mTLS mesh-wide
    enableAutoMtls: true
    # Outbound traffic policy – block any unknown host
    outboundTrafficPolicy:
      mode: REGISTRY_ONLY
    # Access logging to stdout (structured JSON)
    accessLogFile: /dev/stdout
    accessLogFormat: |
      {
        "timestamp": "%START_TIME%",
        "method": "%REQ(:METHOD)%",
        "path": "%REQ(X-ENVOY-ORIGINAL-PATH?:PATH)%",
        "protocol": "%PROTOCOL%",
        "response_code": "%RESPONSE_CODE%",
        "response_flags": "%RESPONSE_FLAGS%",
        "bytes_received": "%BYTES_RECEIVED%",
        "bytes_sent": "%BYTES_SENT%",
        "duration_ms": "%DURATION%",

```

```

    "upstream_service_time": "%RESP(X-ENVOY-UPSTREAM-SERVICE-
    TIME)%",
    "x_forwarded_for": "%REQ(X-FORWARDED-FOR)%",
    "user_agent": "%REQ(USER-AGENT)%",
    "request_id": "%REQ(X-REQUEST-ID)%",
    "authority": "%REQ(:AUTHORITY)%",
    "upstream_host": "%UPSTREAM_HOST%",
    "source_principal": "%CONNECTION_MTLS_PEER_SAN_URI_SANS%"
  }
# Tracing (Jaeger)
tracing:
  zipkin:
    address: jaeger-collector.observability:9411
    sampling: 10.0
components:
  pilot:
    k8s:
      hpaSpec:
        minReplicas: 2
        maxReplicas: 5
  ingressGateways:
    - name: istio-ingressgateway
      enabled: true
      k8s:
        hpaSpec:
          minReplicas: 3
          maxReplicas: 10
        service:
          type: LoadBalancer
          ports:
            - port: 443
              targetPort: 8443
              name: https
            - port: 15443
              targetPort: 15443
              name: tls

```

1.4.2 PeerAuthentication — STRICT mTLS

All namespaces enforce STRICT mTLS — no plaintext traffic is accepted:

```

# Base policy – apply to the entire mesh
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:

```

```
    name: default
    namespace: istio-system
spec:
  mtls:
    mode: STRICT
---
# Namespace-specific policies (belt and suspenders)
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
  namespace: default
spec:
  mtls:
    mode: STRICT
---
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
  namespace: vault
spec:
  mtls:
    mode: STRICT
---
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
  namespace: pki
spec:
  mtls:
    mode: STRICT
---
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
  namespace: audit
spec:
  mtls:
    mode: STRICT
```

1.4.3 AuthorizationPolicy — Least-Privilege per Service

Each service declares the exact source principals allowed to call it:

```
# auth-service: only the API gateway and internal services may
  call it
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: auth-service-authz
  namespace: default
spec:
  selector:
    matchLabels:
      app: auth-service
  action: ALLOW
  rules:
    - from:
      - source:
          principals:
            - "cluster.local/ns/istio-system/sa/istio-
              ingressgateway-service-account"
            - "spiffe://helixterminator.io/ns/default/sa/api-
              gateway"
      to:
      - operation:
          methods: ["POST", "GET"]
          paths:
            - "/v1/auth/*"
            - "/v1/token/*"
            - "/health"
    ---
# vault-service: only auth-service, API gateway, and sync service
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: vault-service-authz
  namespace: vault
spec:
  selector:
    matchLabels:
      app: vault-service
  action: ALLOW
  rules:
    - from:
      - source:
```

```

    principals:
      - "spiffe://helixterminator.io/ns/default/sa/api-
gateway"
      - "spiffe://helixterminator.io/ns/default/sa/auth-
service"
      - "spiffe://helixterminator.io/ns/vault/sa/sync-service"
  to:
  - operation:
      methods: ["GET", "POST", "PUT", "DELETE", "PATCH"]
      paths:
        - "/v1/vaults/*"
        - "/v1/items/*"
---
# PKI CA: only auth-service and session-manager may request
certificates
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: pki-ca-authz
  namespace: pki
spec:
  selector:
    matchLabels:
      app: pki-ca
  action: ALLOW
  rules:
  - from:
    - source:
        principals:
          - "spiffe://helixterminator.io/ns/default/sa/auth-
service"
          - "spiffe://helixterminator.io/ns/default/sa/session-
manager"
    to:
    - operation:
        methods: ["POST"]
        paths:
          - "/v1/sign/user"
          - "/v1/sign/host"
---
# Deny-all default: all namespaces have an explicit deny-all that
# the above ALLOW rules selectively override.
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: deny-all

```

```
namespace: default
spec: {} # Empty spec = deny all traffic
```

1.4.4 Certificate Rotation — 24-Hour TTL

Istio certificates (used for the sidecar proxy mTLS) are configured for 24-hour TTL with automatic rotation:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: istio
  namespace: istio-system
data:
  mesh: |
    defaultConfig:
      # Certificate lifetime for workload certificates
      proxyMetadata:
        SECRET_TTL: "24h"
        SECRET_GRACE_PERIOD_RATIO: "0.5" # Rotate at 12h (50% of
        TTL)
```

The SPIRE-issued SVIDs (1h TTL) rotate more aggressively than Istio's internal certificates. In production, all service-to-service TLS uses SPIRE-issued SVIDs through the SPIFFE CSI driver.

1.5 Identity Hierarchy

HelixTerminator recognizes four distinct identity types, each with different trust characteristics and lifecycle properties:

1.5.1 User Identity (OIDC JWT)

Issued upon successful authentication. Bound to a human user.

Identity attributes:

```
sub:      Stable user UUID (never changes, even after email
change)
email:    User email (mutable, used for display only)
org_id:   Organization UUID
teams:    []string – team UUIDs the user belongs to
roles:    []string – effective roles (computed from RBAC)
amr:      []string – authentication methods used (pwd, otp, hwk)
```

acr: string – authentication context class reference
device_id: UUID of the authenticated device

Lifetime: Access token 15 minutes, refresh token 30 days (rotating).

Binding: Tied to a device and session. Cross-device re-authentication required.

1.5.2 Device Identity (Device Certificate)

Issued upon successful device enrollment. Bound to a specific device instance.

Device Certificate fields:

Subject: CN=<device-uuid>, O=HelixTerminator, OU=<org-id>
SAN: URI: spiffe://helixterminator.io/device/<device-uuid>
URI: spiffe://helixterminator.io/org/<org-id>/device/<device-uuid>
Extensions:
DeviceOS: macOS 14.5 / Windows 11 / Linux (kernel 6.x)
DeviceName: User-assigned friendly name
EnrollmentDate: ISO 8601 timestamp
Compliance: boolean (MDM-reported)
TTL: 90 days
Issuer: HelixTerminator Device CA

Device certificates are stored in the platform's secure enclave (macOS Keychain / Windows CNG / Linux TPM) and are non-exportable. The private key never leaves the device.

1.5.3 Workload Identity (SPIFFE SVID)

Issued by SPIRE to each service instance. Machine-to-machine only.

SVID fields:

Subject: SPIFFE URI SAN
SAN: URI: spiffe://helixterminator.io/ns/<ns>/sa/<service>
Not Before: issuance time
Not After: issuance time + TTL (1h)
Issuer: CN=SPIRE Intermediate CA, O=helixterminator.io
Key Usage: Digital Signature, Key Encipherment
Extended KU: TLS Web Client Authentication, TLS Web Server Authentication

1.5.4 Session Identity (Short-Lived Session Token)

Issued when a user starts an SSH session. Ties together user, device, and target host.

```
{
  "session_id": "sess_01HQZ7VBKF2MXKPTCG4N5Y8R3D",
  "user_id":    "usr_01HQZ7VBKF2MXKPTCG4N5Y8R00",
  "device_id":  "dev_01HQZ7VBKF2MXKPTCG4N5Y8R01",
  "host_id":    "host_01HQZ7VBKF2MXKPTCG4N5Y8R02",
  "org_id":     "org_01HQZ7VBKF2MXKPTCG4N5Y8R03",
  "issued_at":  "2026-06-28T17:00:00Z",
  "expires_at": "2026-06-29T01:00:00Z",
  "cert_serial": "7F3A9C2E1B4D5F6A8C0E",
  "recording_required": true,
  "approval_id": null,
  "ttl": 28800
}
```

Lifetime: Matches the SSH certificate TTL (8 hours for standard user sessions, configurable down to 30 minutes for high-privilege access).

1.6 Continuous Verification

Zero Trust requires that authentication is not a one-time event at login. HelixTerminator implements continuous verification across the following triggers:

Trigger	Action
Policy change (RBAC update, org policy change)	All active sessions validated against new policy within 60 seconds; non-compliant sessions terminated
Device compliance failure (MDM reports non-compliant)	All sessions from that device invalidated immediately; refresh tokens revoked
IP reputation change (UEBA flags anomalous location)	Step-up authentication required before next action
Credential compromise signal (HaveIBeenPwned, threat intelligence)	Forced password reset; all refresh tokens revoked
Inactivity timeout	

Trigger	Action
	Access token re-validation required after 30 minutes of inactivity
Time-based policy (access window expires)	Session terminated; SSH certificate revoked via CRL update
Failed MFA (3 consecutive failures)	Account temporarily locked (30 minutes); admin notified

The continuous verification subsystem is implemented as a policy evaluation service that subscribes to policy change events via Kafka and evaluates active sessions:

```
// pkg/continuousverify/verifier.go
package continuousverify

import (
    "context"
    "sync"
    "time"

    "github.com/helixterm/platform/pkg/session"
    "github.com/helixterm/platform/pkg/policy"
    "github.com/helixterm/platform/pkg/audit"
    "go.uber.org/zap"
)

// Verifier continuously evaluates active sessions against current
// policy.
type Verifier struct {
    sessions    session.Store
    policies    policy.Engine
    auditor     audit.Logger
    log         *zap.Logger
    mu          sync.Mutex
}

// OnPolicyChange is called when any policy changes. It evaluates
// all active
// sessions against the new policy and terminates those that are
// no longer
// authorized.
func (v *Verifier) OnPolicyChange(ctx context.Context, event
    policy.ChangeEvent) error {
```

```

sessions, err := v.sessions.ActiveSessions(ctx)
if err != nil {
    return err
}

for _, sess := range sessions {
    allowed, reason, err := v.policies.Evaluate(ctx,
        policy.Request{
            Subject:    sess.UserID,
            Resource:  policy.Resource{Type: "session", ID:
                sess.SessionID},
            Action:    "maintain",
            Context: policy.Context{
                DeviceID: sess.DeviceID,
                IPAddress: sess.RemoteAddr,
                Time:      time.Now(),
            },
        })
    if err != nil {
        v.log.Error("policy evaluation failed",
            zap.Error(err), zap.String("session_id", sess.SessionID))
        continue
    }
    if !allowed {
        v.log.Warn("terminating session due to policy change",
            zap.String("session_id", sess.SessionID),
            zap.String("user_id", sess.UserID),
            zap.String("reason", reason),
        )
        if err := v.sessions.Terminate(ctx, sess.SessionID,
            reason); err != nil {
            v.log.Error("failed to terminate session",
                zap.Error(err))
        }
        v.auditor.Log(ctx, audit.Event{
            Type:      audit.EventSessionTerminatedByPolicy,
            SessionID: sess.SessionID,
            UserID:    sess.UserID,
            Reason:    reason,
        })
    }
}

return nil
}

```

1.7 Network Micro-Segmentation

Every service is isolated in its own Kubernetes namespace with a default-deny NetworkPolicy. Egress is explicitly allowed only to known service endpoints.

```
# Default deny all ingress and egress for all pods in a namespace
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-all
  namespace: vault
spec:
  podSelector: {}
  policyTypes:
    - Ingress
    - Egress
---
# vault-service: allow ingress from the service mesh only (via
  sidecar)
# and egress to its database and the audit service
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: vault-service-netpol
  namespace: vault
spec:
  podSelector:
    matchLabels:
      app: vault-service
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - namespaceSelector:
            matchLabels:
              istio-injection: enabled
      ports:
        - protocol: TCP
          port: 8080
  egress:
    # Vault PostgreSQL
    - to:
        - namespaceSelector:
```

```
      matchLabels:
        name: data
    ports:
      - protocol: TCP
        port: 5432
# Audit service
- to:
  - namespaceSelector:
      matchLabels:
        name: audit
    ports:
      - protocol: TCP
        port: 9090
# DNS
- to:
  - namespaceSelector: {}
    ports:
      - protocol: UDP
        port: 53
# SPIRE agent socket (localhost)
- to:
  - podSelector:
      matchLabels:
        app: spire-agent
    ports:
      - protocol: TCP
        port: 8081
```

2. Authentication Architecture

2.1 Email + Password Authentication

2.1.1 Password Hashing — Argon2id

All passwords are hashed using **Argon2id** (RFC 9106), the winner of the Password Hashing Competition and the current best-practice recommendation for password storage.

Parameters: - Algorithm: Argon2id - Time (iterations): 3 - Memory: 65,536 KiB (64 MB) - Parallelism: 4 threads - Output length: 32 bytes - Salt: 16 bytes (crypto/rand)

These parameters meet the OWASP recommendation and are calibrated to take approximately 500ms–1s on modern server hardware, making brute-force attacks prohibitively expensive.

```
// pkg/auth/password.go
package auth

import (
    "crypto/rand"
    "crypto/subtle"
    "encoding/base64"
    "fmt"
    "strings"

    "golang.org/x/crypto/argon2"
)

const (
    argon2Time      = 3
    argon2Memory    = 64 * 1024 // 64 MB
    argon2Threads   = 4
    argon2KeyLen    = 32
    argon2SaltLen   = 16
)

// HashPassword hashes the given password using Argon2id and
// returns a
// PHC-formatted string suitable for storage.
func HashPassword(password string) (string, error) {
    salt := make([]byte, argon2SaltLen)
    if _, err := rand.Read(salt); err != nil {
        return "", fmt.Errorf("failed to generate salt: %w", err)
    }

    hash := argon2.IDKey(
        []byte(password),
        salt,
        argon2Time,
        argon2Memory,
        argon2Threads,
        argon2KeyLen,
    )

    // PHC format: $argon2id$v=19$m=65536,t=3,p=4$<salt>$<hash>
    return fmt.Sprintf(
```

```

        "$argon2id$v=%d$m=%d,t=%d,p=%d$%s$%s",
        argon2.Version,
        argon2Memory,
        argon2Time,
        argon2Threads,
        base64.RawStdEncoding.EncodeToString(salt),
        base64.RawStdEncoding.EncodeToString(hash),
    ), nil
}

// VerifyPassword verifies a plaintext password against a PHC-
// formatted Argon2id hash.
// Uses constant-time comparison to prevent timing attacks.
func VerifyPassword(password, encoded string) (bool, error) {
    params, salt, hash, err := parseArgon2Hash(encoded)
    if err != nil {
        return false, err
    }

    computed := argon2.IDKey(
        []byte(password),
        salt,
        params.time,
        params.memory,
        params.threads,
        params.keyLen,
    )

    // Constant-time comparison – prevents timing oracle attacks
    return subtle.ConstantTimeCompare(hash, computed) == 1, nil
}

type argon2Params struct {
    time, memory uint32
    threads      uint8
    keyLen       uint32
}

func parseArgon2Hash(encoded string) (*argon2Params, []byte,
    []byte, error) {
    parts := strings.Split(encoded, "$")
    if len(parts) != 6 {
        return nil, nil, nil, fmt.Errorf("invalid hash format:
        expected 6 parts, got %d", len(parts))
    }
}

```

```

var version int
if _, err := fmt.Sscanf(parts[2], "v=%d", &version); err !=
    nil {
    return nil, nil, nil, fmt.Errorf("invalid version: %w",
        err)
}
if version != argon2.Version {
    return nil, nil, nil, fmt.Errorf("unsupported argon2
        version: %d", version)
}

var params argon2Params
if _, err := fmt.Sscanf(parts[3], "m=%d,t=%d,p=%d",
    &params.memory, &params.time, &params.threads); err !=
    nil {
    return nil, nil, nil, fmt.Errorf("invalid parameters:
        %w", err)
}


salt, err := base64.RawStdEncoding.DecodeString(parts[4])
if err != nil {
    return nil, nil, nil, fmt.Errorf("invalid salt: %w", err)
}

hash, err := base64.RawStdEncoding.DecodeString(parts[5])
if err != nil {
    return nil, nil, nil, fmt.Errorf("invalid hash: %w", err)
}

params.keyLen = uint32(len(hash))
return &params, salt, hash, nil
}

```

2.1.2 Login Flow



```

|                                     |-- mTLS + AuthN -----
>|                                     |
|                                     |                               |-- lookup
user --->|
|                                     |                               |<-- user
record ---|
|                                     |
|                                     |
|                                     |                               |-- verify
Argon2id |
|                                     |                               |   hash
(500ms)  |
|                                     |
|                                     |
|                                     |                               |-- check
lockout --|
|                                     |
|                                     |
|<-- 200 {mfa_required: true}--|<-- mfa_challenge
----|
|   {challenge_id, methods}   |
|                               |
|                               |
|                               |
|-- POST /v1/auth/mfa ----->|
|                               |
|   {challenge_id, totp_code}  |-- mTLS -----
>|                               |
|                               |                               |-- verify
TOTP ----|
|                               |
|                               |
|<-- 200 {access_token,      |<-- tokens
-----|
|       refresh_token,      |
|                               |
|       device_token}       |
|                               |

```

Login rate limiting: - 5 failed attempts per 15 minutes per IP →
temporary block - 10 failed attempts per 24 hours per account →
account lock + admin alert - Exponential backoff response: 200ms →
400ms → 800ms → 1600ms

2.2 TOTP Authentication (RFC 6238)

2.2.1 TOTP Implementation

```
// pkg/auth/totp.go
package auth

import (
    "crypto/hmac"
    "crypto/rand"
    "crypto/sha1"
    "encoding/base32"
    "encoding/binary"
    "fmt"
    "math"
    "time"
)

const (
    totpWindow = 1    // ±1 step tolerance (allows for clock
                      drift)
    totpDigits = 6
    totpStep   = 30   // seconds per TOTP step
    secretBytes = 20  // 160-bit TOTP secret
)

// GenerateTOTPSecret generates a new cryptographically random
// TOTP secret.
func GenerateTOTPSecret() (string, error) {
    secret := make([]byte, secretBytes)
    if _, err := rand.Read(secret); err != nil {
        return "", fmt.Errorf("failed to generate TOTP secret:
        %w", err)
    }
    return base32.StdEncoding.EncodeToString(secret), nil
}

// TOTPProvisioningURI returns the provisioning URI for
// authenticator apps.
func TOTPProvisioningURI(secret, account, issuer string) string {
    return fmt.Sprintf(
        "otpauth://totp/%s:%s?
        secret=%s&issuer=%s&algorithm=SHA1&digits=6&period=30",
        issuer, account, secret, issuer,
    )
}
```

```

// VerifyTOTP verifies a TOTP code against the given secret,
// allowing for
// ±1 step drift to compensate for clock skew.
func VerifyTOTP(secret, code string) (bool, error) {
    decoded, err := base32.StdEncoding.DecodeString(secret)
    if err != nil {
        return false, fmt.Errorf("invalid TOTP secret encoding:
        %w", err)
    }

    counter := time.Now().Unix() / int64(totpStep)

    for offset := -totpWindow; offset <= totpWindow; offset++ {
        expected := generateHOTP(decoded,
            uint64(counter+int64(offset)), totpDigits)
        if expected == code {
            return true, nil
        }
    }
    return false, nil
}

// generateHOTP generates an HMAC-based OTP per RFC 4226.
func generateHOTP(secret []byte, counter uint64, digits int)
    string {
    // Step 1: Generate HMAC-SHA1
    msg := make([]byte, 8)
    binary.BigEndian.PutUint64(msg, counter)

    mac := hmac.New(sha1.New, secret)
    mac.Write(msg)
    h := mac.Sum(nil)

    // Step 2: Dynamic truncation
    offset := h[len(h)-1] & 0x0f
    code := (uint32(h[offset]&0x7f) << 24) |
        (uint32(h[offset+1]) << 16) |
        (uint32(h[offset+2]) << 8) |
        uint32(h[offset+3])

    // Step 3: Compute OTP
    otp := code % uint32(math.Pow10(digits))
    return fmt.Sprintf("%0*d", digits, otp)
}

```

TOTP Backup Codes: - 10 single-use backup codes, 10 characters each -
Stored as Argon2id hashes (same parameters as password hashing) -
Each code invalidated after single use - User must re-generate all codes
at once (regeneration invalidates remaining codes)

2.3 FIDO2/WebAuthn


2.3.1 Architecture Overview

FIDO2 is the strongest authentication factor supported by HelixTerminator. It is: - **Phishing-resistant:** the challenge is cryptographically bound to the origin (`app.helixterminator.io`), preventing replay across origins - **Credential-bound:** the private key never leaves the authenticator - **Biometric capable:** supports platform authenticators (TouchID, FaceID, Windows Hello) and roaming authenticators (YubiKey)

Relying Party configuration:

```
{
  "rpId": "helixterminator.io",
  "rpName": "HelixTerminator",
  "origin": "https://app.helixterminator.io",
  "attestationConveyancePreference": "direct",
  "authenticatorSelection": {
    "authenticatorAttachment": "cross-platform",
    "requireResidentKey": true,
    "residentKey": "preferred",
    "userVerification": "required"
  },
  "timeout": 120000
}
```

2.3.2 WebAuthn Registration Flow



```

----->|
|                                     |<-- challenge stored
-----|
|<-- 200
{
|     challenge,
|
|     rpId: "helixterminator.io",
|
|     userId,
|
|     excludeCredentials: [...]
|
| }
-----|
|
|
|     |-- [User performs biometric
|
|         or YubiKey touch]
|
|     |-- Authenticator creates
|
|         key pair, signs clientData
|
|
|
|     |-- POST /v1/webauthn/register/
complete
|
{
|     id: credentialId,
|
|     rawId: ...,
|
|     response:
{
|         attestationObject,
|
|         clientDataJSON
|
|     },
|
|     type: "public-key"

```

```

|
|      |  } -----
>|
|
|                                     |-- verify attestation
----->|
|                                     | 1. Check challenge
matches |
|                                     | 2. Verify origin
bound   |
|                                     | 3. Verify user
presence |
|                                     | 4. Verify user
verification |
|                                     | 5. Parse CBOR
attestation  |
|                                     | 6. Verify attestation
stmt |
|                                     | 7. Extract public
key   |
|                                     |-- store credential
----->|
|<-- 200 {credentialId, backed |<-- stored
-----|
|           up: true/false}
|
|

```

2.3.3 WebAuthn Authentication Flow

```

Client (Browser/App)      API Backend
Database
|
|
|      |-- POST /v1/webauthn/authenticate/
begin |
|      {userId or null for
|
|      discoverable credential}
|
|      {userVerification: required}
|
|
|
|
|      |<-- 200
{
|
|

```

```

        |     challenge,
|
|     timeout: 120000,
|
|     rpId: "helixterminator.io",
|
|     allowCredentials: [...]
|
|   }
|
|
|
|   |-- [User biometric/touch]
|
|   |-- Authenticator signs
|
|     authenticatorData +
hash(   |
|     clientDataJSON) with privkey
|
|
|
|   |-- POST /v1/webauthn/authenticate/
complete
|
{
|     id: credentialId,
|
|     response:
{
|         authenticatorData,
|
|         clientDataJSON,
|
|         signature,
|
|         userHandle
|
|     }
|
|   }
|
|
|
|
|

```

```

|                                     |-- lookup credential
----->|
|                                     |<-- credential (pubkey,
cntr) -|
|                                     |--
verify:                                |
|                                     | 1. rpIdHash
matches          |
|                                     | 2. UP flag
set              |
|                                     | 3. UV flag set
(required)      |
|                                     | 4. signCount >
stored          |
|                                     | 5. signature
over            |
|                                     | authData ||
hash(cdj)       |
|                                     | 6. verify with stored
pubkey|
|                                     |-- update signCount
----->|
|<-- 200 {access_token,
|
| refresh_token}
|

```

2.3.4 Resident Keys (Discoverable Credentials)

HelixTerminator supports WebAuthn resident keys (FIDO2 discoverable credentials), which allow passwordless login without first entering an email address:

```

// pkg/auth/webauthn.go
package auth

import (
    "github.com/go-webauthn/webauthn/webauthn"
)

var WebAuthnConfig = &webauthn.Config{
    RPDisplayName: "HelixTerminator",
    RPID:          "helixterminator.io",
    RPOrigins:     []string{"https://app.helixterminator.io"},

```

```

// Attestation: "direct" to verify authenticator model for
// enterprise
// policy enforcement (e.g., require YubiKey 5 series)
AttestationPreference: "direct",

// Timeout: 2 minutes to accommodate slow hardware tokens
Timeouts: webauthn.TimeoutsConfig{
    Login: webauthn.TimeoutConfig{
        Enforce:    true,
        Timeout:    120000,
    },
    Registration: webauthn.TimeoutConfig{
        Enforce:    true,
        Timeout:    120000,
    },
},
}

```

Supported authenticators: - YubiKey 5 series (NFC, USB-A, USB-C, Nano) - Apple Touch ID / Face ID (platform authenticator, macOS/iOS) - Windows Hello (Face, Fingerprint, PIN — platform authenticator) - Google Titan Security Keys - Any FIDO2-compliant authenticator

2.4 Push Notifications (Duo Security Integration)

For organizations that have deployed Duo Security, HelixTerminator integrates the Duo Auth API for push-based MFA.

```

// pkg/auth/duo.go
package auth

import (
    "context"
    "fmt"
    "net/http"
    "time"

    duoapi "github.com/duosecurity/duo_api_golang"
)

type DuoClient struct {
    api *duoapi.DuoApi
}

```

```

func NewDuoClient(ikey, skey, host string) *DuoClient {
    return &DuoClient{
        api: duoapi.NewDuoApi(ikey, skey, host, "HelixTerminator/
        1.0",
            duoapi.SetTimeout(30*time.Second),
        ),
    }
}

// SendPush sends a Duo push notification and blocks until the
// user approves
// or denies, or until the context is cancelled.
func (d *DuoClient) SendPush(ctx context.Context, username
    string) (bool, error) {
    params := url.Values{
        "username": []string{username},
        "factor":   []string{"push"},
        "device":   []string{"auto"},
        "type":     []string{"HelixTerminator Login"},
        "pushinfo": []string{fmt.Sprintf("Application=HelixTerminator&Reason=SSH+Access")},
    }

    _, body, err := d.api.SignedCall(http.MethodPost, "/auth/v2/
    auth", params, duoapi.UseTimeout)
    if err != nil {
        return false, fmt.Errorf("duo push failed: %w", err)
    }

    var result struct {
        Stat      string `json:"stat"`
        Response struct {
            Result string `json:"result"`
            Status string `json:"status"`
        } `json:"response"`
    }

    if err := json.Unmarshal(body, &result); err != nil {
        return false, fmt.Errorf("failed to parse Duo response:
        %w", err)
    }

    return result.Response.Result == "allow", nil
}

```

2.5 SSO/Federation

2.5.1 SAML 2.0 Service Provider Implementation

HelixTerminator acts as a SAML 2.0 Service Provider (SP). It does not act as an IdP.

SP Metadata endpoint: <https://app.helixterminator.io/saml/metadata> **ACS endpoint:** <https://app.helixterminator.io/saml/acs> **SLO endpoint:** <https://app.helixterminator.io/saml/slo>

```
// pkg/auth/saml.go
package auth

import (
    "crypto/rsa"
    "crypto/tls"
    "net/url"

    "github.com/crewjam/saml/samlsp"
)

func NewSAMLMiddleware(spKeyPair tls.Certificate, idpMetadataURL
    string) (*samlsp.Middleware, error) {
    rootURL, _ := url.Parse("https://app.helixterminator.io")
    idpMetadataURL_parsed, _ := url.Parse(idpMetadataURL)

    idpMetadata, err :=
        samlsp.FetchMetadata(context.Background(),
            http.DefaultClient,
            *idpMetadataURL_parsed)
    if err != nil {
        return nil, fmt.Errorf("failed to fetch IdP metadata:
            %w", err)
    }

    opts := samlsp.Options{
        URL:          *rootURL,
        Key:          spKeyPair.PrivateKey.(*rsa.PrivateKey),
        Certificate:  spKeyPair.Leaf,
        IDPMetadata: idpMetadata,
        // Require signed assertions AND responses
        AllowIDPInitiated: false,
        // Request AuthnContextClassRef:
        PasswordProtectedTransport
        // (IdP must enforce MFA for elevated contexts)
    }
```

```

        AuthnNameIDFormat: saml.EmailAddressNameIDFormat,
    }

    return samlsp.New(opts)
}

```

Required SAML attributes (must be asserted by IdP): - email (NameID or attribute) - firstName - lastName - groups (for JIT group assignment)

2.5.2 OIDC Service Provider

```

// pkg/auth/oidc.go
package auth

import (
    "context"
    "github.com/coreos/go-oidc/v3/oidc"
    "golang.org/x/oauth2"
)

type OIDCProvider struct {
    provider      *oidc.Provider
    oauth2Config  oauth2.Config
    verifier      *oidc.IDTokenVerifier
}

func NewOIDCProvider(ctx context.Context, issuerURL, clientID,
    clientSecret, redirectURL string) (*OIDCProvider, error) {
    provider, err := oidc.NewProvider(ctx, issuerURL)
    if err != nil {
        return nil,
            fmt.Errorf("failed to create OIDC provider for %s: %w",
                issuerURL, err)
    }

    config := oauth2.Config{
        ClientID:     clientID,
        ClientSecret: clientSecret,
        Endpoint:     provider.Endpoint(),
        RedirectURL:  redirectURL,
        // Request profile, email, and group claims
        Scopes: []string{
            oidc.ScopeOpenID,
            "profile",

```

```

        "email",
        "groups",          // Okta, Azure AD, Auth0
        "offline_access",  // Request refresh token
    },
}

verifier := provider.Verifier(&oidc.Config{
    ClientID: clientID,
    // Require issued-at and expiry to be within 5 minutes of
    // server time
    SkewAllowance: 5 * time.Minute,
})

return &OIDCProvider{
    provider:    provider,
    oauth2Config: config,
    verifier:    verifier,
}, nil
}

// VerifyIDToken validates the ID token and returns the claims.
func (o *OIDCProvider) VerifyIDToken(ctx context.Context,
    rawIDToken string) (*oidc.IDToken, error) {
    token, err := o.verifier.Verify(ctx, rawIDToken)
    if err != nil {
        return nil, fmt.Errorf("ID token verification failed:
        %w", err)
    }
    // Additional checks: nonce validation, acr claims, etc.
    return token, nil
}

```

Supported IdPs and their OIDC discovery endpoints:

IdP	Discovery URL
Okta	https://{tenant}.okta.com/.well-known/openid-configuration
Azure AD	https://login.microsoftonline.com/{tenant}/v2.0/.well-known/openid-configuration
Google Workspace	https://accounts.google.com/.well-known/openid-configuration
Auth0	https://{tenant}.auth0.com/.well-known/openid-configuration
Keycloak	

IdP	Discovery URL
	https://{host}/realms/{realm}/.well-known/openid-configuration

2.5.3 SCIM 2.0 User Provisioning

HelixTerminator exposes a SCIM 2.0 endpoint for automated user provisioning and deprovisioning from enterprise IdPs.

SCIM base URL: https://api.helixterminator.io/scim/v2

Supported operations: - GET /Users — list users - POST /Users — provision new user - GET /Users/{id} — fetch user - PUT /Users/{id} — replace user - PATCH /Users/{id} — update user - DELETE /Users/{id} — deprovision user - GET /Groups — list groups - POST /Groups — create group

Deprovisioning triggers immediate: 1. Session invalidation for all active sessions 2. SSH certificate revocation via CRL delta 3. Refresh token revocation 4. Audit log entry

2.5.4 JIT (Just-In-Time) Provisioning

When a user authenticates via SAML or OIDC for the first time, HelixTerminator can automatically create their account based on IdP-asserted attributes:

```
// pkg/auth/jit.go
package auth

// JITProvisionUser creates or updates a user based on IdP
// assertion data.
// Groups from the IdP are mapped to HelixTerminator teams based
// on the
// org's group-to-team mapping configuration.
func JITProvisionUser(ctx context.Context, assertion
    *SAMLAssertion, orgID string) (*User, error) {
    // Look up org's JIT provisioning config
    orgConfig, err := orgs.GetJITConfig(ctx, orgID)
    if err != nil {
        return nil, err
    }
    if !orgConfig.JITEnabled {
        return nil, ErrJITNotEnabled
    }
}
```

```

// Check if user already exists
existing, err := users.GetByExternalID(ctx, assertion.NameID,
    orgID)
if err != nil && !errors.Is(err, ErrNotFound) {
    return nil, err
}

if existing != nil {
    // Update user attributes from IdP (name, groups, etc.)
    return updateUserFromAssertion(ctx, existing, assertion,
        orgConfig)
}

// Create new user
user := &User{
    Email:      assertion.Email,
    FirstName:  assertion.FirstName,
    LastName:   assertion.LastName,
    ExternalID: assertion.NameID,
    OrgID:      orgID,
    Source:     UserSourceSAML,
    // Default role assigned by JIT config
    Role:       orgConfig.DefaultJITRole,
}

// Map IdP groups to HelixTerminator teams
for _, idpGroup := range assertion.Groups {
    if teamID, ok := orgConfig.GroupMapping[idpGroup]; ok {
        user.TeamIDs = append(user.TeamIDs, teamID)
    }
}

return users.Create(ctx, user)
}

```

2.6 Token Architecture

2.6.1 Access Tokens (JWT, EdDSA/Ed25519, 15-minute TTL)

Access tokens are signed JWTs using **EdDSA (Ed25519)** (RFC 8032/RFC 8037) — the canonical signing algorithm for HelixTerminator access tokens. The signing key is an Ed25519 key pair managed by the Auth service. Public keys are exposed at `https://auth.helixterminator.io/.well-known/jwks.json` for token verification

by resource services. (RS256/ES256/HS256 appear elsewhere in this document only as legacy/negative examples, never as the signing choice.)

Complete JWT payload schema:

```
{
  "iss": "https://auth.helixterminator.io",
  "sub": "usr_01HQZ7VBKF2MXKPTCG4N5Y8R00",
  "aud": ["https://api.helixterminator.io", "https://vault.helixterminator.io"],
  "iat": 1751124000,
  "nbf": 1751124000,
  "exp": 1751124900,
  "jti": "tok_01HQZ7VBKF2MXKPTCG4N5Y8R0A",

  "email": "user@example.com",
  "email_verified": true,
  "name": "Jane Smith",

  "org_id": "org_01HQZ7VBKF2MXKPTCG4N5Y8R03",
  "org_name": "Acme Corp",
  "org_plan": "enterprise",

  "teams": [
    "team_01HQZ7VBKF2MXKPTCG4N5Y8R04",
    "team_01HQZ7VBKF2MXKPTCG4N5Y8R05"
  ],

  "roles": ["member", "team_admin"],
  "permissions": [
    "vault:read",
    "vault:write",
    "host:connect",
    "ssh_key:use",
    "session:start"
  ],

  "device_id": "dev_01HQZ7VBKF2MXKPTCG4N5Y8R01",
  "device_verified": true,

  "session_id": "sess_01HQZ7VBKF2MXKPTCG4N5Y8R0B",

  "amr": ["pwd", "hwk"],
  "acr": "aal2",
```

```

"scope": "openid profile email api:full",

"helixterm": {
  "version": "1",
  "client_version": "3.2.1",
  "client_platform": "macos",
  "ip_at_login": "203.0.113.42",
  "country_at_login": "US",
  "mfa_verified_at": 1751123990,
  "password_changed_at": 1748531990,
  "account_locked": false,
  "requires_mfa": true,
  "vault_access_tier": "personal+team",
  "ssh_cert_ttl": 28800,
  "max_session_duration": 28800
}
}

```

JWT Header:

```

{
  "alg": "EdDSA",
  "typ": "JWT",
  "kid": "helixterm-access-2026-06-key-001"
}

```

2.6.2 Refresh Token Architecture

Refresh tokens are opaque (not JWTs), stored in Redis with a reference to the token binding:

```

// pkg/auth/refresh_token.go
package auth

import (
    "context"
    "crypto/rand"
    "encoding/base64"
    "fmt"
    "time"

    "github.com/redis/go-redis/v9"
)

```

```

const (
    RefreshTokenTTL    = 30 * 24 * time.Hour
    RefreshTokenLength = 64
)

type RefreshTokenStore struct {
    redis *redis.Client
}

type RefreshTokenRecord struct {
    UserID      string    `json:"user_id"`
    DeviceID    string    `json:"device_id"`
    OrgID       string    `json:"org_id"`
    IssuedAt    time.Time `json:"issued_at"`
    ExpiresAt   time.Time `json:"expires_at"`
    // Rotation tracking: the previous token's hash, for detecting
    // theft
    PrevTokenHash string `json:"prev_token_hash,omitempty"`
    // Family ID: if any token in this family is detected as
    // replayed,
    // revoke the entire family (refresh token rotation theft
    // detection)
    FamilyID string `json:"family_id"`
}

// Issue generates a new refresh token and stores it in Redis.
func (s *RefreshTokenStore) Issue(ctx context.Context, record
    RefreshTokenRecord) (string, error) {
    raw := make([]byte, RefreshTokenLength)
    if _, err := rand.Read(raw); err != nil {
        return "", fmt.Errorf("failed to generate refresh token:
            %w", err)
    }
    token := base64.URLEncoding.EncodeToString(raw)
    tokenHash := sha256Hex(token)

    key := fmt.Sprintf("rt:%s", tokenHash)
    data, _ := json.Marshal(record)
    if err := s.redis.Set(ctx, key, data, RefreshTokenTTL).Err();
        err != nil {
        return "", fmt.Errorf("failed to store refresh token:
            %w", err)
    }

    // Also add to user's token family set for bulk revocation
    familyKey := fmt.Sprintf("rt:family:%s", record.FamilyID)
    s.redis.SAdd(ctx, familyKey, tokenHash)
}

```

```

s.redis.Expire(ctx, familyKey, RefreshTokenTTL)

return token, nil
}

// Rotate atomically replaces an old refresh token with a new one.
// If the old token has already been used (replay attack), the
// entire
// token family is revoked.
func (s *RefreshTokenStore) Rotate(ctx context.Context, oldToken
    string) (string, *RefreshTokenRecord, error) {
    oldHash := sha256Hex(oldToken)
    key := fmt.Sprintf("rt:%s", oldHash)

    // Use a Lua script for atomic get-and-delete to prevent
    // replay
    val, err := s.redis.GetDel(ctx, key).Result()
    if err == redis.Nil {
        // Token not found – possible replay attack. Check if it
        // was ever valid.
        // If the token hash exists in any family set, this is a
        // replay.
        // Revoke the entire family.
        s.revokeByTokenHash(ctx, oldHash)
        return "", nil, ErrTokenReplayDetected
    }
    if err != nil {
        return "", nil, fmt.Errorf("failed to retrieve refresh
            token: %w", err)
    }

    var record RefreshTokenRecord
    if err := json.Unmarshal([]byte(val), &record); err != nil {
        return "", nil, fmt.Errorf("failed to parse refresh token
            record: %w", err)
    }

    // Issue a new token in the same family
    record.PrevTokenHash = oldHash
    record.IssuedAt = time.Now()
    record.ExpiresAt = time.Now().Add(RefreshTokenTTL)

    newToken, err := s.Issue(ctx, record)
    return newToken, &record, err
}

```

2.6.3 API Keys

API keys are used for programmatic/service-to-service access where OAuth flows are impractical.

Format: htk_<prefix>_<secret> - Prefix: 8 random alphanumeric characters (visible in UI for identification) - Secret: 32 random bytes, base62-encoded (43 characters) - Example:

htk_a4Bx9Km2_3vMpQrY7nZ2wXcT8dFhJ1kL9bNsE4oR5uA

Storage: - Only the SHA-256 hash of the full key is stored in the database - The prefix is stored in plaintext for UI display - The full key is shown to the user exactly once upon creation

```
CREATE TABLE api_keys (  
  id          UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  prefix      VARCHAR(8) NOT NULL,  
  key_hash    CHAR(64) NOT NULL UNIQUE, -- SHA-256 hex  
  user_id     UUID REFERENCES users(id) ON DELETE  
    CASCADE,  
  org_id      UUID REFERENCES organizations(id) ON  
    DELETE CASCADE,  
  name        TEXT NOT NULL,  
  scopes      TEXT[] NOT NULL DEFAULT '{}',  
  last_used_at TIMESTAMPTZ,  
  expires_at  TIMESTAMPTZ,  
  created_at  TIMESTAMPTZ NOT NULL DEFAULT NOW(),  
  revoked_at  TIMESTAMPTZ,  
  ip_allowlist INET[],  
  CONSTRAINT api_keys_user_or_org CHECK (user_id IS NOT NULL OR  
    org_id IS NOT NULL)  
);  
  
CREATE INDEX api_keys_key_hash_idx ON api_keys (key_hash);  
CREATE INDEX api_keys_user_id_idx ON api_keys (user_id) WHERE  
  revoked_at IS NULL;
```

3. SSH Certificate Authority

3.1 Architecture Overview

HelixTerminator operates its own SSH Certificate Authority (helixterminator.io/services/pki). SSH certificates eliminate the need to distribute public keys to every server — instead, servers trust the CA, and users present certificates signed by that CA. This enables:

- **Centralized revocation:** revoke a CA key to instantly block all associated users
- **Short-lived credentials:** 8-hour certificates expire automatically; no cleanup needed
- **Rich authorization metadata:** certificates carry extensions encoding allowed commands, source addresses, and HelixTerminator-specific policy data
- **MITM prevention:** host certificates cryptographically verify server identity

3.2 CA Hierarchy

Root CA (Offline, HSM-backed)

```
├─ User Intermediate CA (Online, monthly rotation)
|   └─ User Signing Key (Daily rotation)
|       └─ (Historical User Signing Keys, retained for verification)
├─ Host Intermediate CA (Online, monthly rotation)
|   └─ Host Signing Key (Weekly rotation)
|       └─ (Historical Host Signing Keys)
```

3.2.1 Root CA

The Root CA is: - **Offline:** the signing key is stored in an HSM (AWS CloudHSM or YubiHSM 2) and is never accessible over the network - **Long-lived:** 10-year certificate validity - **Ceremony-based:** any operation requiring the Root CA requires a formal key ceremony with multiple signatories (N-of-M quorum) - **Air-gapped:** the Root CA signing machine has no network interface

```
# Root CA generation (performed offline, during key ceremony)
# Generate 4096-bit RSA root CA key
ssh-keygen -t rsa -b 4096 -C "HelixTerminator Root CA" \
-f /secure/offline/helixterm_root_ca
```

```
# Generate Ed25519 root CA (preferred; smaller, faster, equally
    secure)
ssh-keygen -t ed25519 -C "HelixTerminator Root CA v1" \
    -f /secure/offline/helixterm_root_ca_ed25519
```

3.2.2 Intermediate CA

```
// services/pki/internal/ca/intermediate.go
package ca

import (
    "context"
    "crypto/rand"
    "crypto/x509"
    "fmt"
    "time"

    "golang.org/x/crypto/ssh"
)

// IntermediateCA manages the online intermediate certificate
// authority.
type IntermediateCA struct {
    signer    ssh.Signer
    serial    *SerialCounter
    notBefore time.Time
    notAfter  time.Time
}

// SignUserCertificate signs a user SSH certificate with the given
// parameters.
// This is the primary certificate issuance function.
func (ca *IntermediateCA) SignUserCertificate(ctx
    context.Context, req *UserCertRequest) (*ssh.Certificate,
    error) {
    if err := req.Validate(); err != nil {
        return nil, fmt.Errorf("invalid cert request: %w", err)
    }

    pubKey, err := ssh.ParsePublicKey(req.PublicKeyBytes)
    if err != nil {
        return nil, fmt.Errorf("failed to parse public key: %w",
            err)
    }

    serial, err := ca.serial.Next(ctx)
```

```

    if err != nil {
        return nil, fmt.Errorf("failed to get serial number: %w",
            err)
    }

    now := time.Now()
    cert := &ssh.Certificate{
        CertType:    ssh.UserCert,
        Key:         pubKey,
        Serial:      serial,
        KeyId:       fmt.Sprintf("user:%s:session:%s",
            req.UserID, req.SessionID),
        ValidPrincipals: req.Principals,
        ValidAfter:    uint64(now.Unix()),
        ValidBefore:   uint64(now.Add(req.TTL).Unix()),
        Permissions:  ssh.Permissions{
            CriticalOptions: buildCriticalOptions(req),
            Extensions:      buildExtensions(req),
        },
    }

    if err := cert.SignCert(rand.Reader, ca.signer); err != nil {
        return nil, fmt.Errorf("failed to sign certificate: %w",
            err)
    }

    return cert, nil
}

func buildCriticalOptions(req *UserCertRequest) map[string]string
{
    opts := make(map[string]string)

    if req.ForceCommand != "" {
        opts["force-command"] = req.ForceCommand
    }

    if len(req.SourceAddresses) > 0 {
        opts["source-address"] =
            strings.Join(req.SourceAddresses, ",")
    }

    return opts
}

func buildExtensions(req *UserCertRequest) map[string]string {

```

```

exts := map[string]string{
    "permit-pty":          "",
    "permit-user-rc":      "",
}

if req.AllowAgentForwarding {
    exts["permit-agent-forwarding"] = ""
}

if req.AllowX11Forwarding {
    exts["permit-X11-forwarding"] = ""
}

if req.AllowPortForwarding {
    exts["permit-port-forwarding"] = ""
}

// HelixTerminator custom extensions (informational, prefixed
// with custom@)
exts["helixterm-user-id@helixterminator.io"] = req.UserID
exts["helixterm-session-id@helixterminator.io"] =
    req.SessionID
exts["helixterm-org-id@helixterminator.io"] = req.OrgID
exts["helixterm-device-id@helixterminator.io"] = req.DeviceID
exts["helixterm-issued-by@helixterminator.io"] =
    "pki.helixterminator.io/v1"
exts["helixterm-policy-version@helixterminator.io"] =
    req.PolicyVersion

return exts
}

```

3.3 OpenSSH Certificate Format

OpenSSH certificates have a specific binary wire format. Understanding this format is essential for debugging and implementing certificate parsing.

User certificate format (RFC wire encoding):

```

string    "ssh-rsa-cert-v01@openssh.com" (or other algorithm)
string    nonce (random, 32 bytes)
--- key-specific fields (public key blob) ---
uint64    serial
uint32    type (1 = user, 2 = host)
string    key id

```

```

string[] valid principals (comma-separated list as a string list)
uint64   valid after (Unix timestamp)
uint64   valid before (Unix timestamp)
--- critical options ---
string[] critical_options (name-value pairs encoded as string
list)
--- extensions ---
string[] extensions (name-value pairs)
string   reserved (empty)
string   signature key (the CA public key)
string   signature (over all preceding bytes)

```

Inspecting a certificate:

```

# View certificate details
ssh-keygen -L -f /path/to/certificate-cert.pub

# Example output:
# /path/to/certificate-cert.pub:
#      Type: ssh-ed25519-cert-v01@openssh.com user certificate
#      Public key: ED25519-CERT SHA256:aBcDeFg...
#      Signing CA: ED25519 SHA256:xYzAbCd...
#      Key ID: "user:usr_01HQZ...:session:sess_01HQZ..."
#      Serial: 42
#      Valid: from 2026-06-28T17:00:00 to 2026-06-29T01:00:00
#      Principals:
#          alice
#          ubuntu
#          ec2-user
#      Critical Options: (none)
#      Extensions:
#          permit-pty
#          permit-user-rc
#          helixterm-user-id@helixterminator.io
#          helixterm-session-id@helixterminator.io

```

3.4 Certificate Issuance API

```

// services/pki/api/v1/sign.go
package v1

import (
    "context"
    "net/http"

```

```

"time"

"github.com/go-chi/chi/v5"
"helixterminator.io/services/pki/internal/ca"
"helixterminator.io/services/pki/internal/policy"
)

type SignUserCertRequest struct {
    PublicKey    string    `json:"public_key"`    // PEM or
                        authorized_keys format
    UserID       string    `json:"user_id"`
    SessionID    string    `json:"session_id"`
    OrgID        string    `json:"org_id"`
    DeviceID     string    `json:"device_id"`
    Principals   []string  `json:"principals"`    // Unix usernames allowed
    TTL          int       `json:"ttl_seconds"`    // Max 28800 (8h)

    // Policy-driven fields
    ForceCommand string    `json:"force_command,omitempty"`
    SourceAddresses []string `json:"source_addresses,omitempty"`
    AllowForwarding bool      `json:"allow_forwarding"`
}

type SignUserCertResponse struct {
    Certificate string    `json:"certificate"`    // SSH
                        certificate in authorized_keys format
    Serial      uint64    `json:"serial"`
    ValidBefore time.Time `json:"valid_before"`
    ValidAfter  time.Time `json:"valid_after"`
    KeyID       string    `json:"key_id"`
}

func (h *Handler) SignUserCert(w http.ResponseWriter, r
    *http.Request) {
    ctx := r.Context()

    // Caller must be auth-service or session-manager (enforced by
    // Istio AuthorizationPolicy)
    // Verify SPIFFE identity of caller
    callerID := extractSPIFFEID(r)
    if err := h.authorizeCAClient(callerID); err != nil {
        http.Error(w, "forbidden", http.StatusForbidden)
        return
    }
}

```

```

var req SignUserCertRequest
if err := json.NewDecoder(r.Body).Decode(&req); err != nil {
    http.Error(w, "bad request", http.StatusBadRequest)
    return
}

// Enforce policy limits
if req.TTL > policy.MaxUserCertTTL {
    req.TTL = policy.MaxUserCertTTL
}
if req.TTL <= 0 {
    req.TTL = policy.DefaultUserCertTTL
}

certReq := &ca.UserCertRequest{
    PublicKeyBytes: []byte(req.PublicKey),
    UserID:         req.UserID,
    SessionID:      req.SessionID,
    OrgID:          req.OrgID,
    DeviceID:       req.DeviceID,
    Principals:     req.Principals,
    TTL:            time.Duration(req.TTL) * time.Second,
    ForceCommand:   req.ForceCommand,
    SourceAddresses: req.SourceAddresses,
    AllowAgentForwarding: req.AllowForwarding,
}

cert, err := h.userCA.SignUserCertificate(ctx, certReq)
if err != nil {
    h.log.Error("failed to sign user certificate",
        zap.Error(err))
    http.Error(w, "internal error",
        http.StatusInternalServerError)
    return
}

// Log to audit service
h.auditor.Log(ctx, audit.CertIssued{
    Serial:    cert.Serial,
    KeyID:     cert.KeyId,
    UserID:    req.UserID,
    OrgID:     req.OrgID,
    ExpiresAt: time.Unix(int64(cert.ValidBefore), 0),
})

```

```

    resp := SignUserCertResponse{
        Certificate: string(ssh.MarshalAuthorizedKey(cert)),
        Serial:      cert.Serial,
        ValidBefore: time.Unix(int64(cert.ValidBefore), 0),
        ValidAfter:  time.Unix(int64(cert.ValidAfter), 0),
        KeyID:       cert.KeyId,
    }
    json.NewEncoder(w).Encode(resp)
}

```

3.5 Certificate Revocation

3.5.1 CRL (Certificate Revocation List)

A CRL is a signed list of revoked certificate serial numbers. HelixTerminator publishes CRLs via an HTTP endpoint that SSH servers can periodically fetch.

```

// services/pki/internal/revocation/crl.go
package revocation

// RevocationList maintains the list of revoked certificates.
type RevocationList struct {
    db      *sql.DB
    signer  ssh.Signer
}

// Revoke adds a certificate to the revocation list.
func (r *RevocationList) Revoke(ctx context.Context, serial
    uint64, reason RevokeReason) error {
    _, err := r.db.ExecContext(ctx, `
        INSERT INTO revoked_certs (serial, revoked_at, reason)
        VALUES ($1, NOW(), $2)
        ON CONFLICT (serial) DO NOTHING`,
        serial, reason,
    )
    return err
}

// GenerateCRL generates a signed OpenSSH-format KRL (Key
    Revocation List).
// OpenSSH uses KRL format rather than X.509 CRL format.
func (r *RevocationList) GenerateKRL(ctx context.Context)
    ([]byte, error) {
    rows, err := r.db.QueryContext(ctx, `

```

```

        SELECT serial FROM revoked_certs ORDER BY serial`)
    if err != nil {
        return nil, err
    }
    defer rows.Close()

    krl := &ssh.KRL{}
    for rows.Next() {
        var serial uint64
        if err := rows.Scan(&serial); err != nil {
            return nil, err
        }
        krl.RevokeKey(serial)
    }

    return krl.Marshal(rand.Reader, r.signer)
}

```

SSH server sshd_config configuration to use HelixTerminator CRL:

```

# Trust HelixTerminator CA for user authentication
TrustedUserCAKeys /etc/ssh/helixterm_user_ca.pub

# Revocation list – fetched and updated every 5 minutes by a cron
job
RevokedKeys /etc/ssh/helixterm_krl.bin

# AuthorizedPrincipalsFile – maps cert principals to Unix users
AuthorizedPrincipalsFile /etc/ssh/auth_principals/%u

```

3.5.2 OCSP Stapling

For X.509 certificates used in the TLS layer, OCSP stapling is configured to eliminate the client-side OCSP lookup latency:

```

# nginx.conf TLS configuration
ssl_stapling on;
ssl_stapling_verify on;
ssl_trusted_certificate /etc/nginx/ssl/chain.pem;
resolver 8.8.8.8 8.8.4.4 valid=300s;
resolver_timeout 10s;

```

3.6 Host Certificates

Host certificates prevent man-in-the-middle attacks by cryptographically verifying the SSH server's identity. When a user connects to a server through HelixTerminator, the client verifies the host certificate against the HelixTerminator Host CA.

```
# Sign a host certificate (performed by PKI service during host
  enrollment)
ssh-keygen -s /path/to/host_ca_signing_key \
  -I "web-01.prod.acme.helixterminator.io" \
  -h \
  -n "web-01.prod.acme.helixterminator.io,10.0.1.42" \
  -V +52w \      # 1 year validity for host certs
/etc/ssh/ssh_host_ed25519_key.pub
```

Client ~/.ssh/known_hosts configuration:

```
@cert-authority *.helixterminator.io ssh-ed25519 AAAA...
HelixTerminator Host CA
```

3.7 Certificate Transparency Log

All issued certificates are logged to an append-only transparency log, enabling detection of unauthorized certificate issuance.

```
// services/pki/internal/transparency/log.go
package transparency

// CTLogEntry records certificate issuance for audit and
  transparency purposes.
type CTLogEntry struct {
    Timestamp    time.Time `json:"timestamp"`
    Serial       uint64    `json:"serial"`
    KeyID        string     `json:"key_id"`
    UserID       string     `json:"user_id"`
    OrgID        string     `json:"org_id"`
    Principals   []string  `json:"principals"`
    ValidBefore  time.Time `json:"valid_before"`
    ValidAfter   time.Time `json:"valid_after"`
    Fingerprint  string     `json:"fingerprint"` // SHA-256 of the
                        public key
    CAKeyID      string     `json:"ca_key_id"`
    PrevHash     string     `json:"prev_hash"`   // SHA-256 of
                        previous entry (hash chain)
    Hash         string     `json:"hash"`         // SHA-256 of this
                        entry
```

```

}

// Append adds a new entry to the CT log and updates the hash
// chain.
func (l *CTLogger) Append(ctx context.Context, cert
    *ssh.Certificate, userID, orgID string) error {
    prevHash, err := l.getLastHash(ctx)
    if err != nil {
        return err
    }

    entry := CTLogEntry{
        Timestamp:    time.Now(),
        Serial:       cert.Serial,
        KeyID:        cert.KeyId,
        UserID:       userID,
        OrgID:        orgID,
        Principals:   cert.ValidPrincipals,
        ValidBefore:  time.Unix(int64(cert.ValidBefore), 0),
        ValidAfter:   time.Unix(int64(cert.ValidAfter), 0),
        Fingerprint:  ssh.FingerprintSHA256(cert.Key),
        CAKeyID:      ssh.FingerprintSHA256(cert.SignatureKey),
        PrevHash:     prevHash,
    }

    data, _ := json.Marshal(entry)
    entry.Hash = sha256Hex(data)

    return l.writeEntry(ctx, entry)
}

```

4. Vault Encryption Architecture

4.1 Zero-Knowledge Design Principle

The HelixTerminator vault is **zero-knowledge** with respect to the server:

The server stores only encrypted data and encrypted keys. The server **cannot** decrypt any vault content, and no HelixTerminator employee can access plaintext user data.

This is achieved through client-side encryption: all encryption and decryption operations happen on the client device. The server only stores and retrieves ciphertext.

Scope of this guarantee (per CANONICAL_FACTS CD-10): zero-knowledge is a **HARD** requirement for *vault items* — client-side encryption/decryption only, with the server relaying ciphertext it cannot open. This guarantee does **not** extend to SSH host credentials under password-based host authentication, nor to server-side SSH key generation (see `07_api_and_database.md` §4 and `POST /keys/generate`): those are a distinct, lower-assurance credential class that requires the server to handle plaintext at connection time. This document, and every sibling document, must stop describing password-based SSH host authentication as zero-knowledge — it explicitly is not.

DEFERRED (next increment): the vault’s zero-knowledge posture stated above is the design target but is not yet fully realized by the implementation shown in §4.6–§4.7 below — see the inline DEFERRED notes there for the specific server-side key-handling gaps (a plaintext `VaultKey` field in a backend-looking struct, and a re-wrap flow that can execute server-side) that must be removed to make this guarantee true rather than aspirational.

4.2 Key Hierarchy

User Password (never stored)

|

▼ Argon2id(password, user_salt, t=3, m=64MB, p=4) → 32 bytes

Master Key (ephemeral, exists only in memory during active session)

|

▼ AES-256-GCM decrypt(Encrypted Master Key, Master Key)

Vault Key (AES-256, one per vault)

|

▼ AES-256-GCM decrypt(Encrypted Item Key, Vault Key)

Item Key (AES-256, one per vault item)

|

▼ AES-256-GCM decrypt(Encrypted Item Data, Item Key)

Plaintext Item Data

Key storage model: | Key | Stored Where | Encrypted With | Notes |

|——|———|———|——| | Master Key | Never stored |

Derived from password | Reconstructed on each unlock | | Vault Key |

Server (encrypted) | Master Key or Device Key | Per-vault | | Item Key

| Server (encrypted) | Vault Key | Per-item | | Item Data | Server
(encrypted) | Item Key | Ciphertext only | | Device Key | Platform
Keychain/KeyStore | Biometric/PIN | For biometric unlock |

4.3 Complete Encryption Implementation

```
// digital.vasic.security/vault/crypto.go
package vault

import (
    "crypto/aes"
    "crypto/cipher"
    "crypto/rand"
    "crypto/sha256"
    "crypto/subtle"
    "encoding/base64"
    "fmt"
    "io"

    "golang.org/x/crypto/argon2"
)

const (
    // Key derivation parameters
    kdTime      = 3
    kdMemory    = 64 * 1024 // 64 MB
    kdThreads   = 4
    kdKeyLen    = 32 // 256-bit key
    kdSaltLen   = 32 // 256-bit salt (longer than password hashing
                    // for KDF use)

    // AES-GCM parameters
    aeadKeyLen   = 32 // AES-256
    aeadNonceLen = 12 // GCM standard nonce
    aeadTagLen   = 16 // GCM authentication tag
)

// EncryptedBlob is the canonical representation of any encrypted
// value
// in the HelixTerminator vault. The format is:
// version(1) || nonce(12) || ciphertext || tag(16)
// This format is opaque and self-contained – no external metadata
// needed.
type EncryptedBlob struct {
    Version byte `json:"v"`
}
```

```

    Nonce      []byte `json:"n"` // base64url
    Ciphertext []byte `json:"c"` // base64url (includes GCM tag)
}

// DeriveKey derives the Master Key from the user's password using
// Argon2id.
// The salt must be stored (it is non-secret).
func DeriveKey(password []byte, salt []byte) ([]byte, error) {
    if len(salt) != kdSaltLen {
        return nil, fmt.Errorf("invalid salt length: expected %d, got %d", kdSaltLen, len(salt))
    }
    return argon2.IDKey(password, salt, kdTime, kdMemory, kdThreads, kdKeyLen), nil
}

// GenerateSalt generates a cryptographically random salt.
func GenerateSalt() ([]byte, error) {
    salt := make([]byte, kdSaltLen)
    if _, err := rand.Read(salt); err != nil {
        return nil, fmt.Errorf("failed to generate salt: %w", err)
    }
    return salt, nil
}

// GenerateKey generates a new random AES-256 key.
func GenerateKey() ([]byte, error) {
    key := make([]byte, aeadKeyLen)
    if _, err := rand.Read(key); err != nil {
        return nil, fmt.Errorf("failed to generate key: %w", err)
    }
    return key, nil
}

// Encrypt encrypts plaintext with AES-256-GCM using the provided
// key.
// Returns an EncryptedBlob.
func Encrypt(key, plaintext []byte) (*EncryptedBlob, error) {
    if len(key) != aeadKeyLen {
        return nil, fmt.Errorf("invalid key length: expected %d bytes (AES-256)", aeadKeyLen)
    }

    block, err := aes.NewCipher(key)
    if err != nil {

```

```

        return nil, fmt.Errorf("failed to create AES cipher: %w",
            err)
    }

    gcm, err := cipher.NewGCM(block)
    if err != nil {
        return nil, fmt.Errorf("failed to create GCM: %w", err)
    }

    nonce := make([]byte, gcm.NonceSize())
    if _, err := rand.Read(nonce); err != nil {
        return nil, fmt.Errorf("failed to generate nonce: %w",
            err)
    }

    // Seal appends the authentication tag to the ciphertext
    ciphertext := gcm.Seal(nil, nonce, plaintext, nil)

    return &EncryptedBlob{
        Version: 1,
        Nonce: nonce,
        Ciphertext: ciphertext,
    }, nil
}

// Decrypt decrypts an EncryptedBlob with AES-256-GCM.
func Decrypt(key []byte, blob *EncryptedBlob) ([]byte, error) {
    if len(key) != aeadKeyLen {
        return nil, fmt.Errorf("invalid key length")
    }

    if blob.Version != 1 {
        return nil, fmt.Errorf("unsupported blob version: %d",
            blob.Version)
    }

    block, err := aes.NewCipher(key)
    if err != nil {
        return nil, fmt.Errorf("failed to create AES cipher: %w",
            err)
    }

    gcm, err := cipher.NewGCM(block)
    if err != nil {
        return nil, fmt.Errorf("failed to create GCM: %w", err)
    }

```

```

plaintext, err := gcm.Open(nil, blob.Nonce, blob.Ciphertext,
    nil)
if err != nil {
    // GCM authentication failure – data tampered or wrong key
    return nil, fmt.Errorf("decryption failed (authentication
        error): %w", err)
}

return plaintext, nil
}

// WrapKey encrypts a key (e.g., Vault Key) with another key
// (e.g., Master Key).
func WrapKey(wrappingKey, keyToWrap []byte) (*EncryptedBlob,
    error) {
    return Encrypt(wrappingKey, keyToWrap)
}

// UnwrapKey decrypts a wrapped key.
func UnwrapKey(wrappingKey []byte, wrapped *EncryptedBlob)
    ([]byte, error) {
    return Decrypt(wrappingKey, wrapped)
}

```

4.4 Multi-Device Key Sharing via X25519 ECDH

When a user registers a second device, the Vault Key is re-wrapped with the new device's public key using X25519 Elliptic Curve Diffie-Hellman key exchange.

```

// digital.vasic.security/vault/multidevice.go
package vault

import (
    "crypto/rand"
    "crypto/sha256"
    "fmt"

    "golang.org/x/crypto/curve25519"
)

// DeviceKeyPair represents an X25519 key pair used for vault key
// wrapping.
type DeviceKeyPair struct {
    PublicKey [32]byte
    PrivateKey [32]byte
}

```

```

}

// GenerateDeviceKeyPair generates a new X25519 key pair for a
// device.
// The private key is stored in the platform secure enclave;
// the public key is uploaded to the server.
func GenerateDeviceKeyPair() (*DeviceKeyPair, error) {
    var privKey [32]byte
    if _, err := rand.Read(privKey[:]); err != nil {
        return nil, fmt.Errorf("failed to generate X25519 private
            key: %w", err)
    }

    pubKey, err := curve25519.X25519(privKey[:],
        curve25519.Basepoint)
    if err != nil {
        return nil, fmt.Errorf("failed to derive X25519 public
            key: %w", err)
    }

    kp := &DeviceKeyPair{}
    copy(kp.PrivateKey[:], privKey[:])
    copy(kp.PublicKey[:], pubKey)
    return kp, nil
}

// WrapVaultKeyForDevice encrypts a vault key for a specific
// device using
// X25519 ECDH + AES-256-GCM.
//
// The ephemeral ECDH key ensures forward secrecy: even if the
// device key
// is compromised later, past vault key ciphertexts remain secure.
func WrapVaultKeyForDevice(devicePublicKey [32]byte, vaultKey
    []byte) (*DeviceWrappedKey, error) {
    // Generate ephemeral sender key pair
    var ephemeralPriv [32]byte
    if _, err := rand.Read(ephemeralPriv[:]); err != nil {
        return nil, err
    }
    ephemeralPub, err := curve25519.X25519(ephemeralPriv[:],
        curve25519.Basepoint)
    if err != nil {
        return nil, err
    }

    // ECDH: shared_secret = X25519(ephemeral_priv, device_pub)

```

```

sharedSecret, err := curve25519.X25519(ephemeralPriv[:],
    devicePublicKey[:])
if err != nil {
    return nil, err
}

// Derive symmetric key from shared secret using HKDF-SHA256
wrappingKey := hkdfDerive(sharedSecret, ephemeralPub,
    []byte("helixterm-vault-key-wrap-v1"))

// Encrypt vault key
encrypted, err := Encrypt(wrappingKey, vaultKey)
if err != nil {
    return nil, err
}

return &DeviceWrappedKey{
    EphemeralPublicKey: ephemeralPub,
    EncryptedKey:       encrypted,
}, nil
}

// UnwrapVaultKeyWithDevice decrypts a device-wrapped vault key
// using the
// device's private key.
func UnwrapVaultKeyWithDevice(devicePrivKey [32]byte, wrapped
    *DeviceWrappedKey) ([]byte, error) {
    // ECDH: shared_secret = X25519(device_priv, ephemeral_pub)
    sharedSecret, err := curve25519.X25519(devicePrivKey[:],
        wrapped.EphemeralPublicKey)
    if err != nil {
        return nil, err
    }

    wrappingKey := hkdfDerive(sharedSecret,
        wrapped.EphemeralPublicKey, []byte("helixterm-vault-key-
        wrap-v1"))

    return Decrypt(wrappingKey, wrapped.EncryptedKey)
}

// hkdfDerive performs HKDF-SHA256 key derivation.
func hkdfDerive(ikm, salt, info []byte) []byte {
    h := hkdf.New(sha256.New, ikm, salt, info)
    key := make([]byte, 32)
    io.ReadFull(h, key)
    return key
}

```

```

}

type DeviceWrappedKey struct {
    EphemeralPublicKey []byte
    `json:"ephemeral_pub"` // 32 bytes, base64
    EncryptedKey        *EncryptedBlob `json:"encrypted_key"`
}

```

4.5 Shamir's Secret Sharing for Multi-Device Recovery

For vault recovery across devices, the Vault Key is split using Shamir's Secret Sharing with a threshold of 2-of-N devices. This means any 2 registered devices can reconstruct the Vault Key, but no single device alone can.

```

// digital.vasic.security/vault/shamir.go
package vault

import (
    "crypto/rand"
    "fmt"

    "github.com/hashicorp/vault/shamir"
)

// SplitVaultKey splits a vault key into N shares with a threshold
// of T.
// At least T shares are required to reconstruct the key.
func SplitVaultKey(vaultKey []byte, threshold, shares int) ([]byte, error) {
    if threshold < 2 {
        return nil, fmt.Errorf("threshold must be at least 2")
    }
    if shares < threshold {
        return nil, fmt.Errorf("shares (%d) must be >= threshold (%d)", shares, threshold)
    }
    return shamir.Split(vaultKey, shares, threshold)
}

// CombineShares reconstructs the vault key from at least T
// shares.
func CombineShares(shares [][]byte) ([]byte, error) {
    return shamir.Combine(shares)
}

```

4.6 Team Vault Key Sharing

```
// digital.vasic.security/vault/team.go
package vault

// TeamVault manages key distribution for shared team vaults.
type TeamVault struct {
    VaultID      string
    VaultKey     []byte    // plaintext vault key (only in
                memory)
    MemberWraps  map[string]*DeviceWrappedKey // user_id ->
                wrapped key
}

// AddMember grants vault access to a new team member by wrapping
// the vault
// key with the member's device public key.
func (tv *TeamVault) AddMember(ctx context.Context, userID
    string, devicePubKey [32]byte) error {
    wrapped, err := WrapVaultKeyForDevice(devicePubKey,
        tv.VaultKey)
    if err != nil {
        return fmt.Errorf("failed to wrap vault key for user %s:
            %w", userID, err)
    }
    tv.MemberWraps[userID] = wrapped
    return saveWrappedKey(ctx, tv.VaultID, userID, wrapped)
}

// RevokeMember removes a member's access to the team vault.
// This is a two-step process:
// 1. Delete their wrapped key (immediate – they can no longer
//    decrypt new data)
// 2. Rotate the vault key (prevents offline decryption of
//    previously synced data)
//
// Step 2 requires that all remaining members re-sync their
// wrapped key,
// which is triggered asynchronously.
func (tv *TeamVault) RevokeMember(ctx context.Context,
    revokedUserID string) error {
    // Step 1: Delete wrapped key
    if err := deleteWrappedKey(ctx, tv.VaultID, revokedUserID);
        err != nil {
        return err
    }
    delete(tv.MemberWraps, revokedUserID)
```

```
// Step 2: Trigger vault key rotation
return scheduleKeyRotation(ctx, tv.VaultID,
    KeyRotationReasonMemberRevoked)
}
```

DEFERRED (next increment): security redesign — TeamVault.VaultKey []byte holds the plaintext vault key in memory inside a vault package with no client-boundary marker, and AddMember/RevokeMember call WrapVaultKeyForDevice/scheduleKeyRotation in a way that reads as server-executable code. Per CANONICAL_FACTS CD-10 (zero-knowledge is HARD for vault items), this must be redesigned so key-wrap and key-combination operations are provably client-side only, with the server acting solely as a relay for opaque ciphertext it cannot open. Not done this increment.

4.7 Key Rotation Procedure

Full key rotation procedure (triggered by member revocation, periodic rotation policy, or manual trigger):

1. Generate new Vault Key (32 random bytes)
2. For each vault item:
 - a. Decrypt Item Key with old Vault Key
 - b. Wrap Item Key with new Vault Key
 - c. Store re-wrapped Item Key (item DATA is not re-encrypted – only the key wrapping changes)
3. For each active member device:
 - a. Wrap new Vault Key with member device public key
 - b. Store new wrapped Vault Key
4. Delete old wrapped Vault Keys
5. Log rotation event to audit log
6. Notify all active sessions of key rotation (triggers client re-sync)

DEFERRED (next increment): security redesign — steps 2a/2b (Item Key decrypt/re-wrap under the new Vault Key) are described without a stated execution boundary and, combined with §4.6's server-callable RevokeMember/scheduleKeyRotation, currently read as server-side operations on decrypted key material. Per CANONICAL_FACTS CD-10, re-wrap must be performed client-side (or via blinded/opaque server relay only); the server must never hold an unwrapped Item Key or Vault Key. Not done this increment.

4.8 Recovery Code Mechanism

Recovery codes provide offline emergency access when all devices are lost:

```
// digital.vasic.security/vault/recovery.go
package vault

// RecoveryKit is generated at account creation and contains
// everything
// needed to recover vault access offline.
type RecoveryKit struct {
    // 12-word BIP39 mnemonic that encodes the recovery key
    Mnemonic string

    // The vault key encrypted with the recovery key
    // AES-256-GCM, key = Argon2id(mnemonic_entropy,
    // recovery_salt)
    EncryptedVaultKey *EncryptedBlob

    // Salt for key derivation from mnemonic
    RecoverySalt []byte

    // Verification hash: SHA-256 of the recovery key (for user to
    // verify
    // they have the right codes without needing the vault key)
    VerificationHash string
}

func GenerateRecoveryKit(vaultKey []byte) (*RecoveryKit, error) {
    // Generate 32 bytes of entropy for the recovery key
    entropy := make([]byte, 32)
    if _, err := rand.Read(entropy); err != nil {
        return nil, err
    }

    // Convert to BIP39 mnemonic (24 words for 256-bit entropy)
    mnemonic, err := bip39.NewMnemonic(entropy)
    if err != nil {
        return nil, err
    }

    // Derive encryption key from mnemonic
    salt := make([]byte, 32)
    rand.Read(salt)
```

```

recoveryKey := argon2.IDKey(entropy, salt, 3, 64*1024, 4, 32)

// Encrypt vault key with recovery key
encryptedVaultKey, err := Encrypt(recoveryKey, vaultKey)
if err != nil {
    return nil, err
}

verificationHash := sha256Hex(recoveryKey)

return &RecoveryKit{
    Mnemonic:      mnemonic,
    EncryptedVaultKey: encryptedVaultKey,
    RecoverySalt:   salt,
    VerificationHash: verificationHash,
}, nil
}

```

5. SSH Key Management Security

5.1 Supported and Rejected Key Types

Approved key types:

Type	Minimum Size	Notes
Ed25519	N/A (fixed 256-bit)	Preferred. Fastest, smallest, most secure
ECDSA P-384	384-bit	Acceptable for FIPS 140-2 compliance
ECDSA P-256	256-bit	Acceptable; weaker than P-384
RSA	4096-bit minimum	Legacy systems only; prefer Ed25519

Rejected key types (will not import or generate):

Type	Reason
DSA	Completely broken; NIST deprecated in 2013
RSA < 2048-bit	Insufficient key length; factoring attacks practical
ECDSA P-192	Removed from NIST SP 800-186
ECDSA secp256k1	Not approved by NIST
RSA 1024-bit	Factored in practice

5.2 Key Generation

```
// digital.vasic.security/sshkeys/generate.go
package sshkeys

import (
    "crypto/ecdsa"
    "crypto/ed25519"
    "crypto/elliptic"
    "crypto/rand"
    "crypto/rsa"
    "fmt"

    "golang.org/x/crypto/ssh"
)

type KeyType int

const (
    KeyTypeEd25519    KeyType = iota
    KeyTypeECDSAP256
    KeyTypeECDSAP384
    KeyTypeRSA4096
)

// GenerateKey generates a new SSH key pair of the specified type.
// All randomness comes from crypto/rand (CSPRNG).
func GenerateKey(keyType KeyType) (ssh.Signer, error) {
    switch keyType {
    case KeyTypeEd25519:
        _, privKey, err := ed25519.GenerateKey(rand.Reader)
        if err != nil {
```

```

        return nil,
        fmt.Errorf("ed25519 key generation failed: %w", err)
    }
    return ssh.NewSignerFromKey(privKey)

case KeyTypeECDSAP256:
    privKey, err := ecdsa.GenerateKey(elliptic.P256(),
    rand.Reader)
    if err != nil {
        return nil, fmt.Errorf("ECDSA P-256 key generation
        failed: %w", err)
    }
    return ssh.NewSignerFromKey(privKey)

case KeyTypeECDSAP384:
    privKey, err := ecdsa.GenerateKey(elliptic.P384(),
    rand.Reader)
    if err != nil {
        return nil, fmt.Errorf("ECDSA P-384 key generation
        failed: %w", err)
    }
    return ssh.NewSignerFromKey(privKey)

case KeyTypeRSA4096:
    privKey, err := rsa.GenerateKey(rand.Reader, 4096)
    if err != nil {
        return nil, fmt.Errorf("RSA-4096 key generation
        failed: %w", err)
    }
    return ssh.NewSignerFromKey(privKey)

default:
    return nil, fmt.Errorf("unsupported key type: %d",
    keyType)
}
}

```

5.3 Passphrase Protection

Private keys are encrypted using the OpenSSH key format v2, which uses bcrypt for key derivation and AES-256-CBC + SHA-512 HMAC for encryption.

```

// digital.vasic.security/sshkeys/protect.go
package sshkeys

import (

```

```

    "crypto/subtle"

    "golang.org/x/crypto/ssh"
)

// EncryptPrivateKey encrypts a private key with a passphrase
// using the
// OpenSSH private key format v2 (the modern format, not the
// legacy PEM format).
func EncryptPrivateKey(signer ssh.Signer, passphrase []byte)
    ([]byte, error) {
    privKey, ok := signer.PublicKey().(ssh.CryptoPublicKey)
    if !ok {
        return nil, fmt.Errorf("signer does not implement
        CryptoPublicKey")
    }

    // MarshalPrivateKeyWithPassphrase uses bcrypt (rounds=16) +
    // AES-256-CTR
    // for the OpenSSH private key format. This is the recommended
    // approach.
    pemBlock, err :=
        ssh.MarshalPrivateKeyWithPassphrase(privKey.CryptoPublicKey(),
        "", passphrase)
    if err != nil {
        return nil, fmt.Errorf("failed to encrypt private key:
        %w", err)
    }

    return pem.EncodeToMemory(pemBlock), nil
}

```

5.4 Biometric Key Unlock via Platform Secure Enclave

On platforms that support it, SSH key material is stored in the platform's secure enclave (hardware-isolated key storage):

```

// digital.vasic.security/platform/keychain_darwin.go
//go:build darwin

package platform

// #cgo LDFLAGS: -framework Security -framework Foundation
// #include <Security/Security.h>
import "C"
import (
    "fmt"

```

```

        "unsafe"
    )

    // StoreKeyInKeychain stores an SSH private key in the macOS
    // Keychain,
    // with access control requiring biometric authentication (Touch
    // ID / Face ID).
    func StoreKeyInKeychain(keyID string, keyData []byte) error {
        label := C.CString("HelixTerminator SSH Key: " + keyID)
        defer C.free(unsafe.Pointer(label))

        // Create access control requiring biometric authentication
        var cfError C.CFErrorRef
        access := C.SecAccessControlCreateWithFlags(
            C.kCFAllocatorDefault,
            C.kSecAttrAccessibleWhenUnlockedThisDeviceOnly,
            C.kSecAccessControlBiometryCurrentSet|
            C.kSecAccessControlPrivateKeyUsage,
            &cfError,
        )
        if access == 0 {
            return fmt.Errorf("failed to create access control")
        }
        defer C.CFRelease(C.CFTypeRef(access))

        // Store in Keychain
        // ... (Keychain API calls)
        return nil
    }
}

```

5.5 SSH Agent Key Forwarding Security

Key forwarding poses significant security risks — a compromised intermediate host could use the forwarded agent to authenticate to other hosts. HelixTerminator enforces:

1. **Explicit confirmation per use:** every authentication request through the forwarded agent requires the user to confirm
2. **Forwarding lifetime limit:** forwarded keys expire after the SSH session ends (not persisted)
3. **Trusted host list:** forwarding is only permitted to hosts tagged as “trusted” in the HelixTerminator host inventory

```

// digital.vasic.security/agent/confirming_agent.go
package agent

```

```

// ConfirmingAgent wraps an ssh.Agent and requires user
    confirmation
// before signing each authentication request.
type ConfirmingAgent struct {
    inner      agent.Agent
    confirmer  Confirmer
    maxSignOps int
    signOps    int
}

func (a *ConfirmingAgent) Sign(key ssh.PublicKey, data []byte)
    (*ssh.Signature, error) {
    if a.maxSignOps > 0 && a.signOps >= a.maxSignOps {
        return nil, fmt.Errorf("maximum sign operations (%d)
            reached for this session", a.maxSignOps)
    }

    // Prompt user for confirmation
    fingerprint := ssh.FingerprintSHA256(key)
    confirmed, err := a.confirmer.Confirm(fmt.Sprintf(
        "Allow SSH authentication using key %s?", fingerprint,
    ))
    if err != nil || !confirmed {
        return nil, fmt.Errorf("agent authentication denied by
            user")
    }

    a.signOps++
    return a.inner.Sign(key, data)
}

```

5.6 Key Audit Trail

Every SSH key use is logged to the audit service:

```

// digital.vasic.security/sshkeys/audit.go
package sshkeys

type KeyUseEvent struct {
    EventType string
    `json:"event_type"` // "sign", "import", "export",
    "generate", "delete"
    KeyID      string    `json:"key_id"`
    KeyFingerprint string `json:"key_fingerprint"` // SHA-256
    UserID     string    `json:"user_id"`
    DeviceID   string    `json:"device_id"`
}

```

```
SessionID    string    `json:"session_id,omitempty"`
TargetHost   string    `json:"target_host,omitempty"`
TargetUser   string    `json:"target_user,omitempty"`
Timestamp    time.Time `json:"timestamp"`
IPAddress    string    `json:"ip_address"`
Success      bool      `json:"success"`
Error        string    `json:"error,omitempty"`
}
```

6. RBAC Model

6.1 Role Definitions

System-Wide Roles

Role	Description
super_admin	Full system access; bypasses org isolation; assigned to HelixTerminator operations team only
org_admin	Full administration of a specific organization: user management, billing, SSO configuration, audit log access
team_admin	Team-scoped administration: invite/remove members, manage team vaults and hosts
member	Standard team member: use vaults, connect to hosts, manage own SSH keys
auditor	Read-only compliance/audit-log access across the organization; cannot create, modify, or delete any resource
api_user	Service account for programmatic API access; no interactive access

DEFERRED (next increment): the canonical 6-role vocabulary (super_admin, org_admin, team_admin, member, auditor, api_user — per CANONICAL_FACTS CD-8) replaces the prior guest role name here. This pass renames the role and grants auditor audit_log read access (§6.2); a full re-derivation of the permission matrix

(removing the old guest’s shared-resource read/connect/start grants, which do not belong to an audit-only role) is deferred to the next increment.

Resource Types and Actions

Resource	Actions
vault	create, read, update, delete, share, list
host	create, read, update, delete, connect, list
group	create, read, update, delete, list, inherit
ssh_key	create, read, update, delete, use, share, export, import
snippet	create, read, update, delete, execute, share
session	start, view, record, terminate, replay
audit_log	view, export
organization	read, update, delete, transfer
team	create, read, update, delete, invite, remove_member
billing	view, manage, export
api_key	create, read, revoke
device	list, revoke, trust

6.2 Permission Matrix

Resource	Action	super_admin	org_admin	team_admin	member	audit
vault	create	✓	✓	✓	✓	×
vault	read	✓	✓	✓	✓ (own/ team)	✓ (share)
vault	update	✓	✓	✓	✓ (own)	×
vault	delete	✓	✓	✓ (team)	✓ (own)	×

Resource	Action	super_admin	org_admin	team_admin	member	audit
vault	share	✓	✓	✓	✓ (own)	×
host	create	✓	✓	✓	✓	×
host	connect	✓	✓	✓	✓	✓ (share)
host	delete	✓	✓	✓ (team)	✓ (own)	×
ssh_key	create	✓	✓	✓	✓	×
ssh_key	use	✓	✓	✓	✓	×
ssh_key	export	✓	✓	✓	✓ (own)	×
session	start	✓	✓	✓	✓	✓ (share)
session	view	✓	✓	✓ (team)	✓ (own)	×
session	record	✓	✓	✓ (team)	×	×
session	terminate	✓	✓	✓ (team)	✓ (own)	×
audit_log	view	✓	✓	×	×	✓
audit_log	export	✓	✓	×	×	×
organization	update	✓	✓	×	×	×
team	invite	✓	✓	✓	×	×
team	remove_member	✓	✓	✓	×	×
billing	view	✓	✓	×	×	×
billing	manage	✓	✓	×	×	×
device	revoke	✓	✓	×	✓ (own)	×

6.3 Complete SQL Schema

```
-- =====
-- HelixTerminator RBAC Schema
-- =====

-- Roles table: system-defined + org-custom roles
CREATE TABLE roles (
```

```

    id            UUID            PRIMARY KEY DEFAULT
                gen_random_uuid(),
    name          VARCHAR(64) NOT NULL,
    display_name  TEXT            NOT NULL,
    description   TEXT,
    -- NULL org_id = system role; non-NULL = org-custom role
    org_id        UUID            REFERENCES organizations(id) ON
                DELETE CASCADE,
    is_system     BOOLEAN         NOT NULL DEFAULT FALSE,
    is_custom     BOOLEAN         NOT NULL DEFAULT FALSE,
    parent_role_id UUID            REFERENCES roles(id), -- role
                inheritance
    created_at    TIMESTAMPTZ     NOT NULL DEFAULT NOW(),
    updated_at    TIMESTAMPTZ     NOT NULL DEFAULT NOW(),
    UNIQUE (name, org_id)
);

-- Seed system roles
INSERT INTO roles (name, display_name, is_system) VALUES
    ('super_admin', 'Super Administrator', TRUE),
    ('org_admin',   'Organization Administrator', TRUE),
    ('team_admin',  'Team Administrator', TRUE),
    ('member',      'Member', TRUE),
    ('auditor',     'Auditor', TRUE),
    ('api_user',    'API User', TRUE);

-- Permissions: atomic action on a resource type
CREATE TABLE permissions (
    id            UUID            PRIMARY KEY DEFAULT
                gen_random_uuid(),
    resource_type VARCHAR(64) NOT NULL, -- e.g. 'vault',
                'host', 'session'
    action        VARCHAR(64) NOT NULL, -- e.g. 'create',
                'read', 'connect'
    description   TEXT,
    UNIQUE (resource_type, action)
);

-- Role-to-permission assignments
CREATE TABLE role_permissions (
    role_id        UUID            NOT NULL REFERENCES roles(id) ON
                DELETE CASCADE,
    permission_id  UUID            NOT NULL REFERENCES
                permissions(id) ON DELETE CASCADE,
    -- Scope: 'own' (user's own resources), 'team', 'org', 'all'
    scope          VARCHAR(16) NOT NULL DEFAULT 'own',
    PRIMARY KEY (role_id, permission_id)

```

```

);

-- User role assignments: assigns a role to a user, optionally
    scoped to
-- an org, team, or specific resource
CREATE TABLE role_assignments (
    id                UUID                PRIMARY KEY DEFAULT
        gen_random_uuid(),
    user_id           UUID                NOT NULL REFERENCES users(id) ON
        DELETE CASCADE,
    role_id           UUID                NOT NULL REFERENCES roles(id) ON
        DELETE CASCADE,
    org_id            UUID                REFERENCES organizations(id) ON
        DELETE CASCADE,
    team_id           UUID                REFERENCES teams(id) ON DELETE
        CASCADE,
    -- Optional resource-level scoping
    resource_type     VARCHAR(64),
    resource_id       UUID,
    -- Grant metadata
    granted_by        UUID                REFERENCES users(id),
    granted_at        TIMESTAMPTZ NOT NULL DEFAULT NOW(),
    expires_at        TIMESTAMPTZ,
    -- Conditions (JSON policy expression)
    conditions        JSONB,
    CONSTRAINT role_assignments_org_or_team
        CHECK (org_id IS NOT NULL OR team_id IS NOT NULL)
);

CREATE INDEX role_assignments_user_idx ON role_assignments
    (user_id);
CREATE INDEX role_assignments_org_idx  ON role_assignments
    (org_id);
CREATE INDEX role_assignments_team_idx ON role_assignments
    (team_id);

-- Resource policies: fine-grained per-resource access control
CREATE TABLE resource_policies (
    id                UUID                PRIMARY KEY DEFAULT
        gen_random_uuid(),
    resource_type     VARCHAR(64) NOT NULL,
    resource_id       UUID                NOT NULL,
    -- Subject: user, team, or role
    subject_type      VARCHAR(16) NOT NULL, -- 'user', 'team',
        'role'
    subject_id        UUID                NOT NULL,
    -- Effect
    effect            VARCHAR(8) NOT NULL DEFAULT 'allow', --
        'allow' | 'deny'

```

```

-- Actions allowed/denied (NULL = all)
actions          TEXT[],
-- Attribute-based conditions
conditions        JSONB,
-- Granting metadata
granted_by        UUID          REFERENCES users(id),
granted_at        TIMESTAMPTZ NOT NULL DEFAULT NOW(),
expires_at        TIMESTAMPTZ,
UNIQUE (resource_type, resource_id, subject_type, subject_id)
);

CREATE INDEX resource_policies_resource_idx
ON resource_policies (resource_type, resource_id);
CREATE INDEX resource_policies_subject_idx
ON resource_policies (subject_type, subject_id);

-- Attribute-based access control conditions (ABAC extension)
-- Stored in resource_policies.conditions and
--   role_assignments.conditions
-- Example conditions JSONB:
-- {
--   "ip_allowlist": ["10.0.0.0/8", "192.168.1.0/24"],
--   "time_windows": [{"days": ["mon", "tue", "wed", "thu", "fri"],
--     "start": "09:00", "end": "17:00", "tz": "UTC"}],
--   "require_mfa": true,
--   "require_approval": {"approvers": ["user_id_1", "user_id_2"],
--     "min_approvals": 1},
--   "require_recording": true,
--   "max_session_duration_seconds": 3600
-- }

-- Approval requests: for privileged access requiring explicit
--   approval
CREATE TABLE access_approvals (
  id          UUID          PRIMARY KEY DEFAULT
    gen_random_uuid(),
  requester_id  UUID          NOT NULL REFERENCES users(id),
  resource_type VARCHAR(64) NOT NULL,
  resource_id   UUID          NOT NULL,
  action        VARCHAR(64) NOT NULL,
  reason        TEXT          NOT NULL,
  status        VARCHAR(16) NOT NULL DEFAULT 'pending', --
    pending|approved|denied|expired
  reviewer_id   UUID          REFERENCES users(id),
  reviewed_at   TIMESTAMPTZ,
  reviewer_note TEXT,

```

```

    expires_at      TIMESTAMPTZ NOT NULL,
    created_at      TIMESTAMPTZ NOT NULL DEFAULT NOW(),

    -- Auto-revoke: if set, access expires this many seconds
    -- after approval
    access_ttl_seconds INT
);

CREATE INDEX access_approvals_requester_idx ON access_approvals
    (requester_id);
CREATE INDEX access_approvals_status_idx ON access_approvals
    (status) WHERE status = 'pending';

-- Session recording requirements
CREATE TABLE session_recording_policies (
    id                UUID          PRIMARY KEY DEFAULT
        gen_random_uuid(),
    org_id            UUID          NOT NULL REFERENCES
        organizations(id),
    -- Trigger conditions
    applies_to_roles TEXT[]        DEFAULT '{}', -- Role names that
        trigger recording
    applies_to_hosts JSONB,        -- Host tags that
        trigger recording
    -- Recording settings
    record_keystrokes BOOLEAN NOT NULL DEFAULT TRUE,
    record_output     BOOLEAN NOT NULL DEFAULT TRUE,
    retention_days    INT      NOT NULL DEFAULT 90,
    created_at        TIMESTAMPTZ NOT NULL DEFAULT NOW()
);

-- Effective permissions view (for API use)
CREATE OR REPLACE VIEW effective_permissions AS
SELECT
    ra.user_id,
    ra.org_id,
    ra.team_id,
    p.resource_type,
    p.action,
    rp.scope,
    ra.conditions,
    ra.expires_at
FROM role_assignments ra
JOIN role_permissions rp ON rp.role_id = ra.role_id
JOIN permissions p ON p.id = rp.permission_id
WHERE (ra.expires_at IS NULL OR ra.expires_at > NOW());

```

6.4 Policy Evaluation Engine

```
// pkg/rbac/engine.go
package rbac

// Evaluate determines whether a subject can perform an action on
// a resource.
// It evaluates role permissions, resource policies, and ABAC
// conditions.
func (e *Engine) Evaluate(ctx context.Context, req
    EvaluationRequest) (bool, string, error) {
    // 1. Check if there's an explicit DENY policy (deny overrides
    // allow)
    denied, reason, err := e.checkDeny(ctx, req)
    if err != nil {
        return false, "", err
    }
    if denied {
        return false, reason, nil
    }

    // 2. Check role-based permissions
    allowed, err := e.checkRolePermissions(ctx, req)
    if err != nil {
        return false, "", err
    }

    // 3. Check resource-specific policies
    if !allowed {
        allowed, err = e.checkResourcePolicies(ctx, req)
        if err != nil {
            return false, "", err
        }
    }

    if !allowed {
        return false, "no matching allow policy", nil
    }

    // 4. Evaluate ABAC conditions (IP, time, MFA requirements)
    conditionsMet, reason, err := e.evaluateConditions(ctx, req)
    if err != nil {
        return false, "", err
    }
    if !conditionsMet {
        return false, reason, nil
    }
}
```

```

    }

    return true, "", nil
}

// evaluateConditions checks ABAC conditions for the request
// context.
func (e *Engine) evaluateConditions(ctx context.Context, req
    EvaluationRequest) (bool, string, error) {
    conditions := req.PolicyConditions

    // IP allowlist check
    if len(conditions.IPAllowlist) > 0 {
        if !ipInAllowlist(req.Context.IPAddress,
            conditions.IPAllowlist) {
            return false, fmt.Sprintf("IP %s not in allowlist",
                req.Context.IPAddress), nil
        }
    }

    // Time window check
    if len(conditions.TimeWindows) > 0 {
        inWindow := false
        for _, window := range conditions.TimeWindows {
            if timeInWindow(req.Context.Time, window) {
                inWindow = true
                break
            }
        }
        if !inWindow {
            return false, "access not permitted at this time", nil
        }
    }

    // MFA requirement
    if conditions.RequireMFA && !req.Context.MFAVerified {
        return false, "MFA required for this resource", nil
    }

    // Approval requirement
    if conditions.RequireApproval != nil {
        approved, err := e.checkApproval(ctx, req)
        if err != nil {
            return false, "", err
        }
        if !approved {

```

```

        return false, "approval required", nil
    }
}

return true, "", nil
}

```

7. Audit Logging

7.1 Architecture

The audit log is the authoritative record of all security-relevant events in HelixTerminator. Its design requirements are:

- **Immutable:** events cannot be modified or deleted after writing
- **Append-only:** the log only grows; no deletions
- **Tamper-evident:** hash chain ensures that any tampering is detectable
- **Highly available:** a Kafka-backed pipeline ensures no events are lost even under high load
- **Searchable:** indexed in PostgreSQL for querying, with optional SIEM export

Pipeline:

```

Service → Audit SDK → Kafka (audit-events topic) → Audit Consumer
→ PostgreSQL
                                                    → SIEM (Splunk/
Datadog)
                                                    → S3 (long-term
archive)

```

7.2 Complete PostgreSQL Audit Schema

```

-- =====
-- HelixTerminator Audit Log Schema
-- =====

CREATE TABLE audit_events (
    -- Identity
    id          UUID          PRIMARY KEY DEFAULT
        gen_random_uuid(),

```

```

event_id          VARCHAR(32) NOT NULL UNIQUE, -- e.g.,
                  "evt_01HQZ7..."

-- Temporal
occurred_at       TIMESTAMPTZ NOT NULL,
received_at       TIMESTAMPTZ NOT NULL DEFAULT NOW(),

-- Event classification
event_type        VARCHAR(128) NOT NULL, -- e.g.,
                  "auth.login.success"
event_category    VARCHAR(64) NOT NULL, -- "auth", "vault",
                  "session", "key", "admin"
severity          VARCHAR(16) NOT NULL DEFAULT 'info', --
                  debug|info|warning|error|critical

-- Actor (who performed the action)
actor_type        VARCHAR(32) NOT NULL, -- "user", "service",
                  "system"
actor_id          UUID,                  -- user_id or service
                  account ID
actor_email       TEXT,
actor_name        TEXT,
actor_ip          INET,
actor_country     CHAR(2),              -- ISO 3166-1 alpha-2
actor_user_agent  TEXT,
actor_device_id   UUID,

-- Session context
session_id        VARCHAR(32),
access_token_jti  VARCHAR(32),

-- Organization context
org_id            UUID REFERENCES organizations(id),
team_id           UUID REFERENCES teams(id),

-- Target (what was acted upon)
target_type       VARCHAR(64), -- "vault", "host", "ssh_key",
                  "user", etc.
target_id         UUID,
target_name       TEXT,

-- Outcome
outcome           VARCHAR(16) NOT NULL DEFAULT 'success', --
                  "success" | "failure" | "partial"
error_code        VARCHAR(64),
error_message     TEXT,

```

```

-- Metadata (event-type-specific details)
metadata          JSONB          NOT NULL DEFAULT '{}',

-- Tamper detection: SHA-256 hash of (prev_hash || event data)
prev_hash         CHAR(64)       NOT NULL,    -- hash of the
previous event
event_hash        CHAR(64)       NOT NULL,
-- hash of (prev_hash || this event's data)

-- Retention
retention_class   VARCHAR(32)    NOT NULL DEFAULT 'standard', --
standard|compliance|critical
purge_after       TIMESTAMPTZ,
-- NULL = never purge (compliance events)

-- Kafka provenance
kafka_topic       VARCHAR(128),
kafka_partition   INT,
kafka_offset      BIGINT
);

-- Partitioned by month for efficient archival and query
performance
CREATE TABLE audit_events_2026_06 PARTITION OF audit_events
FOR VALUES FROM ('2026-06-01') TO ('2026-07-01');

-- Indexes for common query patterns
CREATE INDEX audit_events_actor_idx ON audit_events
(actor_id, occurred_at DESC);
CREATE INDEX audit_events_org_idx ON audit_events (org_id,
occurred_at DESC);
CREATE INDEX audit_events_session_idx ON audit_events
(session_id) WHERE session_id IS NOT NULL;
CREATE INDEX audit_events_target_idx ON audit_events
(target_type, target_id, occurred_at DESC);
CREATE INDEX audit_events_event_type_idx ON audit_events
(event_type, occurred_at DESC);
CREATE INDEX audit_events_occurred_at_idx ON audit_events
(occurred_at DESC);
CREATE INDEX audit_events_severity_idx ON audit_events
(severity) WHERE severity IN ('error', 'critical');

-- Full-text search on metadata
CREATE INDEX audit_events_metadata_gin ON audit_events USING GIN
(metadata);

-- The hash chain anchor (genesis record)
INSERT INTO audit_events (

```

```

    id, event_id, occurred_at, event_type, event_category,
    severity,
    actor_type, actor_id, org_id, outcome, metadata, prev_hash,
    event_hash
) VALUES (
    gen_random_uuid(), 'evt_genesis', NOW(),
    'system.log.initialized', 'system', 'info',
    'system', NULL, NULL, 'success', '{"genesis": true}',
    '0000000000000000000000000000000000000000000000000000000000000000',
    encode(sha256('genesis'), 'hex')
);

```

```

-- Immutability enforcement: prevent UPDATE and DELETE
CREATE RULE no_audit_update AS ON UPDATE TO audit_events DO
    INSTEAD NOTHING;
CREATE RULE no_audit_delete AS ON DELETE TO audit_events DO
    INSTEAD NOTHING;

```

```

-- Alternatively, use Row Level Security to enforce append-only
ALTER TABLE audit_events ENABLE ROW LEVEL SECURITY;
CREATE POLICY audit_insert_only ON audit_events FOR INSERT TO
    audit_writer WITH CHECK (TRUE);
CREATE POLICY audit_select ON audit_events FOR SELECT TO
    audit_reader USING (TRUE);
-- No UPDATE or DELETE policies = those operations are blocked

```

```

-- Event type reference table
CREATE TABLE audit_event_types (
    event_type      VARCHAR(128) PRIMARY KEY,
    description     TEXT NOT NULL,
    severity        VARCHAR(16) NOT NULL DEFAULT 'info',
    retention_class VARCHAR(32) NOT NULL DEFAULT 'standard',
    retention_years INT NOT NULL DEFAULT 1
);

```

```

-- Seed event types
INSERT INTO audit_event_types (event_type, description, severity,
    retention_class, retention_years) VALUES
    -- Authentication events
    ('auth.login.success', 'Successful user login',
    'info', 'standard', 1),
    ('auth.login.failure', 'Failed login attempt',
    'warning', 'standard', 1),
    ('auth.login.blocked', 'Login blocked (rate limit/
    lockout)', 'warning', 'standard', 1),
    ('auth.logout', 'User logout', 'info',
    'standard', 1),
    ('auth.mfa.success', 'MFA verification
    successful', 'info', 'standard', 1),

```

```

('auth.mfa.failure', 'MFA verification failed',
 'warning', 'standard', 1),
('auth.mfa.enrolled', 'New MFA factor enrolled',
 'info', 'compliance', 7),
('auth.mfa.revoked', 'MFA factor removed',
 'warning', 'compliance', 7),
('auth.token.issued', 'Access token issued',
 'debug', 'standard', 1),
('auth.token.refreshed', 'Token refreshed', 'debug',
 'standard', 1),
('auth.token.revoked', 'Token explicitly revoked',
 'info', 'standard', 1),
('auth.password.changed', 'Password changed', 'info',
 'compliance', 7),
('auth.password.reset', 'Password reset performed',
 'info', 'compliance', 7),
('auth.device.enrolled', 'New device enrolled',
 'info', 'compliance', 7),
('auth.device.revoked', 'Device access revoked',
 'warning', 'compliance', 7),
-- Vault events
('vault.created', 'Vault created', 'info',
 'standard', 1),
('vault.deleted', 'Vault deleted', 'warning',
 'compliance', 7),
('vault.item.accessed', 'Vault item decrypted and
accessed', 'info', 'standard', 1),
('vault.item.created', 'New vault item created',
 'info', 'standard', 1),
('vault.item.updated', 'Vault item modified',
 'info', 'standard', 1),
('vault.item.deleted', 'Vault item deleted',
 'warning', 'compliance', 7),
('vault.shared', 'Vault shared with user/
team', 'info', 'compliance', 7),
('vault.access.revoked', 'Vault access revoked',
 'info', 'compliance', 7),
('vault.key.rotated', 'Vault encryption key
rotated', 'info', 'compliance', 7),
-- SSH session events
('session.started', 'SSH session started',
 'info', 'standard', 1),
('session.terminated', 'SSH session ended', 'info',
 'standard', 1),
('session.terminated.policy', 'Session terminated by policy
enforcement', 'warning', 'compliance', 7),
('session.command.executed', 'Command executed in SSH
session', 'debug', 'standard', 1),
('session.file.transferred', 'File transferred over SFTP/
SCP', 'info', 'standard', 1),
('session.recording.started', 'Session recording started',
 'info', 'standard', 1),
('session.recording.stopped', 'Session recording stopped',
 'info', 'standard', 1),

```

```

-- SSH key events
('sshkey.generated',          'SSH key pair generated',
 'info', 'compliance', 7),
('sshkey.imported',          'SSH public key imported',
 'info', 'compliance', 7),
('sshkey.exported',          'SSH private key exported',
 'critical', 'compliance', 7),
('sshkey.used',              'SSH key used for
 authentication', 'debug', 'standard', 1),
('sshkey.deleted',           'SSH key deleted', 'info',
 'compliance', 7),
-- Certificate events
('cert.issued',              'SSH certificate issued by PKI CA', 'info', 'compliance',
 7),
('cert.revoked',              'SSH certificate revoked',
 'warning', 'compliance', 7),
-- Admin events
('admin.user.created',        'User account created',
 'info', 'compliance', 7),
('admin.user.deleted',        'User account deleted',
 'warning', 'compliance', 7),
('admin.user.role.changed',    'User role assignment
 changed', 'warning', 'compliance', 7),
('admin.org.settings.changed', 'Organization settings
 modified', 'warning', 'compliance', 7),
('admin.sso.configured',      'SSO integration configured',
 'warning', 'compliance', 7),
('admin.scim.provisioned',     'User provisioned via SCIM',
 'info', 'compliance', 7),
('admin.scim.deprovisioned',   'User deprovisioned via
 SCIM', 'warning', 'compliance', 7);

```

7.3 Hash Chain Implementation

```

// pkg/audit/hashchain.go
package audit

import (
    "crypto/sha256"
    "encoding/hex"
    "encoding/json"
    "fmt"
    "time"
)

// EventRecord is the canonical audit event representation used
// for hashing.
// Fields are lexicographically ordered JSON keys for
// deterministic serialization.

```

```

type EventRecord struct {
    EventID      string          `json:"event_id"`
    EventType    string          `json:"event_type"`
    OccurredAt   time.Time       `json:"occurred_at"`
    ActorID      string          `json:"actor_id"`
    TargetType   string          `json:"target_type,omitempty"`
    TargetID     string          `json:"target_id,omitempty"`
    OrgID        string          `json:"org_id"`
    Outcome      string          `json:"outcome"`
    Metadata     json.RawMessage `json:"metadata"`
    PrevHash     string          `json:"prev_hash"`
}

// ComputeHash computes the tamper-detection hash for an audit
// event.
// The hash is: SHA-256(prev_hash || canonical_json(event))
func ComputeHash(event *EventRecord) (string, error) {
    // Produce a canonical (deterministic) JSON representation
    data, err := json.Marshal(event)
    if err != nil {
        return "", fmt.Errorf("failed to marshal event for
        hashing: %w", err)
    }

    h := sha256.New()
    h.Write([]byte(event.PrevHash))
    h.Write(data)
    return hex.EncodeToString(h.Sum(nil)), nil
}

// VerifyChain verifies the integrity of the audit log hash chain
// between
// startSerial and endSerial. Returns the first invalid entry if
// found.
func VerifyChain(ctx context.Context, db *sql.DB, startID, endID
    string) (*ChainVerificationResult, error) {
    rows, err := db.QueryContext(ctx, `
        SELECT event_id, prev_hash, event_hash, occurred_at,
               event_type, actor_id, target_type, target_id,
               org_id, outcome, metadata
        FROM audit_events
        WHERE occurred_at BETWEEN
            (SELECT occurred_at FROM audit_events WHERE event_id =
            $1) AND
            (SELECT occurred_at FROM audit_events WHERE event_id =
            $2)
        ORDER BY occurred_at ASC, id ASC`,

```

```

        startID, endID,
    )
    if err != nil {
        return nil, err
    }
    defer rows.Close()

    var prevHash string
    var count int
    for rows.Next() {
        var rec EventRecord
        var storedHash string
        if err := rows.Scan(
            &rec.EventID, &rec.PrevHash, &storedHash,
            &rec.OccurredAt,
            &rec.EventType, &rec.ActorID, &rec.TargetType,
            &rec.TargetID,
            &rec.OrgID, &rec.Outcome, &rec.Metadata,
        ); err != nil {
            return nil, err
        }

        // Verify prev_hash linkage
        if prevHash != "" && rec.PrevHash != prevHash {
            return &ChainVerificationResult{
                Valid:          false,
                BreakAtEventID:  rec.EventID,
                ExpectedPrevHash: prevHash,
                ActualPrevHash:   rec.PrevHash,
            }, nil
        }

        // Verify event hash
        computed, err := ComputeHash(&rec)
        if err != nil {
            return nil, err
        }
        if computed != storedHash {
            return &ChainVerificationResult{
                Valid:          false,
                BreakAtEventID:  rec.EventID,
                HashMismatch:     true,
            }, nil
        }
    }

```

```

        prevHash = storedHash
        count++
    }

    return &ChainVerificationResult{Valid: true, EventsVerified:
        count}, nil
}

```

7.4 SIEM Integration

```

// pkg/audit/siem/splunk.go
package siem

// SplunkForwarder forwards audit events to Splunk HEC (HTTP Event
// Collector).
type SplunkForwarder struct {
    hecURL    string
    token     string
    client    *http.Client
    batchSize int
    buffer    chan *audit.EventRecord
}

// SplunkHECEvent is the Splunk HEC event format.
type SplunkHECEvent struct {
    Time        float64    `json:"time"`
    Host        string     `json:"host"`
    Source      string     `json:"source"`
    Sourcetype  string     `json:"sourcetype"`
    Index       string     `json:"index"`
    Event       interface{} `json:"event"`
}

func (s *SplunkForwarder) ForwardEvent(event *audit.EventRecord) {
    s.buffer <- event
}

func (s *SplunkForwarder) batchWorker(ctx context.Context) {
    ticker := time.NewTicker(5 * time.Second)
    batch := make([]*audit.EventRecord, 0, s.batchSize)

    for {
        select {
        case event := <-s.buffer:
            batch = append(batch, event)

```

```

        if len(batch) >= s.batchSize {
            s.flushBatch(ctx, batch)
            batch = batch[:0]
        }
    case <-ticker.C:
        if len(batch) > 0 {
            s.flushBatch(ctx, batch)
            batch = batch[:0]
        }
    case <-ctx.Done():
        return
    }
}
}

```

7.5 Retention and Archival

Classification	Examples	Retention	Storage
debug	Token refreshes, key uses	90 days	Hot (PostgreSQL)
standard	Logins, sessions, vault access	1 year	Hot → Cold (S3)
compliance	Role changes, MFA changes, key exports	7 years	Hot (1y) + Archive (S3 Glacier)
critical	Key ceremonies, CA operations, breach events	10 years	Permanent archive

8. Network Security

8.1 TLS Configuration

External endpoints (API, web client, SCIM): - TLS 1.3 only — TLS 1.2 and below are rejected - No cipher suite negotiation for TLS 1.3 (the protocol mandates AEAD ciphers) - OCSP stapling enabled - HSTS: Strict-Transport-Security: max-age=63072000; includeSubDomains; preload

Internal service mesh: - mTLS via Istio/SPIRE (see Section 1.4) - TLS 1.3 between all service pairs - Certificate TTL: 24 hours

Nginx TLS configuration for external endpoints:

```
server {
    listen 443 ssl;
    server_name api.helixterminator.io;

    ssl_certificate      /etc/nginx/ssl/api.helixterminator.io.crt;
    ssl_certificate_key  /etc/nginx/ssl/api.helixterminator.io.key;

    # TLS 1.3 only
    ssl_protocols TLSv1.3;

    # TLS 1.3 cipher suites (informational – TLS 1.3 doesn't allow
    explicit config)
    # TLS_AES_256_GCM_SHA384
    # TLS_CHACHA20_POLY1305_SHA256
    # TLS_AES_128_GCM_SHA256 (disable for higher security)

    # HSTS – 2 years, include subdomains, preload
    add_header Strict-Transport-Security "max-age=63072000;
    includeSubDomains; preload" always;

    # OCSP Stapling
    ssl_stapling on;
    ssl_stapling_verify on;
    ssl_trusted_certificate /etc/nginx/ssl/chain.pem;

    # Session resumption via session tickets (disabled – impairs
    forward secrecy in some configs)
    ssl_session_tickets off;
    ssl_session_cache shared:SSL:10m;
```

```

    ssl_session_timeout 1d;

    # DH parameters (not used with TLS 1.3, but included for
    safety)
    ssl_dhparam /etc/nginx/ssl/dhparam4096.pem;
}

```

8.2 SSH Protocol Hardening

SSH connections managed by HelixTerminator to target servers enforce strict protocol parameters:

```

// pkg/ssh/client_config.go
package ssh

import (
    "golang.org/x/crypto/ssh"
)

// SecureClientConfig returns a hardened ssh.ClientConfig that
// only accepts
// strong algorithms and enforces host certificate verification.
func SecureClientConfig(signer ssh.Signer, hostKeyCallback
    ssh.HostKeyCallback) *ssh.ClientConfig {
    return &ssh.ClientConfig{
        Auth: []ssh.AuthMethod{
            ssh.PublicKeys(signer),
        },
        HostKeyCallback: hostKeyCallback,
        HostKeyAlgorithms: []string{
            // Ed25519 host keys preferred
            ssh.KeyAlgoED25519,
            // ECDSA P-256 acceptable
            ssh.KeyAlgoECDSA256,
            // RSA with SHA-256 or SHA-512 signature (not SHA-1)
            ssh.KeyAlgoRSASHA256,
            ssh.KeyAlgoRSASHA512,
            // SSH certificates
            ssh.CertAlgoED25519v01,
            ssh.CertAlgoECDSA256v01,
            ssh.CertAlgoRSASHA256v01,
            ssh.CertAlgoRSASHA512v01,
        },
        // Cipher order: ChaCha20-Poly1305 preferred (more cache-
        attack resistant

```

```

        // than AES-GCM on systems without AES-NI hardware
        acceleration)
        // AES-256-GCM next (hardware-accelerated on most modern
        CPUs)
        // AES-128-GCM acceptable
        // NOTE: golang.org/x/crypto/ssh picks its own order; use
        a patched version
        // or negotiate post-connection to enforce this.
    }
}

```

SSH server sshd_config hardening (applied to HelixTerminator-managed bastion hosts):

```

# Protocol
Protocol 2

# Ciphers – ChaCha20 and AES-256-GCM only
Ciphers chacha20-poly1305@openssh.com,aes256-gcm@openssh.com

# MACs – Encrypt-then-MAC, SHA-2 family only
MACs hmac-sha2-512-etm@openssh.com,hmac-sha2-256-etm@openssh.com

# Key Exchange – curve25519 and ECDH with NIST P-256
KexAlgorithms curve25519-sha256,curve25519-sha256@libssh.org,ecdh-
sha2-nistp256

# Host keys – Ed25519 preferred, ECDSA P-256 fallback
HostKey /etc/ssh/ssh_host_ed25519_key
HostKey /etc/ssh/ssh_host_ecdsa_key

# Authentication
PasswordAuthentication no
PubkeyAuthentication yes
AuthorizedKeysFile /dev/null      # Don't use static
authorized_keys
TrustedUserCAKeys /etc/ssh/helixterm_user_ca.pub  # Trust
HelixTerminator CA
RevokedKeys /etc/ssh/helixterm_krl.bin

# Sessions
MaxAuthTries 3
MaxSessions 10
LoginGraceTime 30

# Logging

```

```
LogLevel VERBOSE
SyslogFacility AUTH
```

8.3 CORS Policy

```
// pkg/middleware/cors.go
package middleware

func CORSMiddleware() func(http.Handler) http.Handler {
    return cors.Handler(cors.Options{
        AllowedOrigins: []string{
            "https://app.helixterminator.io",
            "https://admin.helixterminator.io",
        },
        AllowedMethods: []string{
            "GET", "POST", "PUT", "PATCH", "DELETE", "OPTIONS",
        },
        AllowedHeaders: []string{
            "Authorization",
            "Content-Type",
            "X-Request-ID",
            "X-Device-ID",
        },
        ExposedHeaders: []string{
            "X-Request-ID",
            "X-RateLimit-Limit",
            "X-RateLimit-Remaining",
            "X-RateLimit-Reset",
        },
        AllowCredentials: true,
        MaxAge:           86400, // 24 hours preflight cache
    })
}
```

8.4 Content Security Policy

The HelixTerminator web client sends the following CSP headers:

```
Content-Security-Policy:
    default-src 'self';
    script-src 'self' 'nonce-{RANDOM_NONCE}';
    style-src 'self' 'nonce-{RANDOM_NONCE}';
    img-src 'self' data: https://secure.gravatar.com;
    font-src 'self';
```

```
connect-src 'self' https://api.helixterminator.io https://
auth.helixterminator.io wss://terminal.helixterminator.io;
frame-src 'none';
object-src 'none';
base-uri 'self';
form-action 'self';
upgrade-insecure-requests;
block-all-mixed-content;
require-trusted-types-for 'script';
```

Additional security headers:

```
X-Content-Type-Options: nosniff
X-Frame-Options: DENY
X-XSS-Protection: 1; mode=block
Referrer-Policy: strict-origin-when-cross-origin
Permissions-Policy: camera=(), microphone=(), geolocation=()
Cross-Origin-Opener-Policy: same-origin
Cross-Origin-Embedder-Policy: require-corp
Cross-Origin-Resource-Policy: same-site
```

8.5 DDoS Protection

Layer 1 — Cloudflare: - DDoS protection at edge (L3/L4 and L7) - Rate limiting at CDN level - Bot management (challenge suspicious traffic) - IP reputation blocking

Layer 2 — Application-level rate limiting:

```
// pkg/middleware/ratelimit.go
package middleware

// Rate limits by category (requests per window)
var rateLimits = map[string]RateLimit{
    "auth.login": {Requests: 5, Window: 15 *
        time.Minute, Key: "ip"},
    "auth.mfa": {Requests: 10, Window: 15 *
        time.Minute, Key: "ip"},
    "auth.password_reset": {Requests: 3, Window: 1 *
        time.Hour, Key: "email"},
    "api.general": {Requests: 1000, Window: 1 *
        time.Minute, Key: "user"},
    "api.vault.read": {Requests: 500, Window: 1 *
        time.Minute, Key: "user"},
    "cert.sign": {Requests: 60, Window: 1 *
        time.Minute, Key: "user"},
```

```

        "websocket.terminal": {Requests: 10, Window: 1 *
            time.Minute, Key: "user"},
    }

```

8.6 Certificate Pinning (Mobile Apps)

Mobile applications (iOS and Android) pin to the HelixTerminator TLS certificate chain:

```

// iOS: HelixTerminator certificate pinning
class PinnedURLSession: NSObject, URLSessionDelegate {
    // SHA-256 hashes of the Subject Public Key Info (SPKI) of
    // pinned certificates
    private let pinnedHashes: Set<String> = [
        "sha256/AAAA...=", // api.helixterminator.io leaf cert
        "sha256BBBB...=", // intermediate CA
        "sha256CCCC...=", // root CA (backup pin)
    ]

    func urlSession(_ session: URLSession,
                    didReceive challenge:
                        URLAuthenticationChallenge,
                    completionHandler: @escaping
                        (URLSession.AuthChallengeDisposition, URLCredential?) ->
                        Void) {
        guard challenge.protectionSpace.authenticationMethod ==
            NSURLAuthenticationMethodServerTrust,
            let serverTrust =
                challenge.protectionSpace.serverTrust else {
            completionHandler(.cancelAuthenticationChallenge, nil)
            return
        }

        guard verifyServerTrust(serverTrust, against:
            pinnedHashes) else {
            completionHandler(.cancelAuthenticationChallenge, nil)
            return
        }

        completionHandler(.useCredential, URLCredential(trust:
            serverTrust))
    }
}

```

8.7 Subresource Integrity

All third-party assets loaded by the web client use SRI hashes:

```
<script
  src="https://cdn.example.com/library.js"
  integrity="sha384-oqVuAfXRKap7fdgcCY5uykM6+R9GqQ8K/ux..."
  crossorigin="anonymous">
</script>
```

9. Compliance

9.1 SOC 2 Type II Controls Mapping

HelixTerminator maps its controls to the AICPA Trust Services Criteria (SOC 2):

Control Domain	SOC 2 Criterion	HelixTerminator Implementation
CC6.1 — Logical access controls	Security	RBAC + ABAC (Section 6), mTLS (Section 1.4)
CC6.2 — Prior to issuance of system credentials	Security	User provisioning workflow, SCIM, JIT
CC6.3 — Access reviews	Security	Quarterly automated role review, SCIM deprovisioning
CC6.6 — Security threats via the internet	Security	TLS 1.3, WAF, DDoS protection (Section 8)
CC6.7 — Transmission of data	Security	TLS everywhere, vault E2E encryption (Section 4)
CC6.8 — Unauthorized software	Security	Container image signing (cosign), SBOM generation
CC7.1 — Detection of threats	Availability	SIEM integration, anomaly detection
	Availability	

Control Domain	SOC 2 Criterion	HelixTerminator Implementation
CC7.2 — Monitoring for indicators of compromise	Security	Audit log alerts (Section 7)
CC9.2 — Vendor risk management		Third-party software SCA scanning, SBOM
A1.2 — Environmental protections	Availability	Multi-AZ deployment, RTO < 1 hour
PI1.5 — Data retention and disposal	Processing Integrity	Configurable retention, GDPR erasure
C1.1 — Confidentiality of customer data	Confidentiality	Zero-knowledge vault (Section 4)

9.2 GDPR Compliance

Data minimization: - User profile data: name, email, organization — no tracking or behavioral analytics beyond what is necessary for the service - IP addresses in audit logs anonymized after 90 days (last octet zeroed) - Session recording requires explicit opt-in and is disclosed to users

Right to erasure (GDPR Article 17):

```
// pkg/gdpr/erasure.go
package gdpr

// EraseUserData implements GDPR Article 17 right to erasure.
// It soft-deletes the user and schedules hard deletion of
//   personal data
// after the legally required retention period.
func EraseUserData(ctx context.Context, userID string) error {
    // 1. Revoke all active sessions immediately
    if err := sessions.RevokeAll(ctx, userID); err != nil {
        return err
    }

    // 2. Revoke all tokens (refresh tokens, device tokens, API
    //   keys)
    if err := tokens.RevokeAllForUser(ctx, userID); err != nil {
```

```

        return err
    }

    // 3. Delete personal vault items (user explicitly requested
    //      erasure)
    if err := vaults.DeletePersonalVault(ctx, userID); err != nil
    {
        return err
    }

    // 4. Anonymize audit log entries (replace PII with anonymized
    //      placeholders)
    // Note: audit log events cannot be deleted (immutable), but
    //      PII fields
    //      are overwritten with anonymized values while preserving
    //      event integrity
    if err := audit.AnonymizeUser(ctx, userID); err != nil {
        return err
    }

    // 5. Mark user as erased (for compliance log)
    return users.MarkErased(ctx, userID)
}

```

Data portability (GDPR Article 20):

Users can export their data at any time via GET /v1/me/export. The export includes: - Profile information - Vault items (plaintext, after authentication) - SSH key fingerprints (not private keys) - Audit log of their own activities - Connected devices list

Data processing agreements: Available for enterprise customers at legal.helixterminator.io/dpa.

Data residency: Customer data can be pinned to EU (Frankfurt + Dublin), US (us-east-1 + us-west-2), or APAC (ap-southeast-1 + ap-northeast-1). Data residency is enforced at the infrastructure level via AWS region selection and Kafka partition key routing.

9.3 HIPAA Compliance

For healthcare customers, HelixTerminator offers a HIPAA-eligible deployment configuration:

- **Business Associate Agreement (BAA):** Available upon request
- **Encryption at rest:** All PostgreSQL data encrypted with pgcrypto + volume encryption (AES-256)

- **Encryption in transit:** TLS 1.3 required (see Section 8.1)
- **Access controls:** RBAC with time-based and IP-based restrictions
- **Audit logging:** All access to PHI-containing systems logged (Section 7)
- **Minimum necessary:** ABAC conditions enforce need-to-know access
- **Workforce training:** Required annually for all employees with system access
- **Incident response:** 60-day breach notification timeline to HHS (Section 10)

Note: HelixTerminator stores SSH credentials, not PHI directly. However, because HelixTerminator provides access to systems that may process PHI, the platform is classified as a HIPAA business associate when used by covered entities.

9.4 PCI DSS (Payment Card Industry)

HelixTerminator's payment processing is handled by Stripe (a PCI DSS Level 1 certified processor). HelixTerminator does not store card numbers. Relevant PCI DSS controls:

Requirement	Implementation
Req 2 — Do not use vendor-supplied defaults	All default passwords changed; hardened OS configurations
Req 4 — Encrypt transmission	TLS 1.3 for all data in transit
Req 7 — Restrict access	RBAC (Section 6)
Req 8 — Identify and authenticate users	MFA required (Section 2)
Req 10 — Track and monitor access	Audit logging (Section 7)
Req 11 — Regularly test security	Quarterly penetration testing

9.5 ISO 27001 Alignment

HelixTerminator's security controls align with ISO/IEC 27001:2022 Annex A:

ISO 27001 Control	HelixTerminator Control
A.5.1 — Policies for information security	This document + internal security policy
A.8.2 — Privileged access rights	super_admin and org_admin role governance
A.8.5 — Secure authentication	MFA (Section 2), FIDO2, biometric
A.8.7 — Protection against malware	Container scanning, SBOM, runtime security (Falco)
A.8.12 — Data leakage prevention	Vault E2E encryption (Section 4), DLP policies
A.8.15 — Logging	Audit logging (Section 7)
A.8.24 — Use of cryptography	Key hierarchy, AES-256-GCM, Argon2id (Sections 3, 4)

9.6 Encryption at Rest

PostgreSQL:

```
-- Column-level encryption for sensitive fields using pgcrypto
-- Master key stored in AWS KMS, fetched at startup
UPDATE users SET totp_secret = pgp_sym_encrypt(
    totp_secret::text,
    current_setting('app.encryption_key')
);
```

Volume encryption: - All AWS EBS volumes use AES-256 encryption with CMKs in AWS KMS - All AWS S3 buckets use SSE-KMS with customer-managed keys - PostgreSQL RDS instances use AES-256 encryption

10. Incident Response

10.1 Incident Classification

Priority	Name	Description	Initial Response SLA
P0	Critical	Active breach, data exfiltration in	15 minutes

Priority	Name	Description	Initial Response SLA
P1	High	progress, CA compromise Compromised credentials in active use, unauthorized access confirmed	1 hour
P2	Medium	Suspected account compromise, unusual access patterns detected	4 hours
P3	Low	Failed attack attempt, vulnerability report (no active exploitation)	24 hours
P4	Informational	Security finding with no immediate impact	72 hours

10.2 P0 — Critical Incident Response Procedure

0–15 minutes (Detection and Initial Containment):

- 1. Alert:** PagerDuty fires; on-call security engineer and CTO paged
- 2. War Room:** Incident channel #incident-p0-YYYYMMDD-NN created in Slack
- 3. Preserve Evidence:** AWS CloudTrail logs, application logs, Kafka offsets frozen
- 4. Initial Scoping:** Determine affected systems, users, data types
- 5. Immediate Containment:**
 - If credential compromise: `POST /admin/v1/users/{id}/emergency-lockout`
 - If service compromise: `kubectl rollout undo deployment/{service}`
 - If CA compromise: activate CA revocation procedure (10.4)

15–60 minutes (Containment and Eradication):

1. **Revoke Affected Credentials:** All tokens, API keys, device certificates for affected users/services
2. **Rotate Secrets:** Rotate affected service credentials, database passwords
3. **Block Attack Vector:** Firewall rules, WAF rules, Cloudflare blocks
4. **Forensic Collection:** Memory dumps, disk images of affected nodes

1–24 hours (Recovery):

1. **Deploy Clean Instances:** Replace potentially compromised pods/nodes
2. **Re-validate Data Integrity:** Check audit log hash chain (Section 7.3)
3. **Controlled Service Restoration:** Bring services back online with enhanced monitoring
4. **Affected User Notification:** In-app + email notification to affected users

24–72 hours (Post-Incident):

1. **Root Cause Analysis:** Timeline reconstruction, attack vector identification
2. **Regulatory Notification:** If applicable (GDPR: 72 hours; HIPAA: 60 days; state breach laws vary)
3. **Remediation:** Fix the vulnerability, verify fix, deploy
4. **Post-Mortem:** Blameless post-mortem published to internal wiki within 5 business days

10.3 P1 — High Incident Response

Hour 0-1: Security engineer paged; initial scoping
Hour 1-2: Targeted containment (affected user/service only)
Hour 2-4: Evidence collection; vendor notification if third-party involved
Hour 4-8: Eradication and recovery
Hour 8-24: Enhanced monitoring; root cause analysis begins
Day 2-5: Post-mortem; remediation deployed

10.4 CA Compromise Emergency Procedure

A compromise of the SSH Certificate Authority is a P0-critical incident that invalidates all issued certificates. The procedure:

1. IMMEDIATE (within 15 minutes):
 - a. Take PKI service offline (`kubectl scale deployment/pki-ca --replicas=0`)
 - b. Activate emergency CA revocation: publish KRL with all active serials revoked
 - c. Push KRL to all known SSH servers (Ansible playbook: `emergency-krl-push.yml`)
 - d. Alert all customers via status page and email
2. SHORT-TERM (within 2 hours):
 - a. Conduct key ceremony with new Root CA (pre-generated and stored offline)
 - b. Issue new Intermediate CA signed by new Root CA
 - c. Distribute new CA public key to all SSH servers
 - d. Bring PKI service back online with new keys
3. USER RECOVERY (within 4 hours):
 - a. All users must re-authenticate (existing sessions invalidated)
 - b. New SSH certificates issued automatically upon next connection
 - c. Communicate via in-app notification and email
4. POST-INCIDENT (within 24 hours):
 - a. Full forensic analysis of how CA was compromised
 - b. Revoke and rotate any credentials that had access to CA signing keys
 - c. Audit all certificates issued in the 30 days prior to compromise

10.5 Account Takeover Response

```
// pkg/incident/ato.go
package incident

// RespondToAccountTakeover handles a suspected or confirmed
// account takeover.
// This is the automated response; human review follows.
func RespondToAccountTakeover(ctx context.Context, userID string,
    confidence ATOConfidence) error {
    switch confidence {
    case ATOConfidenceHigh:
        // Full lockout: revoke everything
        if err := revokeAllSessions(ctx, userID); err != nil {
            return err
        }
    }
}
```

```

    }
    if err := revokeAllTokens(ctx, userID); err != nil {
        return err
    }
    if err := lockAccount(ctx, userID,
        "suspected_account_takeover"); err != nil {
        return err
    }
    if err := notifyUser(ctx, userID,
        "security_alert_ato_high"); err != nil {
        return err
    }
    if err := notifyOrgAdmins(ctx, userID, "ato_alert"); err !=
    nil {
        return err
    }
    return notifySecurityTeam(ctx, userID, ATOIncidentP1)

case ATOConfidenceMedium:
    // Step-up auth required: force MFA on next access
    if err := requireStepUpAuth(ctx, userID); err != nil {
        return err
    }
    return notifyUser(ctx, userID,
        "security_alert_ato_medium")

case ATOConfidenceLow:
    // Flag for review; no immediate action
    return flagForSecurityReview(ctx, userID, "ato_signal")
}
return nil
}

```

10.6 Vulnerability Disclosure Policy

HelixTerminator operates a coordinated vulnerability disclosure program:

Disclosure endpoint: security@helixterminator.io / <https://security.helixterminator.io/report>

PGP Key: Published at <https://security.helixterminator.io/pgp-key.txt>

Response SLAs: - Initial acknowledgment: 24 hours - Triage and severity assessment: 72 hours - Remediation timeline provided: 7 days - Fix deployed (critical): 30 days - Fix deployed (high): 60 days - Fix deployed (medium/low): 90 days - Public disclosure (coordinated): 90 days from report (may be extended by mutual agreement)

Safe Harbor: HelixTerminator will not pursue legal action against security researchers who: 1. Report through the coordinated disclosure program 2. Do not access, exfiltrate, or modify customer data 3. Do not perform denial-of-service attacks 4. Do not disclose findings publicly before the remediation window expires

Bug Bounty: HelixTerminator operates a private bug bounty program through HackerOne. Payouts range from \$100 (informational) to \$20,000 (critical, RCE, CA compromise).

10.7 Breach Notification

Regulatory requirements by jurisdiction:

Regulation	Notification Target	Deadline
GDPR Art. 33	Supervisory authority (DPA)	72 hours
GDPR Art. 34	Affected individuals	“Without undue delay”
HIPAA Breach Rule	HHS	60 days (>500 individuals: 60 days + media)
CCPA	California AG (>500 CA residents)	“Expedient” (30 days practical)
NY SHIELD Act	NY AG	“Most expedient time possible”

Internal notification chain:

Security Incident → CTO → CEO → Legal Counsel → Board (P0 only)
→ Customer Success
(affected customers)
→ Marketing (public statement if needed)

Customer notification template: All breach notifications include: -
What happened (factual, no speculation) - What data was affected (specific categories) - What HelixTerminator has done in response -
What the customer should do (action items) - Contact information for questions - Reference to regulatory rights (GDPR, CCPA, etc.)

Appendix A: Cryptographic Algorithm Summary

Use Case	Algorithm	Key Size	Notes
Password hashing	Argon2id	32-byte output	t=3, m=64MB, p=4
Vault encryption	AES-256-GCM	256-bit	Per-item keys
Key derivation (vault)	HKDF-SHA256	256-bit output	From X25519 shared secret
Key exchange	X25519	255-bit	For device key wrapping
SSH key generation	Ed25519	255-bit	Preferred
SSH key generation	ECDSA P-384	384-bit	FIPS alternative
SSH key generation	RSA	4096-bit	Legacy systems only
JWT signing	EdDSA (Ed25519)	256-bit	Access tokens; RFC 8037. (RS256/ ES256/ HS256 used only in negative/ attack-test

Use Case	Algorithm	Key Size	Notes
			examples elsewhere in this document.)
TLS (external)	TLS 1.3	N/A	AEAD ciphers mandated
Workload identity	ECDSA P-256	256-bit	SPIFFE SVIDs
Audit hash chain	SHA-256	256-bit	Tamper detection
TOTP	HMAC-SHA1	160-bit HMAC	RFC 6238
API key entropy	CSPRNG	256-bit	Via crypto/rand

Appendix B: Security Dependencies

Package	Version	Purpose
golang.org/x/crypto	latest	Argon2id, SSH, ECDH, HKDF
github.com/spiffe/go-spiffe/v2	v2.x	SPIFFE workload identity
github.com/go-webauthn/webauthn	latest	FIDO2/WebAuthn
github.com/crewjam/saml	latest	SAML 2.0 SP
github.com/coreos/go-oidc/v3	v3.x	OIDC SP
github.com/hashicorp/vault/shamir	latest	Shamir's Secret Sharing
golang.org/x/crypto/hkdf	latest	HKDF key derivation
github.com/redis/go-redis/v9	v9.x	Refresh token store
	latest	

Package	Version	Purpose
github.com/twmb/ franz-go		Kafka client (audit log)

Appendix C: Security Contact Information

Role	Contact	PGP
Security Team	security@helixterminator.io	Published at security.helixterminator.io
Vulnerability Disclosure	security@helixterminator.io	Same
Data Protection Officer	dpo@helixterminator.io	Published at legal.helixterminator.io
Legal	legal@helixterminator.io	—
Status Page	status.helixterminator.io	—

*Document maintained by the HelixTerminator Platform Security Team.
For questions or proposed changes, open a PR against the platform-security repository or email security@helixterminator.io.*

Last updated: 2026-06-28 | Next review: 2026-09-28 | Document version: 1.0.0