

HelixTerminator – 09_performance_analysis

HelixTerminator: Performance Analysis Specification

Document: 09_performance_analysis.md

Version: 1.0.0

Classification: Internal – Performance Engineering

Last Updated: 2026-06-28

Owner: Platform Performance Team

Table of Contents

- [1. Performance Architecture Overview](#)
 - [2. Critical Path Analysis](#)
 - [3. Identified Bottlenecks & Danger Zones](#)
 - [4. Gap Analysis – What's Missing vs. Production-Ready](#)
 - [5. Database Performance Deep-Dive](#)
 - [6. Caching Strategy & Cache Architecture](#)
 - [7. SSH Proxy Performance Engineering](#)
 - [8. Kafka Performance Tuning](#)
 - [9. API Gateway & Rate Limiting Performance](#)
 - [10. Flutter Client Performance](#)
 - [11. Observability for Performance](#)
 - [12. Danger Zones – Security-Performance Tradeoffs](#)
 - [13. Load Testing Results & Benchmarks](#)
 - [14. Performance Improvement Roadmap](#)
-

1. Performance Architecture Overview

1.1 Performance Philosophy

HelixTerminator is engineered under a **performance-first, correctness-concurrent** philosophy. Every architectural decision is evaluated against three axes:

- Latency budget exhaustion** – Does this operation push us past our per-tier budget?

2. **Throughput ceiling** — Does this design path block horizontal scale-out?
3. **Tail latency amplification** — Does p99 grow faster than p50 as load increases?

The system targets **millions of requests per second** aggregate across its 25 microservices, with a non-negotiable **sub-millisecond p50 latency** for all hot-path operations. These targets are not aspirational—they are encoded as SLO burn-rate alerts and will page on-call engineers if violated.

Core tenets:

- **Measure before optimising.** Every identified bottleneck must be accompanied by a profiling artifact (pprof flamegraph, EXPLAIN ANALYZE, or Kafka consumer lag chart) before engineering time is allocated.
- **Design for degradation.** Every service must have a defined graceful degradation path that maintains $\geq 70\%$ user-facing functionality under 3× normal load.
- **Allocate zero work to the critical path.** Hot paths must avoid heap allocations. Use `sync.Pool`, pre-allocated byte slices, and stack-local buffers wherever possible.
- **Async by default.** Operations that do not require synchronous user feedback (audit logging, metrics emission, Kafka publishes for non-transactional events) must be off the critical path.

1.2 SLO/SLA Definitions per Microservice

All latency SLOs are measured at the **service egress point** (not client-perceived), excluding external network propagation. Availability SLAs are measured as rolling 30-day windows.

Service	p50 Latency	p95 Latency	p99 Latency	Throughput Target	Availability SLA
api-gateway	≤ 2 ms	≤ 8 ms	≤ 20 ms	500,000 req/s	99.99%
auth-service	≤ 5 ms	≤ 15 ms	≤ 40 ms	100,000 req/s	99.99%
ssh-proxy	≤ 1 ms (frame forward)	≤ 5 ms	≤ 15 ms	100,000 concurrent sessions	99.95%
vault-service	≤ 3 ms	≤ 10 ms	≤ 25 ms	50,000 req/s	99.99%
host-service	≤ 4 ms	≤ 12 ms	≤ 30 ms	30,000 req/s	99.99%
	≤ 2 ms	≤ 8 ms	≤ 20 ms		99.95%

Service	p50 Latency	p95 Latency	p99 Latency	Throughput Target	Availability SLA
session-service				200,000 req/s	
audit-service	≤10 ms (async write)	≤50 ms	≤200 ms	500,000 events/s	99.9%
notification-service	≤20 ms	≤100 ms	≤500 ms	10,000 req/s	99.5%
billing-service	≤50 ms	≤200 ms	≤500 ms	1,000 req/s	99.9%
team-service	≤5 ms	≤20 ms	≤50 ms	20,000 req/s	99.99%
rbac-service	≤1 ms	≤3 ms	≤8 ms	1,000,000 req/s	99.99%
certificate-service	≤50 ms	≤150 ms	≤500 ms	5,000 req/s	99.95%
config-service	≤2 ms	≤5 ms	≤10 ms	100,000 req/s	99.99%
search-service	≤10 ms	≤40 ms	≤100 ms	10,000 req/s	99.9%
sync-service	≤5 ms	≤15 ms	≤40 ms	50,000 req/s	99.9%
websocket-gateway	≤500 μs (frame)	≤2 ms	≤5 ms	200,000 concurrent conns	99.95%
key-distribution	≤5 ms	≤20 ms	≤60 ms	10,000 req/s	99.99%
mfa-service	≤20 ms	≤80 ms	≤200 ms	5,000 req/s	99.99%
metrics-aggregator	≤5 ms	≤20 ms	≤80 ms	internal only	99.5%
log-ingester	≤2 ms	≤8 ms	≤20 ms	1,000,000 events/s	99.5%

Service	p50 Latency	p95 Latency	p99 Latency	Throughput Target	Availability SLA
identity-proxy	≤1 ms	≤3 ms	≤8 ms	500,000 req/s	99.99%
scheduler-service	≤100 ms	≤500 ms	≤2000 ms	100 jobs/s	99.5%
export-service	≤500 ms	≤2000 ms	≤10000 ms	10 req/s	99.0%
plugin-runtime	≤10 ms	≤50 ms	≤200 ms	5,000 req/s	99.5%
ai-assist	≤200 ms	≤800 ms	≤2000 ms	500 req/s	99.0%

SLA Error Budget Calculation:

For a 99.99% SLA over 30 days:

```
Error budget = (1 - 0.9999) × 30 × 24 × 3600 = 259.2 seconds/month
Burn rate 1× = exhausted in 30 days
Burn rate 14.4× = exhausted in 2 days (5-minute window alert threshold)
Burn rate 6× = exhausted in 5 days (1-hour window alert threshold)
```

1.3 Latency Budget Breakdown

The end-to-end latency budget for a **typical authenticated API request** is decomposed into its constituent stages. Each stage has a hard ceiling; exceeding it triggers the circuit breaker for that tier.

Total E2E Budget: 25 ms (p99 for hot path API calls)

Stage	Budget	Cumulative
Client → CDN/Edge	1 ms	1 ms
Edge → API Gateway	1 ms	2 ms
mTLS Handshake (cached session)	0 ms	2 ms ← resumed
JWT Validation (cache hit)	0.5 ms	2.5 ms
Rate Limit Check (Redis)	0.5 ms	3 ms

Request Routing + MW	0.5 ms	3.5 ms
RBAC Policy Eval (cache)	0.5 ms	4 ms
Service Call (gRPC)	1 ms	5 ms
Redis L2 Cache (hit)	0.5 ms	5.5 ms
PostgreSQL Query (indexed)	2 ms	7.5 ms
Kafka Publish (async, fire)	0 ms	7.5 ms ← off-path
Response Serialise (proto)	0.5 ms	8 ms
Service → Gateway	0.5 ms	8.5 ms
Gateway → Client	1 ms	9.5 ms
Nominal hot-path total:	9.5 ms	
Slack to p99 budget:	15.5 ms	

For **SSH terminal frame forwarding** (interactive latency-critical):

Total E2E Budget: 5 ms (p99 round-trip, same datacenter)

Stage	Budget	Cumulative
Client → WebSocket Gateway	0.5 ms	0.5 ms
WS Frame Decode + Routing	0.2 ms	0.7 ms
SSH Proxy goroutine dispatch	0.1 ms	0.8 ms
TCP write to remote host	1.0 ms	1.8 ms
Remote host processing	1.0 ms	2.8 ms
TCP read from remote host	1.0 ms	3.8 ms
SSH Proxy → WS Gateway (pipe)	0.2 ms	4.0 ms
WS Gateway → Client	0.5 ms	4.5 ms
Nominal hot-path total:	4.5 ms	
Slack to p99 budget:	0.5 ms	


1.4 Throughput Targets per Service

Service	req/s (sustained)	req/s (peak burst)	Concurrent Conns	Msgs/s (Kafka in/out)
api-gateway	500,000	1,200,000	50,000	200,000 / 200,000
ssh-proxy	N/A	N/A	100,000 sessions	100,000 / 100,000

Service	req/s (sustained)	req/s (peak burst)	Concurrent Conns	Msgs/s (Kafka in/ out)
websocket- gateway	N/A	N/A	200,000 persistent	200,000 / 200,000
auth- service	100,000	300,000	10,000	50,000 / 50,000
rbac- service	1,000,000	3,000,000	20,000	—
vault- service	50,000	150,000	5,000	20,000 / 5,000
audit- service	—	—	—	— / 500,000
log- ingester	—	—	—	— / 1,000,000

1.5 Performance Testing Pyramid





CI Integration:

```
# .github/workflows/bench.yml
name: Performance Benchmarks
on:
  pull_request:
    paths:
      - '**/*.go'
jobs:
  benchmark:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-go@v5
        with:
          go-version: '1.25'
      - name: Run Benchmarks
        run: |
          go test -bench=. -benchmem -count=5 \
            -benchtime=3s \
            ./... 2>&1 | tee bench-new.txt
      - name: Compare with baseline
        uses: benchmark-action/github-action-benchmark@v1
        with:
          tool: 'go'
          output-file-path: bench-new.txt
          alert-threshold: '105%' # 5% regression triggers failure
          fail-on-alert: true
          github-token: ${ secrets.GITHUB_TOKEN }
```

2. Critical Path Analysis

2.1 SSH Connection Establishment Critical Path

The SSH connection establishment path is the most latency-sensitive user-facing flow. A user perceives "slow SSH" as the time from clicking "Connect" to seeing the first prompt. Target: **< 500 ms** for a warm connection, **< 2 s** for a cold connection (with new certificate issuance).

```
Step 1: Client initiates WebSocket upgrade [0 ms baseline]
├─ HTTP/1.1 → WebSocket upgrade handshake
├─ Budget: 2 ms (same-region), 20 ms (cross-region)
└─ Component: websocket-gateway

Step 2: Session token validation [+2 ms]
├─ JWT decode (Ed25519, cached JWKS): 0.3 ms
├─ Redis session lookup (L2 hit): 0.5 ms
├─ RBAC check for host access (cached): 0.2 ms
└─ Total: 1 ms

Step 3: SSH Proxy slot allocation [+3 ms]
├─ Select least-loaded proxy pod (via registry): 0.2 ms
├─ Allocate goroutine from pool: 0.05 ms
├─ Reserve file descriptor: 0.05 ms
└─ Total: 0.3 ms

Step 4: Retrieve SSH credentials from Vault [+5 ms]
├─ vault-service gRPC call: 1 ms
├─ Redis cache lookup for user's vault key: 0.5 ms
├─ Decrypt credential (AES-256-GCM, 256-byte key): 0.1 ms
└─ Total: 1.6 ms

Step 5: Fetch SSH certificate (if using cert auth) [+10 ms]
├─ certificate-service gRPC call: 2 ms
├─ Check certificate cache (Redis TTL-based): 0.5 ms
│   └─ CACHE HIT: return immediately (0.5 ms total)
│       └─ CACHE MISS: sign new certificate
│           ├── Generate ephemeral key pair: 1 ms
│           ├── HSM signing request: 5 ms (p95: 15 ms)
│           └─ Cache result: 0.2 ms
└─ Total (cache hit): 1 ms | (cache miss): 8.7 ms

Step 6: TCP connect to target host [+30 ms]
├─ DNS resolution (cached): 0.1 ms
├─ TCP SYN/SYN-ACK/ACK: 1–100 ms (network RTT dependent)
├─ SO_KEEPALIVE configured: yes
└─ Typical intra-DC: 1 ms | inter-DC: 20–80 ms

Step 7: SSH handshake (KEX + Auth) [+100 ms]
├─ Protocol version exchange: 1 RTT (~1 ms intra-DC)
```

```
├─ KEX_INIT + algorithm negotiation: 1 RTT
├─ ECDH key exchange (curve25519-sha256): 1 RTT + 0.5 ms
CPU
├─ Server host key verification: 0.1 ms (known_hosts cache)
├─ User authentication (publickey + cert): 1 RTT
└─ Total SSH handshake: 5 RTTs = ~5 ms intra-DC, ~100 ms
cross-region
```

Step 8: Shell request + PTY allocation [+110 ms]

```
├─ SSH_MSG_CHANNEL_REQUEST (pty-req): 1 RTT
├─ SSH_MSG_CHANNEL_REQUEST (shell): 1 RTT
└─ First prompt received: 1 RTT for stdout
```

Step 9: First byte to Flutter client [+112 ms]

```
├─ ssh-proxy → websocket-gateway (gRPC stream): 0.5 ms
├─ websocket-gateway → Flutter client (WS frame): 1 ms
└─ Total: 1.5 ms
```

TOTAL WARM PATH (cert cached, intra-DC):
~115 ms from click to first prompt

TOTAL COLD PATH (no cert cache, cross-region, 50 ms RTT):
~400-600 ms from click to first prompt

Go instrumentation for this path:

```
// helixterm.io/services/ssh-proxy/internal/connection/
// establish.go

package connection

import (
    "context"
    "time"

    "go.opentelemetry.io/otel"
    "go.opentelemetry.io/otel/attribute"
    "go.opentelemetry.io/otel/trace"
    "digital.vasic.observability/metrics"
)

var (
    tracer = otel.Tracer("ssh-proxy.connection.establish")

    connEstablishHistogram = metrics.NewHistogram(
        "helixterm_ssh_connection_establish_duration_seconds",
        "Duration of SSH connection establishment by stage",
        []string{"stage", "cache_status"},
    )
)
```

```

        []float64{0.001, 0.005, 0.01, 0.025, 0.05, 0.1,
        0.25, 0.5, 1.0, 2.0},
    )
)

func EstablishConnection(ctx context.Context, req
    *ConnectRequest) (*Session, error) {
    ctx, span := tracer.Start(ctx, "ssh.establish",
        trace.WithAttributes(
            attribute.String("host.id", req.HostID),
            attribute.String("user.id", req.UserID),
        ),
    )
    defer span.End()

    stages := []struct {
        name string
        fn    func(context.Context) error
    }{
        {"validate_session", func(ctx context.Context) error {
            return validateSession(ctx, req)
        }},
        {"allocate_proxy_slot", func(ctx context.Context)
            error {
                return allocateProxySlot(ctx, req)
            }},
        {"fetch_credentials", func(ctx context.Context)
            error {
                return fetchCredentials(ctx, req)
            }},
        {"fetch_certificate", func(ctx context.Context)
            error {
                return fetchCertificate(ctx, req)
            }},
        {"tcp_connect", func(ctx context.Context) error {
            return tcpConnect(ctx, req)
        }},
        {"ssh_handshake", func(ctx context.Context) error {
            return sshHandshake(ctx, req)
        }},
        {"shell_request", func(ctx context.Context) error {
            return shellRequest(ctx, req)
        }},
    }

    for _, stage := range stages {
        stageCtx, stageSpan := tracer.Start(ctx,
            "ssh.stage."+stage.name)
        start := time.Now()
        if err := stage.fn(stageCtx); err != nil {
            stageSpan.RecordError(err)

```

```

        stageSpan.End()
        connEstablishHistogram.Observe(
            time.Since(start).Seconds(),
            stage.name, "error",
        )
        return nil, fmt.Errorf("stage %s: %w",
            stage.name, err)
    }
    elapsed := time.Since(start)
    connEstablishHistogram.Observe(elapsed.Seconds(),
        stage.name, "ok")
    stageSpan.End()
}

return buildSession(req), nil
}

```

2.2 Vault Sync Critical Path (CRDT Merge + Encryption)

The vault sync path handles bidirectional synchronisation of encrypted vault items using CRDTs (Last-Write-Wins with vector clocks). The performance challenge is that every write requires decryption → merge → re-encryption, which is CPU-bound.

Client triggers sync [0 ms]



Step 1: Compute local diff (client-side, Dart) [0–5 ms]

- ├ Iterate vault items: $O(n)$ where n = item count
- ├ Compare vector clocks: $O(n)$
- ├ Typical vault: 1,000 items → 2 ms
- └ Large vault: 10,000 items → 20 ms

Step 2: gRPC streaming upload of delta [2–20 ms]

- ├ Serialise changed items (Protobuf): $O(k)$ where k = changed items
- ├ Compress payload (Snappy): 0.3 ms per MB
- ├ Stream to sync-service: 1 ms per 10 items (MTU batching)
- └ Budget: ≤ 5 ms for ≤ 100 items

Step 3: sync-service receives and validates [+3 ms]

- ├ JWT + RBAC validation: 1 ms (cached)
- ├ Schema validation (protobuf): 0.2 ms
- └ Total: 1.2 ms

Step 4: Decrypt server-side vault items [+8 ms]

- ├ Fetch user's vault encryption key from vault-service: 1 ms
- └ AES-256-GCM decrypt: 0.1 ms per item

- └ For 100 items: 10 ms total
- └ Parallelised across 8 goroutines: 1.25 ms effective

Step 5: CRDT merge algorithm [+10 ms]

- └ Compare vector clocks $O(\text{items} \times \text{clock_size})$: 0.5 ms
- └ Identify conflicts (LWW resolution): 0.2 ms
- └ Merge winning values: 0.3 ms
- └ Total: 1 ms for 100 items

Step 6: Re-encrypt merged items [+11 ms]

- └ AES-256-GCM encrypt: 0.1 ms per item
- └ Parallelised: 1.25 ms for 100 items
- └ Total: 1.25 ms

Step 7: Write merged items to PostgreSQL [+14 ms]

- └ BEGIN transaction: 0.1 ms
- └ UPSERT vault_items (batch insert): 2 ms for 100 rows
- └ Update vector_clock record: 0.2 ms
- └ COMMIT: 0.5 ms (WAL fsync)
- └ Total: 2.8 ms

Step 8: Publish sync event to Kafka [+15 ms]

- └ Produce to vault.sync.completed topic: 0.5 ms (async)
- └ Consumed by notification-service (async, off-path)
- └ Off critical path – fire and forget

Step 9: Return delta to client [+17 ms]

- └ Query server-side changes since client's last vector clock: 1 ms
- └ Serialise response: 0.3 ms
- └ gRPC stream to client: 1 ms

TOTAL SYNC PATH (100 items, warm cache): ~17 ms
 TOTAL SYNC PATH (1000 items, warm cache): ~65 ms
 TOTAL SYNC PATH (cold start, no cache): ~120 ms

2.3 API Gateway Critical Path

Inbound HTTP/2 request arrives [0 ms]

▼

Step 1: TLS termination (Istio Envoy sidecar) [0.2 ms]

- └ TLS 1.3 session resumption (0-RTT where safe): 0.1 ms
- └ New TLS session: 0.5–2 ms (budget exceeded – use session tickets)
- └ mTLS certificate verification: 0.1 ms (SPIFFE SVID cached 1h)

Step 2: Envoy → api-gateway service [0.3 ms]

- └ HTTP/2 multiplexed stream forwarding: negligible
- └ L7 routing decision in Envoy: 0.1 ms

Step 3: Gin router match [0.1 ms]

- └ Trie-based routing (httprouter under Gin): $O(\log n)$
- └ Method + path match: ~200 ns
- └ Parameter extraction: ~50 ns

Step 4: Middleware pipeline execution [1.0 ms]

- └ Request ID injection: 50 ns
- └ Structured logging setup: 100 ns (deferred actual write)
- └ OpenTelemetry trace context extraction: 150 ns
- └ JWT verification (Ed25519, JWKS cached): 300 μ s
- └ Rate limiter check (Redis sliding window): 500 μ s
- └ RBAC coarse-grained check (token claims): 200 μ s
- └ Total middleware: ~1.3 ms

Step 5: gRPC upstream call [2.0 ms]

- └ Connection from pool (keep-alive): 0 ms
- └ Protobuf marshal: 50–200 μ s
- └ gRPC send: 0.1 ms (H2 frame)
- └ Upstream service processing: 1–5 ms (budget per service)
- └ gRPC receive + unmarshal: 50–200 μ s

Step 6: Response serialisation [0.2 ms]

- └ JSON marshal (encoding/json or sonic): 100–500 μ s
- └ HTTP/2 response write: 0.1 ms

TOTAL API GATEWAY PATH: ~4–8 ms (within 8 ms p95 budget)

2.4 WebSocket Terminal Critical Path

User keypress in Flutter terminal [0 ms]

└

▼

Step 1: Flutter encodes keystroke [0.1 ms]

- └ VT100/ANSI sequence encode: 10 μ s
- └ Dart WebSocket send: 50 μ s
- └ Frame serialisation: 40 μ s

Step 2: Network transit to WebSocket gateway [0.5–50 ms]

- └ Network RTT (datacenter-dependent)

Step 3: WebSocket gateway receives frame [0.2 ms]

- └ WebSocket frame decode (RFC 6455): 50 μ s
- └ Session lookup (in-memory map, $O(1)$): 10 μ s
- └ Route to ssh-proxy goroutine: 20 μ s
- └ Total: 80 μ s

```

Step 4: ssh-proxy goroutine processes frame [0.3 ms]
├─ Write to SSH channel (in-memory pipe): 30 µs
├─ golang.org/x/crypto/ssh channel write: 100 µs
└─ TCP write syscall to remote host: 50 µs

Step 5: Remote host processes input → generates output [1–100 ms]
└─ Shell command execution (network + CPU dependent)

Step 6: ssh-proxy reads SSH channel data [0.2 ms]
├─ TCP read from remote: blocking read (goroutine parked)
├─ SSH channel receive: 50 µs
└─ io.Copy to websocket pipe: 30 µs

Step 7: WebSocket gateway sends to client [0.3 ms]
├─ Frame encode: 40 µs
├─ WebSocket write: 100 µs
└─ Network transit back: 0.5–50 ms

TOTAL ROUND-TRIP (intra-DC, excluding remote processing):
~1.5 ms overhead (HelixTerminator processing only)
~3–5 ms including network (intra-DC)

```

2.5 Kafka Consumer Lag Critical Path

```

Producer publishes message [T+0]
├─ Client serialise (Protobuf): 100 µs
├─ Kafka producer batch (linger.ms=5): 0–5 ms
├─ Network to broker: 1 ms (intra-DC)
└─ Broker ack (acks=1): 1 ms | (acks=all): 3–5 ms

Broker stores message [T+5ms]
├─ Write to partition log (sequential I/O): 0.5 ms
├─ Replicate to ISR (min.insync.replicas=2): 2 ms
└─ Available to consumers: T+3–7 ms

Consumer poll [T+7ms]
├─ Consumer poll interval: 0–100 ms (fetch.min.bytes config)
├─ Network to broker: 1 ms
├─ Fetch response (up to fetch.max.bytes): 1–5 ms
└─ Consumer receives message: T+10–20 ms typical

Consumer processes message [T+20ms]
├─ Deserialise (Protobuf): 100 µs
├─ Business logic: 1–50 ms (service-dependent)
├─ DB write: 2–10 ms
└─ Offset commit: 1 ms (async commit = 0 ms on critical path)

```

Consumer lag definition:

$\text{lag} = \text{latest_offset} - \text{consumer_group_committed_offset}$

Acceptable lag: < 1,000 messages (< 10 ms at 100k msg/s)

Warning threshold: 10,000 messages

Critical threshold: 100,000 messages (signals consumer group falling behind)

2.6 Database Query Critical Path

Go service executes query [T+0]



Step 1: Acquire connection from pgx pool [T+0.1ms]

├ Pool has idle connections: ~0 ms wait

├ Pool exhausted, new connection needed: 5–20 ms (TCP + auth)

└ Pool exhausted + max connections: queue wait (DANGER ZONE)

Step 2: Network transit to PostgreSQL [T+0.2ms]

├ Unix socket (same pod): < 0.1 ms

├ TCP intra-DC (PgBouncer): 0.1–0.5 ms

└ TCP cross-AZ: 1–3 ms

Step 3: PgBouncer connection reuse [T+0.3ms]

├ Transaction-mode pooling: picks server connection from pool

└ Overhead: 0.05–0.1 ms

Step 4: Query parse + plan [T+0.5ms]

├ Parse: < 0.1 ms for prepared statements (cached)

├ Plan: 0.1–2 ms (complex joins can be 10 ms+)

└ Prepared statements CRITICAL for performance

Step 5: Index seek [T+1ms]

├ B-tree index seek $O(\log n)$: < 0.1 ms for 1M rows

├ Index-only scan: < 0.5 ms

└ Sequential scan on 10M rows: 500 ms (UNACCEPTABLE)

Step 6: Row fetch + deserialise [T+2ms]

├ Fetch 100 rows: 0.5 ms

├ pgx row scan: 0.2 ms

└ Total: 0.7 ms

Step 7: Network return [T+3ms]

└ Same as Step 2 inbound

TOTAL HOT PATH (prepared stmt, index hit, 100 rows): ~2–3 ms
TOTAL COLD PATH (sequential scan, 10M rows): 500–1000 ms

3. Identified Bottlenecks & Danger Zones

3.1 SSH Proxy: Goroutine Leaks in Long-Lived Connections

Description: Each SSH session uses at minimum 3 goroutines: one for reading from the SSH channel, one for writing, and one for managing the WebSocket pipe. Under connection churn, goroutines can leak if the cleanup path panics or if `context.Done()` is not propagated correctly through all channel operations.

Detection Method:

```
// Expose goroutine count in Prometheus
// helixterm.io/services/ssh-proxy/internal/metrics/
//      goroutines.go

package metrics

import (
    "runtime"
    "time"
    "github.com/prometheus/client_golang/prometheus"
    "github.com/prometheus/client_golang/prometheus/promauto"
)

var (
    goroutineGauge = promauto.NewGauge(prometheus.GaugeOpts{
        Namespace: "helixterm",
        Subsystem: "ssh_proxy",
        Name:      "goroutines_total",
        Help:      "Current number of goroutines in ssh-proxy process",
    })

    goroutineLeakAlert =
        promauto.NewCounter(prometheus.CounterOpts{
            Namespace: "helixterm",
            Subsystem: "ssh_proxy",
            Name:      "goroutine_leak_detected_total",
            Help:      "Number of times goroutine leak detection
                triggered",
        })
)

func StartGoroutineMonitor(expectedPerSession int) {
    go func() {
```

```

        ticker := time.NewTicker(10 * time.Second)
        defer ticker.Stop()
        for range ticker.C {
            current := runtime.NumGoroutine()
            goroutineGauge.Set(float64(current))
        }
    }()
}

```

Alerting rule:

```

# prometheus/rules/ssh-proxy.yaml
groups:
- name: ssh-proxy-goroutines
  rules:
  - alert: GoroutineLeakDetected
    expr: |
      rate(helixterm_ssh_proxy_goroutines_total[5m]) > 10
      AND
      helixterm_ssh_proxy_active_sessions_total * 3
      < helixterm_ssh_proxy_goroutines_total - 200
    for: 2m
    labels:
      severity: critical
    annotations:
      summary: "Goroutine leak in ssh-proxy
      ({{ $labels.pod }})"
      description: >
        Goroutine count {{ $value }} exceeds expected 3x
        session count.
        Potential goroutine leak detected.

```

Root Cause: Missing `defer cancel()` in goroutine launch sites; `select` statements missing `<-ctx.Done()` arm; error paths that `return` without closing the SSH channel.

Solution:

```

// helixterm.io/services/ssh-proxy/internal/session/
// session.go

package session

import (
    "context"
    "fmt"
    "io"

```

```

"sync"

"golang.org/x/crypto/ssh"
"digital.vasic.concurrency/goroutinepool"
)

// Session encapsulates a single SSH proxy session with
// proper lifecycle management.
type Session struct {
    id            string
    sshConn       ssh.Conn
    channel       ssh.Channel
    wsWriter      io.WriteCloser
    wsReader      io.Reader
    cancel        context.CancelFunc
    wg            sync.WaitGroup
    closeOnce     sync.Once
}

// Run starts the bidirectional forwarding and blocks until
// the session ends.
// It guarantees all goroutines are cleaned up before
// returning.
func (s *Session) Run(parent context.Context) error {
    ctx, cancel := context.WithCancel(parent)
    s.cancel = cancel
    defer s.cleanup() // ALWAYS runs, even on panic

    errCh := make(chan error, 2)

    // Goroutine 1: SSH → WebSocket
    s.wg.Add(1)
    go func() {
        defer s.wg.Done()
        defer cancel() // propagate termination to sibling
        // goroutine
        _, err := copyWithContext(ctx, s.wsWriter, s.channel)
        errCh <- fmt.Errorf("ssh→ws: %w", err)
    }()

    // Goroutine 2: WebSocket → SSH
    s.wg.Add(1)
    go func() {
        defer s.wg.Done()
        defer cancel()
        _, err := copyWithContext(ctx, s.channel, s.wsReader)
        errCh <- fmt.Errorf("ws→ssh: %w", err)
    }()

    // Wait for either direction to fail or context
    // cancellation
    select {

```

```

        case err := <-errCh:
            return err
        case <-ctx.Done():
            return ctx.Err()
    }
}

// cleanup is idempotent and safe to call multiple times.
func (s *Session) cleanup() {
    s.closeOnce.Do(func() {
        if s.cancel != nil {
            s.cancel()
        }
        s.channel.Close()
        s.wsWriter.Close()
        s.wg.Wait() // Wait for goroutines to fully exit
        sessionActiveGauge.Dec()
    })
}

// copyWithContext wraps io.Copy with context cancellation
// support.
func copyWithContext(ctx context.Context, dst io.Writer, src
    io.Reader) (int64, error) {
    buf := copyBufPool.Get().([]byte)
    defer copyBufPool.Put(buf)

    type result struct {
        n    int64
        err error
    }
    ch := make(chan result, 1)
    go func() {
        n, err := io.CopyBuffer(dst, src, buf)
        ch <- result{n, err}
    }()

    select {
    case r := <-ch:
        return r.n, r.err
    case <-ctx.Done():
        return 0, ctx.Err()
    }
}

var copyBufPool = sync.Pool{
    New: func() interface{} {
        buf := make([]byte, 32*1024) // 32 KB, optimal for
        SSH channel reads
        return buf
    }
}

```

```
    },  
  }  
}
```

Impact: Goroutine leak = memory growth of ~8 KB per leaked goroutine stack. At 100k sessions/hour churn rate, a 1% leak rate = 1,000 leaked goroutines/hour = 8 MB/hour steady-state growth → OOMKilled within 24 hours.

3.2 SSH Proxy: File Descriptor Exhaustion

Description: Each SSH connection consumes 3–5 file descriptors (TCP socket to client, TCP socket to remote host, PTY master, PTY slave, signal pipe). At 100,000 concurrent sessions, the proxy requires ≥ 500,000 file descriptors per pod.

Detection:

```
# Monitor FD usage per pod  
kubectl exec -n helixterm ssh-proxy-xxx -- cat /proc/$(pgrep  
    ssh-proxy)/fdinfo | wc -l  
  
# Prometheus metric  
process_open_fds{job="ssh-proxy"} # exposed by Go runtime  
    automatically
```

Root Cause: Default Linux `ulimit -n` is 1,024. Default Kubernetes container FD limit is 65,536. Both are insufficient for 100k sessions.

Solution:

```
# kubernetes/deployments/ssh-proxy.yaml  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: ssh-proxy  
  namespace: helixterm  
spec:  
  template:  
    spec:  
      initContainers:  
        - name: sysctl-tuning  
          image: busybox:1.36  
          command:  
            - sh  
            - -c  
            - |  
              sysctl -w fs.file-max=2000000  
              sysctl -w net.core.somaxconn=65535
```

```

        sysctl -w net.ipv4.tcp_max_syn_backlog=65535
        sysctl -w net.ipv4.ip_local_port_range="1024
65535"
        securityContext:
          privileged: true
containers:
  - name: ssh-proxy
    image: helixterm/ssh-proxy:latest
    resources:
      requests:
        cpu: "4"
        memory: "8Gi"
      limits:
        cpu: "8"
        memory: "16Gi"
    securityContext:
      capabilities:
        add:
          - NET_ADMIN
    # FD limits set via pod security policy
  # Node-level: ensure LimitNOFILE is set in systemd

```

```

# /etc/systemd/system/kubelet.service.d/override.conf
[Service]
LimitNOFILE=2000000
LimitNPROC=unlimited
LimitCORE=infinity

```

```

// helixterm.io/services/ssh-proxy/cmd/main.go
// Programmatic FD limit check at startup
package main

import (
    "log/slog"
    "syscall"
)

func enforceFileDescriptorLimits() {
    var rLimit syscall.Rlimit
    if err := syscall.Getrlimit(syscall.RLIMIT_NOFILE,
        &rLimit); err != nil {
        slog.Error("Failed to get RLIMIT_NOFILE", "error",
            err)
        return
    }
}

```

```

    required := uint64(600_000) // 100k sessions × 5 FDs +
    headroom
    if rLimit.Cur < required {
        slog.Warn("FD limit below recommended",
            "current", rLimit.Cur,
            "recommended", required,
        )
        rLimit.Cur = min(required, rLimit.Max)
        if err := syscall.Setrlimit(syscall.RLIMIT_NOFILE,
            &rLimit); err != nil {
            slog.Error("Failed to set RLIMIT_NOFILE",
                "error", err)
        }
    }
    slog.Info("File descriptor limits configured",
        "soft", rLimit.Cur, "hard", rLimit.Max)
}

```

3.3 SSH Proxy: TCP Backlog Overflow

Description: During SSH connection storms (e.g., mass reconnect after network partition), the TCP `accept()` backlog can fill up, causing new connections to be RST'd before the application even sees them.

Detection:

```

# Check TCP listen overflow counters
ss -lnt | grep :2222
netstat -s | grep "SYNs to LISTEN"
# Prometheus: node_netstat_TcpExt_ListenDrops

```

Solution:

```

// helixterm.io/services/ssh-proxy/internal/server/
// listener.go

package server

import (
    "net"
    "syscall"
)

func newSSHListener(addr string) (net.Listener, error) {
    lc := net.ListenConfig{

```

```

Control: func(network, address string, c
syscall.RawConn) error {
    return c.Control(func(fd uintptr) {
        // SO_REUSEPORT: allow multiple goroutines to
accept on same port
        syscall.SetsockoptInt(int(fd),
syscall.SOL_SOCKET, syscall.SO_REUSEPORT, 1)
        // Increase socket receive buffer
        syscall.SetsockoptInt(int(fd),
syscall.SOL_SOCKET, syscall.SO_RCVBUF, 4*1024*1024)
        // Increase TCP backlog via the kernel param
(set at node level)
        // net.core.somaxconn=65535 (set in
initContainer above)
    })
},
}

return lc.Listen(context.Background(), "tcp", addr)
}

```

3.4 Database Layer: N+1 Queries

Description: The host listing endpoint in `host-service` was found to execute 1 query to list hosts, then N queries to fetch the associated SSH key for each host — one query per host.

Detection:

```

// Detect N+1 at test time using sqlmock query counter
// helixterm.io/services/host-service/internal/repo/
host_repo_test.go

func TestListHostsNoNPlusOne(t *testing.T) {
    db, mock, _ :=
        sqlmock.New(sqlmock.QueryMatcherOption(sqlmock.QueryMatcherRegexp))
    repo := NewHostRepository(db)

    // Expect exactly 1 query (JOIN, not N+1)
    mock.ExpectQuery(`SELECT.*FROM hosts.*JOIN`).
        WillReturnRows(mockHostRows(100))

    _, err := repo.ListByUser(context.Background(),
        "user-123", 100, 0)
    require.NoError(t, err)

    // Assert no additional queries were executed
}

```

```
    assert.Equal(t, 0, mock.ExpectationsWereMet() != nil,
        "N+1 queries detected")
}
```

Root Cause: ORM-style code iterating over results and calling `GetSSHKey(hostID)` in a loop.

Solution:

```
// helixterm.io/services/host-service/internal/repo/
// host_repo.go

package repo

import (
    "context"
    "github.com/jackc/pgx/v5/pgxpool"
)

// ListHostsWithKeys fetches hosts and their SSH keys in a
// single JOIN query.
// Uses a CTE for clarity and avoids N+1.
func (r *HostRepository) ListHostsWithKeys(
    ctx context.Context,
    userID string,
    limit, offset int,
) ([]*HostWithKey, error) {
    const query = `
        WITH user_hosts AS (
            SELECT
                h.id,
                h.hostname,
                h.port,
                h.username,
                h.label,
                h.group_id,
                h.created_at,
                h.updated_at,
                sk.id      AS key_id,
                sk.public_key AS key_public,
                sk.label    AS key_label,
                sk.type     AS key_type
            FROM hosts h
            LEFT JOIN ssh_keys sk ON sk.id = h.ssh_key_id
            WHERE h.user_id = $1
                AND h.deleted_at IS NULL
            ORDER BY h.label ASC
            LIMIT $2 OFFSET $3
        )
    `
}
```

```

        SELECT * FROM user_hosts
    `

    rows, err := r.pool.Query(ctx, query, userID, limit,
        offset)
    if err != nil {
        return nil, fmt.Errorf("ListHostsWithKeys: %w", err)
    }
    defer rows.Close()

    var results []*HostWithKey
    for rows.Next() {
        h := &HostWithKey{}
        if err := rows.Scan(
            &h.ID, &h.Hostname, &h.Port, &h.Username,
            &h.Label,
            &h.GroupID, &h.CreatedAt, &h.UpdatedAt,
            &h.Key.ID, &h.Key.PublicKey, &h.Key.Label,
            &h.Key.Type,
        ); err != nil {
            return nil, fmt.Errorf("scan: %w", err)
        }
        results = append(results, h)
    }
    return results, rows.Err()
}

```

3.5 Database Layer: Connection Pool Starvation

Description: Under high load, all pgx pool connections are in use. New requests queue indefinitely (or until timeout). This manifests as sudden latency spikes despite PostgreSQL itself being healthy.

Detection:

```

// Expose pool stats as Prometheus metrics
// helixterm.io/services/vault-service/internal/db/
// pool_metrics.go

package db

import (
    "github.com/jackc/pgx/v5/pgxpool"
    "github.com/prometheus/client_golang/prometheus"
    "github.com/prometheus/client_golang/prometheus/promauto"
)

var (

```

```

poolAcquireCount =
    promauto.NewCounterVec(prometheus.CounterOpts{
        Namespace: "helixterm",
        Subsystem: "db",
        Name:       "pool_acquire_total",
        Help:       "Total number of pool connection
                    acquisition attempts",
    }, []string{"service", "status"})

poolAcquireDuration =
    promauto.NewHistogramVec(prometheus.HistogramOpts{
        Namespace: "helixterm",
        Subsystem: "db",
        Name:       "pool_acquire_duration_seconds",
        Help:       "Time to acquire a connection from the
                    pool",
        Buckets:    []float64{0.001, 0.005, 0.01, 0.05, 0.1,
                               0.5, 1.0},
    }, []string{"service"})

poolSize = promauto.NewGaugeVec(prometheus.GaugeOpts{
    Namespace: "helixterm",
    Subsystem: "db",
    Name:       "pool_size",
    Help:       "Current pool size (total, idle, in-use)",
}, []string{"service", "state"})
)

func MonitorPool(pool *pgxpool.Pool, serviceName string) {
    go func() {
        ticker := time.NewTicker(5 * time.Second)
        defer ticker.Stop()
        for range ticker.C {
            stats := pool.Stat()
            poolSize.WithLabelValues(serviceName,
                "total").Set(float64(stats.TotalConns()))
            poolSize.WithLabelValues(serviceName,
                "idle").Set(float64(stats.IdleConns()))
            poolSize.WithLabelValues(serviceName,
                "in_use").Set(float64(stats.AcquiredConns()))
        }
    }()
}

```

Solution:

```

// helixterm.io/services/vault-service/internal/db/pool.go

package db

```

```

import (
    "context"
    "time"

    "github.com/jackc/pgx/v5/pgxpool"
)

func NewPool(ctx context.Context, dsn string)
    (*pgxpool.Pool, error) {
    config, err := pgxpool.ParseConfig(dsn)
    if err != nil {
        return nil, err
    }

    config.MaxConns = 100                // Per pod; total =
        pods × 100
    config.MinConns = 10                // Maintain warm
        connections
    config.MaxConnLifetime = 30 * time.Minute
    config.MaxConnIdleTime = 5 * time.Minute
    config.HealthCheckPeriod = 30 * time.Second

    // CRITICAL: acquire timeout prevents indefinite queuing
    config.ConnConfig.ConnectTimeout = 5 * time.Second

    // After-connect hook: set statement timeout per
    // connection
    config.AfterConnect = func(ctx context.Context, conn
        *pgx.Conn) error {
        _, err := conn.Exec(ctx, "SET statement_timeout =
            '30s'")
        return err
    }

    pool, err := pgxpool.NewWithConfig(ctx, config)
    if err != nil {
        return nil, err
    }

    MonitorPool(pool, "vault-service")
    return pool, nil
}

```

3.6 Database Layer: Lock Contention on vault_items

Description: The `vault_items` table uses an advisory lock for CRDT merge operations. Under high concurrency (many devices syncing simultaneously), this creates a serialisation bottleneck.

Detection:

```
-- Detect lock contention in real-time
SELECT
    pid,
    username,
    application_name,
    state,
    wait_event_type,
    wait_event,
    query_start,
    now() - query_start AS duration,
    left(query, 100) AS query_snippet
FROM pg_stat_activity
WHERE wait_event_type = 'Lock'
    AND query NOT LIKE '%pg_stat_activity%'
ORDER BY duration DESC;

-- Check lock waits on vault_items
SELECT
    blocked_locks.pid AS blocked_pid,
    blocking_locks.pid AS blocking_pid,
    blocked_activity.query AS blocked_query,
    blocking_activity.query AS blocking_query
FROM pg_catalog.pg_locks blocked_locks
JOIN pg_catalog.pg_stat_activity blocked_activity
    ON blocked_activity.pid = blocked_locks.pid
JOIN pg_catalog.pg_locks blocking_locks
    ON blocking_locks.locktype = blocked_locks.locktype
    AND blocking_locks.database IS NOT DISTINCT FROM
        blocked_locks.database
    AND blocking_locks.relation IS NOT DISTINCT FROM
        blocked_locks.relation
    AND blocking_locks.granted
    AND NOT blocked_locks.granted
JOIN pg_catalog.pg_stat_activity blocking_activity
    ON blocking_activity.pid = blocking_locks.pid
WHERE blocked_locks.relation::regclass::text = 'vault_items';
```

Solution — Partition-Level Locking with User-Scoped Advisory Locks:

```
// helixterm.io/services/vault-service/internal/repo/
// vault_sync.go

package repo

// MergeVaultItems uses a user-scoped advisory lock instead
// of table-level lock.
```

```

// This allows concurrent syncs for different users while
// serialising per-user.
func (r *VaultRepository) MergeVaultItems(
    ctx context.Context,
    userID string,
    items []*VaultItem,
) error {
    // Hash user ID to a 64-bit advisory lock key
    lockKey := fnv64(userID)

    return r.pool.BeginTxFunc(ctx, pgx.TxOptions{
        IsoLevel:    pgx.ReadCommitted,
        AccessMode: pgx.ReadWrite,
    }, func(tx pgx.Tx) error {
        // Acquire per-user advisory lock (non-blocking
        // timeout)
        var acquired bool
        err := tx.QueryRow(ctx,
            "SELECT pg_try_advisory_xact_lock($1)", lockKey,
        ).Scan(&acquired)
        if err != nil {
            return fmt.Errorf("advisory lock: %w", err)
        }
        if !acquired {
            return ErrSyncConflict // caller can retry after
            backoff
        }

        // Now safe to merge for this user – no other
        // transaction
        // can acquire the same lock until this transaction
        // commits
        return r.mergeItemsInTx(ctx, tx, userID, items)
    })
}

func fnv64(s string) int64 {
    h := fnv.New64a()
    h.Write([]byte(s))
    return int64(h.Sum64())
}

```

3.7 Database Layer: Vacuum Bloat on audit_logs

Description: The `audit_logs` table receives 500,000+ rows per second under peak load. PostgreSQL's MVCC writes new row versions for every UPDATE (rare for audit logs, but DELETE for retention policy causes massive dead-tuple bloat).

Detection:

```

-- Check table bloat
SELECT
    schemaname,
    tablename,
    pg_size_pretty(pg_total_relation_size(relid)) AS
        total_size,
    pg_size_pretty(pg_relation_size(relid)) AS table_size,
    pg_size_pretty(pg_total_relation_size(relid) -
        pg_relation_size(relid)) AS index_size,
    n_dead_tup,
    n_live_tup,
    round(n_dead_tup::numeric / nullif(n_live_tup +
        n_dead_tup, 0) * 100, 2) AS dead_pct,
    last_vacuum,
    last_autovacuum,
    last_analyze,
    last_autoanalyze
FROM pg_stat_user_tables
WHERE tablename = 'audit_logs'
ORDER BY n_dead_tup DESC;

```

Solution — Declarative Partitioning with DROP PARTITION:

```

-- Partition audit_logs by month; DROP old partitions instead
-- of DELETE
-- This eliminates VACUUM requirement entirely for expired
-- data.

CREATE TABLE audit_logs (
    id          UUID          NOT NULL DEFAULT
        gen_random_uuid(),
    user_id     UUID          NOT NULL,
    service     VARCHAR(64)   NOT NULL,
    action      VARCHAR(128)  NOT NULL,
    resource_id VARCHAR(255),
    ip_address  INET,
    metadata    JSONB,
    created_at  TIMESTAMPTZ   NOT NULL DEFAULT now()
) PARTITION BY RANGE (created_at);

-- Create partitions for current + next 3 months
CREATE TABLE audit_logs_2026_06 PARTITION OF audit_logs
    FOR VALUES FROM ('2026-06-01') TO ('2026-07-01');

CREATE TABLE audit_logs_2026_07 PARTITION OF audit_logs
    FOR VALUES FROM ('2026-07-01') TO ('2026-08-01');

-- Retention: drop partitions older than 90 days

```

```
-- Run as a scheduled job (scheduler-service)
-- DROP TABLE audit_logs_2026_03; -- instantaneous, no VACUUM
needed
```

3.8 Kafka: Consumer Lag Under Burst

Description: When SSH connection rate spikes (e.g., after scheduled maintenance window ends and thousands of users reconnect), Kafka producers flood the `ssh.events` topic. Consumer groups fall behind, causing audit and session data to lag.

Detection:

```
# prometheus/rules/kafka.yaml
groups:
- name: kafka-consumer-lag
  rules:
    - alert: KafkaConsumerLagHigh
      expr: |
        kafka_consumer_group_lag{topic="ssh.events"} >
        10000
      for: 2m
      labels:
        severity: warning
      annotations:
        summary: "Kafka consumer lag {{ $value }} on
        {{ $labels.topic }}"

    - alert: KafkaConsumerLagCritical
      expr: |
        kafka_consumer_group_lag{topic="ssh.events"} >
        100000
      for: 5m
      labels:
        severity: critical
```

Solution — Dynamic Consumer Scaling:

```
// helixterm.io/services/audit-service/internal/kafka/
consumer_manager.go

package kafka

import (
    "context"
    "sync"
```

```

        "time"

        "github.com/IBM/sarama"
        "digital.vasic.concurrency/workerpool"
    )

    // ConsumerManager dynamically scales consumers based on lag.
    type ConsumerManager struct {
        client      sarama.Client
        group       sarama.ConsumerGroup
        topic       string
        handler     sarama.ConsumerGroupHandler
        minWorkers  int
        maxWorkers  int
        currentW    int
        workerPool  *workerpool.Pool
        mu          sync.Mutex
    }

    // AutoScale monitors lag and scales consumers accordingly.
    func (cm *ConsumerManager) AutoScale(ctx context.Context) {
        ticker := time.NewTicker(30 * time.Second)
        defer ticker.Stop()

        for {
            select {
            case <-ctx.Done():
                return
            case <-ticker.C:
                lag, err := cm.measureLag(ctx)
                if err != nil {
                    continue
                }

                cm.mu.Lock()
                desired := cm.calculateDesiredWorkers(lag)
                if desired != cm.currentW {
                    cm.scaleWorkers(ctx, desired)
                }
                cm.mu.Unlock()
            }
        }
    }

    func (cm *ConsumerManager) calculateDesiredWorkers(lag
        int64) int {
        switch {
        case lag < 1_000:
            return cm.minWorkers
        case lag < 10_000:
            return cm.minWorkers * 2

```

```
    case lag < 100_000:
        return cm.minWorkers * 4
    default:
        return cm.maxWorkers
    }
}
```

3.9 Kafka: Rebalancing Storms

Description: Frequent Kafka consumer group rebalances (triggered by pod restarts, HPA scale-in, or GC pauses exceeding `max.poll.interval.ms`) cause all consumers in the group to stop processing during the rebalance.

Root Cause: Default `max.poll.interval.ms=300000` (5 min) is too lenient; `session.timeout.ms=45000` (45s) is too short for GC-paused Go pods.

Solution:

```
// helixterm.io/services/audit-service/internal/kafka/
// config.go

package kafka

import "github.com/IBM/sarama"

func NewConsumerConfig() *sarama.Config {
    cfg := sarama.NewConfig()

    // Use cooperative sticky rebalancing (Kafka 2.4+)
    // Only reassigns partitions that need to move, not all
    // partitions
    cfg.Consumer.Group.Rebalance.GroupStrategies =
        []sarama.BalanceStrategy{
            sarama.NewBalanceStrategyCooperativeSticky(),
        }

    // Increase session timeout to tolerate GC pauses
    cfg.Consumer.Group.Session.Timeout = 90 * time.Second

    // Heartbeat must be < session.timeout / 3
    cfg.Consumer.Group.Heartbeat.Interval = 20 * time.Second

    // Max poll interval: time between successive calls to
    // poll()
    // Must account for worst-case batch processing time
    cfg.Consumer.MaxProcessingTime = 60 * time.Second
}
```

```

    // Static group membership: prevents rebalance on planned
    // restart
    // Requires Kafka 2.3+
    cfg.Consumer.Group.Member.UserData = []byte("ssh-proxy-
        pod-0") // per-pod static ID

    return cfg
}

```

3.10 Kafka: Large Message Impact

Description: Terminal session recordings published to the `terminal.recordings` topic can exceed 1 MB when flushed. Large messages cause broker memory pressure and increase replica lag.

Solution:

```

// helixterm.io/services/ssh-proxy/internal/recording/
// publisher.go

package recording

const (
    MaxKafkaMessageSize = 900 * 1024 // 900 KB hard ceiling
    ChunkSize           = 256 * 1024 // 256 KB per chunk
)

// PublishRecordingChunked splits large recordings into
// chunks.
func (p *RecordingPublisher) PublishRecordingChunked(
    ctx context.Context,
    sessionID string,
    data []byte,
) error {
    if len(data) <= MaxKafkaMessageSize {
        return p.publishSingle(ctx, sessionID, data, 0, 1)
    }

    totalChunks := (len(data) + ChunkSize - 1) / ChunkSize
    for i := 0; i < totalChunks; i++ {
        start := i * ChunkSize
        end := min(start+ChunkSize, len(data))
        chunk := data[start:end]

        if err := p.publishSingle(ctx, sessionID, chunk, i,
            totalChunks); err != nil {
            return fmt.Errorf("chunk %d/%d: %w", i,
                totalChunks, err)
        }
    }
}

```

```
    }  
  }  
  return nil  
}
```

3.11 RabbitMQ: Queue Depth Explosion Under SSH Storm

Description: The `ssh.connect.requests` RabbitMQ queue accumulates millions of messages when the SSH proxy worker pool is saturated during a connection storm.

Detection:

```
# Monitor via RabbitMQ management API  
curl -s -u guest:guest http://rabbitmq:15672/api/queues/%2F/  
ssh.connect.requests | \  
jq '{messages: .messages, consumers: .consumer_count,  
    rate: .messages_details.rate}'
```

Solution — Queue Overflow with Dead-Letter Exchange:

```
// helixterm.io/services/ssh-proxy/internal/rabbitmq/  
queue_setup.go  
  
package rabbitmq  
  
import amqp "github.com/rabbitmq/amqp091-go"  
  
func DeclareSSHConnectQueue(ch *amqp.Channel) error {  
    // Dead-letter exchange for overflow  
    if err := ch.ExchangeDeclare(  
        "ssh.connect.dlx", "direct",  
        true, false, false, false, nil,  
    ); err != nil {  
        return err  
    }  
  
    _, err := ch.QueueDeclare(  
        "ssh.connect.requests",  
        true, // durable  
        false, // auto-delete  
        false, // exclusive  
        false, // no-wait  
        amqp.Table{  
            "x-max-length":  
                100k messages, int32(100_000), // max  
        },  
    )  
    if err != nil {  
        return err  
    }  
}
```

```

        "x-overflow": "reject-publish", //
        backpressure to producer
        "x-dead-letter-exchange": "ssh.connect.dlx",
        "x-message-ttl": int32(30_000), // 30s
        TTL; stale connect requests discarded
        "x-queue-type":
        "quorum", // Quorum queues for HA
    },
)
return err
}

```

3.12 Redis: Hot Key Problem on Session Tokens

Description: Session tokens are validated on every API request. If all requests are for the same (or few) user sessions, a single Redis node receives all traffic — a classic hot key problem.

Detection:

```

# Enable Redis hotkeys monitor (Redis 4+)
redis-cli --hotkeys -i 0.1

# Prometheus: monitor per-key request rate (requires Redis
# keyspace notifications)
# helixterm_redis_key_hit_rate{key_prefix="sess:"} > 10000/s
# on single node

```

Solution — Local L1 Cache for Session Tokens:

```

// helixterm.io/services/api-gateway/internal/auth/
// session_cache.go

package auth

import (
    "time"
    "github.com/dgraph-io/ristretto"
    "digital.vasic.cache"
)

var sessionL1Cache *ristretto.Cache

func init() {
    var err error

```

```

    sessionL1Cache, err =
        ristretto.NewCache(&ristretto.Config{
            NumCounters: 1_000_000, // 1M counters for
            frequency estimation
            MaxCost:      256 * 1024 * 1024, // 256 MB L1 session
            cache
            BufferItems: 64,
            Metrics:      true,
        })
    if err != nil {
        panic(err)
    }
}

// GetSession implements L1→L2 fallback with short TTL to
// limit stale reads.
func GetSession(ctx context.Context, tokenHash string)
(*Session, error) {
    // L1 hit: avoid Redis call entirely
    if val, ok := sessionL1Cache.Get(tokenHash); ok {
        sessionCacheHits.WithLabelValues("l1").Inc()
        return val.(*Session), nil
    }

    // L2 miss: fetch from Redis
    sess, err := redisGetSession(ctx, tokenHash)
    if err != nil {
        return nil, err
    }

    // Populate L1 with short TTL (30s) to cap staleness
    sessionL1Cache.SetWithTTL(tokenHash, sess, 1,
        30*time.Second)
    sessionCacheHits.WithLabelValues("l2").Inc()
    return sess, nil
}

```

3.13 Redis: Lua Script Blocking

Description: Distributed rate limiter Lua scripts block the Redis event loop for the duration of their execution. Scripts exceeding 5 ms block all other Redis commands.

Solution — Async Rate Limiter with Local Token Bucket:

See Section 9.2 for the full rate limiter implementation. Key principle: execute Lua only once per 100ms window (batched updates) rather than on every request.

3.14 API Gateway: JWT Verification CPU Cost

Description: Under 500,000 req/s, JWT verification using RS256 (RSA-2048) costs ~2 ms per verification and saturates 4 CPU cores. Ed25519 reduces this to ~0.1 ms, but key rotation must be handled.

Solution — Ed25519 + JWKS Caching:

```
// helixterm.io/services/auth-service/internal/jwt/
// verifier.go

package jwt

import (
    "crypto/ed25519"
    "time"
    "github.com/golang-jwt/jwt/v5"
    "github.com/dgraph-io/ristretto"
)

type CachingVerifier struct {
    jwksURL    string
    keyCache   *ristretto.Cache
    parser     *jwt.Parser
}

func NewCachingVerifier(jwksURL string) *CachingVerifier {
    cache, _ := ristretto.NewCache(&ristretto.Config{
        NumCounters: 100,
        MaxCost:     1024 * 1024, // 1 MB key cache
        BufferItems: 64,
    })

    return &CachingVerifier{
        jwksURL: jwksURL,
        keyCache: cache,
        parser: jwt.NewParser(
            jwt.WithValidMethods([]string{"EdDSA"}), //
            Ed25519 only
            jwt.WithExpirationRequired(),
            jwt.WithIssuedAt(),
        ),
    }
}

func (v *CachingVerifier) Verify(tokenStr string) (*Claims, error) {
    token, err := v.parser.ParseWithClaims(
        tokenStr,
        &Claims{},
    )
}
```

```

        v.keyFunc,
    )
    if err != nil {
        return nil, err
    }
    return token.Claims.(*Claims), nil
}

func (v *CachingVerifier) keyFunc(token *jwt.Token)
(interface{}, error) {
    kid, _ := token.Header["kid"].(string)

    // Cache hit: avoid JWKS fetch
    if key, ok := v.keyCache.Get(kid); ok {
        return key.(ed25519.PublicKey), nil
    }

    // Cache miss: fetch JWKS and cache all keys
    keys, err := fetchJWKS(v.jwksURL)
    if err != nil {
        return nil, err
    }

    for id, k := range keys {
        v.keyCache.SetWithTTL(id, k, 1, 15*time.Minute)
    }

    if key, ok := keys[kid]; ok {
        return key, nil
    }
    return nil, fmt.Errorf("key ID %q not found in JWKS",
        kid)
}

```

Benchmark:

```

// helixterm.io/services/auth-service/internal/jwt/
// verifier_bench_test.go

func BenchmarkJWTVerifyRS256(b *testing.B) {
    token := generateTestTokenRS256()
    verifier := newRS256Verifier()
    b.ResetTimer()
    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            if _, err := verifier.Verify(token); err != nil {
                b.Fatal(err)
            }
        }
    })
}

```

```

    })
}
// Result: 450 ns/op (Ed25519 cached) vs 2,100,000 ns/op
// (RS256, no cache)

func BenchmarkJWTVerifyEd25519Cached(b *testing.B) {
    token := generateTestTokenEd25519()
    verifier := NewCachingVerifier(testJWKSServer.URL)
    // warm cache
    verifier.Verify(token)
    b.ResetTimer()
    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            if _, err := verifier.Verify(token); err != nil {
                b.Fatal(err)
            }
        }
    })
}
// Result: 380 ns/op, 2 allocs/op

```

3.15 Istio: sidecar Proxy CPU Overhead

Description: Each Envoy sidecar intercepts all inbound and outbound traffic, adding L7 processing overhead. At high request rates, sidecar CPU usage can exceed the application pod's own CPU.

Detection:

```

# Compare app vs sidecar CPU usage
kubectl top pods -n helixterm --containers | awk '{print $1,
    $2, $3}'
# Expected: sidecar should be < 30% of total pod CPU

```

Solution — Selective Sidecar Injection + Traffic Exclusion:

```

# For ultra-high-throughput services (ssh-proxy, websocket-
# gateway),
# bypass Envoy for data-plane traffic while keeping control-
# plane mTLS.

apiVersion: networking.istio.io/v1alpha3
kind: Sidecar
metadata:
  name: ssh-proxy-sidecar
  namespace: helixterm

```

```
spec:
  workloadSelector:
    labels:
      app: ssh-proxy
  ingress:
    - port:
        number: 2222
        protocol: TCP
        name: ssh
      defaultEndpoint: "127.0.0.1:2222"
      captureMode: NONE # Bypass Envoy for SSH data plane
  egress:
    - captureMode: IPTABLES # Keep Envoy for outbound
      service calls
    hosts:
      - "helixterm/*"
```

3.16 Flutter Client: Excessive Terminal Widget Rebuild Cycles

Description: The terminal widget re-renders on every keystroke and every output byte received. At 10,000 characters/second output from a chatty remote process, this causes 10,000 `setState` calls per second, exceeding Flutter's 60/120 fps refresh budget.

Solution: See Section 10 for full Flutter performance analysis.

4. Gap Analysis — What's Missing vs. Production-Ready

4.1 Missing Circuit Breakers

Circuit breakers are absent from the following critical inter-service call paths. Without them, a single slow downstream service causes cascading timeouts that exhaust all upstream goroutine pools.

Call Path	Calling Service	Called Service	Risk	Circuit Breaker Type
Vault decrypt	sync-service	vault-service	HIGH	Half-open with 3-probe
Certificate issue	ssh-proxy	certificate-service	HIGH	Fallback to cached cert

Call Path	Calling Service	Called Service	Risk	Circuit Breaker Type
RBAC eval	api-gateway	rbac-service	CRITICAL	Fallback to last-known policy
Session validate	api-gateway	auth-service	CRITICAL	Deny new sessions, allow renewals
Kafka produce	all services	Kafka broker	HIGH	Drop non-critical events
Redis get	all services	Redis cluster	MEDIUM	Fallback to DB
HSM sign	certificate-service	HSM	HIGH	Fallback to software signing

Implementation:

```
// helixterm.io/services/api-gateway/internal/circuitbreaker/
registry.go

package circuitbreaker

import (
    "time"
    "github.com/sony/gobreaker/v2"
    "digital.vasic.recovery"
)

type Registry struct {
    breakers map[string]*gobreaker.CircuitBreaker[any]
}

func NewRegistry() *Registry {
    r := &Registry{
        breakers:
            make(map[string]*gobreaker.CircuitBreaker[any]),
    }

    // RBAC service: fail-open with cached policy
    r.register("rbac-service", gobreaker.Settings{
```

```

        Name:      "rbac-service",
        MaxRequests: 3,           // probe requests in half-
        open
        Interval:   60 * time.Second,
        Timeout:    30 * time.Second,
        ReadyToTrip: func(counts gobreaker.Counts) bool {
            return counts.ConsecutiveFailures > 5 ||
                (counts.Requests > 100 &&
                    float64(counts.TotalFailures)/
                    float64(counts.Requests) > 0.5)
        },
        OnStateChange: func(name string, from, to
            gobreaker.State) {
            circuitBreakerStateGauge.WithLabelValues(name,
                to.String()).Set(1)
            if to == gobreaker.StateOpen {
                // Notify on-call
                alertCircuitBreakerOpen(name)
            }
        },
    })
})

// Certificate service: fallback to cached cert on open
r.register("certificate-service", gobreaker.Settings{
    Name:      "certificate-service",
    MaxRequests: 1,
    Interval:   30 * time.Second,
    Timeout:    60 * time.Second,
    ReadyToTrip: func(counts gobreaker.Counts) bool {
        return counts.ConsecutiveFailures > 3
    },
})

return r
}

func (r *Registry) Execute(
    name string,
    fn func() (any, error),
    fallback func() (any, error),
) (any, error) {
    cb, ok := r.breakers[name]
    if !ok {
        return fn() // no breaker configured: fail through
    }

    result, err := cb.Execute(fn)
    if err != nil && fallback != nil {
        // Circuit open or function failed: use fallback
        return fallback()
    }
}

```

```
    return result, err
}
```

4.2 Missing Read Replicas for Specific Services

The following services perform read-heavy workloads against primary PostgreSQL but have not been configured to use read replicas.

Service	Read/Write Ratio	Missing Replica Routing	Estimated Primary Load Reduction
host-service (list hosts)	95% read	List, search operations	60%
vault-service (fetch items)	80% read	Get vault items, list folders	50%
audit-service (reporting)	99% read	All reporting queries	80%
search-service (full-text)	100% read	All search queries	100%
session-service (active list)	70% read	List active sessions	40%

Solution:

```
// helixterm.io/services/host-service/internal/db/
// replica_router.go

package db

import (
    "context"
    "github.com/jackc/pgx/v5/pgxpool"
)

type ReplicaRouter struct {
    primary *pgxpool.Pool
    replicas []*pgxpool.Pool
    rr      atomic.Uint64 // round-robin counter
}

func (r *ReplicaRouter) ReadPool(ctx context.Context)
    *pgxpool.Pool {
```

```

    // Check if caller explicitly requested primary
    if ForceUsePrimary(ctx) {
        return r.primary
    }
    // Round-robin across replicas
    if len(r.replicas) == 0 {
        return r.primary
    }
    idx := r.rr.Add(1) % uint64(len(r.replicas))
    return r.replicas[idx]
}

func (r *ReplicaRouter) WritePool() *pgxpool.Pool {
    return r.primary
}

// ForceUsePrimary is used after a write to ensure read-your-
// writes consistency.
type contextKey int
const forcePrimaryKey contextKey = 1

func WithForcePrimary(ctx context.Context) context.Context {
    return context.WithValue(ctx, forcePrimaryKey, true)
}

func ForceUsePrimary(ctx context.Context) bool {
    v, _ := ctx.Value(forcePrimaryKey).(bool)
    return v
}

```

4.3 Missing Cache Warming Strategies

Upon pod restart or initial deployment, all caches are cold. This causes a thundering herd against the database as all pods simultaneously miss on every cache key.

Cache	Service	Missing Warm Strategy	Cold Start Impact
Session token cache	api-gateway	Pre-populate from Redis on start	10× DB load spike for 5 min
JWKS key cache	api-gateway	Fetch JWKS at startup	JWT verify fails until fetch completes
RBAC policy cache	rbac-service	Load all active policies at boot	Policy eval falls to DB for 2 min

Cache	Service	Missing Warm Strategy	Cold Start Impact
Host list cache	host-service	Warm top-1000 hosts per tenant	Cold DB queries on every request
Vault key cache	vault-service	Pre-load vault encryption keys	Decrypt fails until warmed

Solution:

```
// helixterm.io/services/rbac-service/internal/startup/
// cache_warmer.go

package startup

import (
    "context"
    "time"
    "golang.org/x/sync/errgroup"
)

type CacheWarmer struct {
    policyRepo    PolicyRepository
    policyCache   PolicyCache
    batchSize     int
    parallelism   int
}

// WarmCaches pre-populates critical caches before the
// service starts accepting traffic.
// Called during readiness probe delay window.
func (w *CacheWarmer) WarmCaches(ctx context.Context) error {
    g, gctx := errgroup.WithContext(ctx)
    g.SetLimit(w.parallelism)

    // Warm RBAC policies (most critical)
    g.Go(func() error {
        return w.warmRBACPolicies(gctx)
    })

    // Warm role assignments
    g.Go(func() error {
        return w.warmRoleAssignments(gctx)
    })

    if err := g.Wait(); err != nil {
        return fmt.Errorf("cache warming failed: %w", err)
    }
}
```

```

        cacheWarmupCompleted.Set(1)
        return nil
    }

    func (w *CacheWarmer) warmRBACPolicies(ctx context.Context)
        error {
        offset := 0
        for {
            policies, err := w.policyRepo.ListAll(ctx,
                w.batchSize, offset)
            if err != nil {
                return err
            }
            if len(policies) == 0 {
                break
            }
            for _, p := range policies {
                w.policyCache.Set(p.ID, p)
            }
            offset += len(policies)
            time.Sleep(1 * time.Millisecond) // yield to other
            goroutines
        }
        return nil
    }
}

```

4.4 Missing Connection Pooling Configurations

Service	Missing Config	Default (Bad)	Recommended
audit-service → Kafka	Producer pool	New producer per goroutine	Shared producer, 16 goroutines
sync-service → Redis	Pool size	10	100 (matches request concurrency)
ssh-proxy → gRPC upstreams	Connection pool	1 per target	5 per target (round-robin)
api-gateway → gRPC backends	Max idle	2	20 per upstream
All services → PgBouncer	Pool mode	session	transaction (95% reduction in PG connections)

4.5 Missing Backpressure Mechanisms

Component	Missing Backpressure	Risk
WebSocket gateway	No inbound frame rate limiting per connection	OOM from fast producer clients
SSH proxy goroutine pool	No queue with bounded depth	Unbounded goroutine growth
Kafka producer buffer	No blocking on <code>buffer.memory</code> exhaustion	Silently drops events
audit-service consumer	No admission control on DB write queue	DB overload cascades
Flutter client	Unbounded terminal output buffer	Client OOM on <code>cat /dev/urandom</code>

Solution — SSH Proxy Bounded Queue:

```
// helixterm.io/services/ssh-proxy/internal/pool/
// goroutine_pool.go

package pool

import (
    "context"
    "errors"
    "sync"
    "digital.vasic.concurrency/semaphore"
)

var ErrPoolFull = errors.New("goroutine pool at capacity:
    backpressure applied")

type BoundedPool struct {
    sem      *semaphore.Weighted
    maxSize  int
    metrics  *poolMetrics
}

func NewBoundedPool(maxSize int) *BoundedPool {
    return &BoundedPool{
        sem:      semaphore.NewWeighted(int64(maxSize)),
        maxSize:  maxSize,
        metrics:  newPoolMetrics(),
    }
}
```

```

    }
}

// Submit attempts to run fn in the pool, returning
// ErrPoolFull if at capacity.
func (p *BoundedPool) Submit(ctx context.Context, fn func())
    error {
    // Non-blocking acquire: apply backpressure immediately
    // rather than queueing
    if !p.sem.TryAcquire(1) {
        p.metrics.rejected.Inc()
        return ErrPoolFull
    }

    go func() {
        defer p.sem.Release(1)
        defer func() {
            if r := recover(); r != nil {
                p.metrics.panics.Inc()
                // Log panic but don't re-panic (pool must
                // stay healthy)
            }
        }()
        fn()
    }()
    return nil
}

```

4.6 Missing Graceful Degradation Paths

Feature	Degraded Mode	Trigger Condition	Degraded Behaviour
Vault sync	Read-only mode	vault-service circuit open	Return last-known vault, block writes
SSH authentication	Cert-auth fallback	certificate-service down	Fall back to pubkey auth
RBAC evaluation	Permissive mode	rbac-service down	Use cached last-known policy
Audit logging	Best-effort write	audit-service overloaded	Buffer in Redis, retry async
Full-text search	Prefix search	search-service slow	Fall back to PostgreSQL LIKE

Feature	Degraded Mode	Trigger Condition	Degraded Behaviour
Real-time terminal	Buffered terminal	websocket-gateway degraded	Store in memory, flush when recovered

4.7 Missing Bulkhead Patterns

Bulkheads isolate resource pools so that a surge in one traffic type doesn't starve another.

```
// helixterm.io/services/api-gateway/internal/bulkhead/
// registry.go

package bulkhead

import "digital.vasic.concurrency/semaphore"

// Bulkheads prevent one API category from starving another.
type Registry struct {
    SSH          *semaphore.Weighted // SSH-related API calls
    Vault        *semaphore.Weighted // Vault CRUD operations
    Admin        *semaphore.Weighted // Admin operations
    Background   *semaphore.Weighted // Batch/background jobs
}

func NewRegistry() *Registry {
    return &Registry{
        SSH:          semaphore.NewWeighted(50_000), // 50k
                concurrent SSH API calls
        Vault:        semaphore.NewWeighted(20_000), // 20k
                vault operations
        Admin:        semaphore.NewWeighted(1_000),  // 1k
                admin calls (low priority)
        Background:  semaphore.NewWeighted(500),    // 500
                batch operations
    }
}
```

4.8 Missing Timeout Hierarchies

Every network call must have a timeout. Currently, many services use the default Go HTTP client (no timeout = indefinite hang).

Operation	Missing Timeout	Required Timeout	Implementation
External DNS resolution	None	2 s	<code>net.Dialer{Timeout: 2*time.Second}</code>
TCP connect to SSH target	None	10 s	<code>net.DialTimeout</code>
Redis command	None	500 ms	<code>redis.NewClient(&redis.Options{ReadTimeout: 500ms})</code>
gRPC call (per-RPC)	None	5 s	<code>ctx, cancel = context.WithTimeout(ctx, 5*time.Second)</code>
PostgreSQL query	None	30 s	<code>SET statement_timeout = '30s'</code>
Kafka produce ack	None	5 s	<code>Producer.Timeout = 5*time.Second</code>
HTTP client calls	None	10 s	<code>http.Client{Timeout: 10*time.Second}</code>
SSH KEX timeout	None	30 s	<code>ssh.ClientConfig{Timeout: 30*time.Second}</code>
Full request end-to-end	None	60 s	Top-level request context

4.9 Missing Retry Budgets

Unbounded retries under load amplify failures into retry storms. Each service needs a capped retry budget.

```
// helixterm.io/services/api-gateway/internal/retry/budget.go

package retry

import (
    "context"
    "time"
    "golang.org/x/time/rate"
)
```

```

// Budget enforces a rate-limited retry policy.
// Prevents retry storms by capping the aggregate retry rate
// across all callers.
type Budget struct {
    maxAttempts int
    baseDelay   time.Duration
    maxDelay    time.Duration
    limiter     *rate.Limiter // aggregate retry rate limit
    jitter      float64       // 0.0-1.0
}

func NewBudget(maxAttempts int, base, max time.Duration, rps
    float64) *Budget {
    return &Budget{
        maxAttempts: maxAttempts,
        baseDelay:   base,
        maxDelay:    max,
        limiter:     rate.NewLimiter(rate.Limit(rps),
            maxAttempts),
        jitter:      0.2,
    }
}

func (b *Budget) Do(ctx context.Context, fn func() error)
    error {
    var lastErr error
    for attempt := 0; attempt < b.maxAttempts; attempt++ {
        if err := fn(); err == nil {
            return nil
        } else {
            lastErr = err
            if !IsRetryable(err) {
                return err
            }
        }

        delay := b.backoff(attempt)

        // Check retry budget before sleeping
        if !b.limiter.Allow() {
            return fmt.Errorf("retry budget exhausted: %w",
                lastErr)
        }

        select {
        case <-ctx.Done():
            return ctx.Err()
        case <-time.After(delay):
        }
    }
}

```

```

    return fmt.Errorf("max attempts (%d) exceeded: %w",
        b.maxAttempts, lastErr)
}

func (b *Budget) backoff(attempt int) time.Duration {
    delay := b.baseDelay * time.Duration(1<<attempt)
    if delay > b.maxDelay {
        delay = b.maxDelay
    }
    // Add jitter: ±20%
    jitter := time.Duration(float64(delay) * b.jitter *
        (rand.Float64()*2 - 1))
    return delay + jitter
}

```

5. Database Performance Deep-Dive

5.1 PostgreSQL Tuning Parameters

The following configuration is tuned for a **dedicated PostgreSQL server** with 64 CPU cores, 256 GB RAM, and NVMe SSD storage (XFS filesystem, 3.5 GB/s sequential write).

```

# /etc/postgresql/17/main/postgresql.conf
# HelixTerminator Production PostgreSQL Configuration
# Server: 64 vCPU, 256 GB RAM, 4× 2 TB NVMe SSD (RAID-10
#   effective: 2× write throughput)

#-----
# MEMORY
#-----
shared_buffers
    = 64GB                                # 25% of RAM; primary data
    cache
effective_cache_size = 192GB              # Hint for planner:
    75% of RAM
work_mem = 256MB                          # Per sort/hash
    operation; 64 vCPU × 8 max_parallel × 256MB = 131GB
    worst case
                                           # Monitor: SET
    work_mem per session for bulk ops
maintenance_work_mem = 4GB               # For VACUUM, CREATE
    INDEX, REINDEX
huge_pages = try                          # Enable transparent
    huge pages for shared_buffers
temp_buffers = 64MB                      # Per-session
    temporary table cache

```

```

#-----
# CONNECTION MANAGEMENT
#-----
max_connections = 500                # PgBouncer sits in
    front; real app sees pooled conns
superuser_reserved_connections = 10

#-----
# WAL & DURABILITY
#-----
wal_level = replica                  # Required for
    streaming replication
wal_buffers = 64MB                   # Should be ≥
    (max_wal_size / 100); auto is usually fine
max_wal_size = 8GB                   # Larger reduces
    checkpoint frequency
min_wal_size = 1GB
checkpoint_completion_target = 0.9    # Spread checkpoint I/
    0 over 90% of checkpoint interval
checkpoint_timeout = 15min
synchronous_commit = on              # Set to 'local' for
    async replicas if latency required
wal_compression = lz4                # Compress WAL;
    reduces I/O at cost of CPU
full_page_writes = on                # Required for crash
    safety

#-----
# QUERY PLANNER
#-----
default_statistics_target = 500      # More stats → better
    plans; default is 100
random_page_cost = 1.1               # NVMe: near
    sequential cost; default 4.0 is for HDD
seq_page_cost = 1.0
effective_io_concurrency = 200        # NVMe can handle
    200+ concurrent I/O requests
parallel_tuple_cost = 0.1
parallel_setup_cost = 1000
max_parallel_workers_per_gather
    = 8    # Use up to 8 parallel workers per query
max_parallel_workers = 32
max_parallel_maintenance_workers = 8

#-----
# AUTOVACUUM
#-----
autovacuum = on
autovacuum_max_workers = 8           # More workers for
    high-write workloads
autovacuum_naptime = 30s             # Check for needed
    vacuums every 30s (default 1min)

```

```
autovacuum_vacuum_threshold
    = 1000      # Minimum dead tuples before vacuum
                (default 50)
autovacuum_vacuum_scale_factor = 0.01 # 1% dead tuples
                triggers vacuum (default 0.20)
autovacuum_analyze_threshold = 500
autovacuum_analyze_scale_factor = 0.005 # 0.5% changed rows
                triggers analyze
autovacuum_vacuum_cost_delay
    = 2ms      # Less throttling for hot tables
autovacuum_vacuum_cost_limit = 800      # Default is 200;
                higher = faster vacuum
log_autovacuum_min_duration = 250ms    # Log slow
                autovacuums
```

```
#-----
# LOGGING
#-----
log_min_duration_statement = 100ms      # Log queries slower
                than 100 ms
log_checkpoints = on
log_connections = off                  # Too verbose with
                PgBouncer
log_disconnections = off
log_lock_waits = on
log_temp_files = 0                    # Log all temp file
                creation
log_line_prefix = '%t [%p]: [%l-1]
                user=%u,db=%d,app=%a,client=%h '
log_statement = 'ddl'                  # Log all DDL
                statements
```

```
#-----
# LOCK MANAGEMENT
#-----
lock_timeout = 10s                     # Fail rather than
                wait indefinitely
deadlock_timeout = 500ms                # Detect deadlocks
                quickly
```

```
#-----
# pg_stat_statements
#-----
shared_preload_libraries = 'pg_stat_statements,auto_explain'
pg_stat_statements.max = 10000
pg_stat_statements.track = all
pg_stat_statements.save = on

auto_explain.log_min_duration = 500ms
auto_explain.log_analyze = on
auto_explain.log_buffers = on
```

```
auto_explain.log_timing = on
auto_explain.log_nested_statements = on
```

5.2 Index Strategy for Critical Tables

vault_items

```
-- Primary access patterns:
-- 1. Fetch all items for a user_id (most common)
-- 2. Get specific item by id + user_id (security check)
-- 3. Search by label (full-text)
-- 4. Get items modified since last sync (CRDT sync)
-- 5. List by folder

CREATE TABLE vault_items (
  id                UUID                NOT NULL DEFAULT
    gen_random_uuid(),
  user_id           UUID                NOT NULL,
  folder_id         UUID,
  type              VARCHAR(32) NOT NULL,
  label             TEXT                NOT NULL,
  encrypted_data    BYTEA               NOT NULL,
  encryption_iv     BYTEA               NOT NULL,
  vector_clock      JSONB               NOT NULL DEFAULT '{}',
  checksum          VARCHAR(64) NOT NULL,
  deleted_at        TIMESTAMPTZ,
  created_at        TIMESTAMPTZ NOT NULL DEFAULT now(),
  updated_at        TIMESTAMPTZ NOT NULL DEFAULT now(),
  CONSTRAINT vault_items_pkey PRIMARY KEY (id)
) PARTITION BY HASH (user_id); -- Hash partition by user_id
  for even distribution

-- 16 hash partitions (adjust based on user count)
CREATE TABLE vault_items_0 PARTITION OF vault_items FOR
  VALUES WITH (MODULUS 16, REMAINDER 0);
-- ... create vault_items_1 through vault_items_15

-- Index 1: Primary lookup – fetch all items for a user (with
  soft-delete filter)
CREATE INDEX CONCURRENTLY idx_vault_items_user_active
  ON vault_items (user_id, updated_at DESC)
  WHERE deleted_at IS NULL
  INCLUDE (id, label, type, folder_id); -- Covering index
  for list views

-- Index 2: Security check – verify item ownership before
  decrypt
CREATE UNIQUE INDEX CONCURRENTLY idx_vault_items_id_user
  ON vault_items (id, user_id);
```

```

-- Index 3: Folder lookup
CREATE INDEX CONCURRENTLY idx_vault_items_folder
    ON vault_items (user_id, folder_id)
    WHERE deleted_at IS NULL;

-- Index 4: Sync delta – items updated since timestamp
-- CRITICAL: This is the most performance-sensitive query
--           during sync
CREATE INDEX CONCURRENTLY idx_vault_items_sync
    ON vault_items (user_id, updated_at)
    WHERE deleted_at IS NULL;

-- Index 5: Full-text search on label
CREATE INDEX CONCURRENTLY idx_vault_items_label_fts
    ON vault_items USING gin (to_tsvector('english', label))
    WHERE deleted_at IS NULL;

```

audit_logs

```

-- Partitioned table (see Section 3.7)
-- Access patterns: time-range queries, user activity,
--                 service activity

-- Index per partition (PostgreSQL 10+ automatically creates
--                       on partitions)
CREATE INDEX CONCURRENTLY idx_audit_logs_user_time
    ON audit_logs (user_id, created_at DESC);

CREATE INDEX CONCURRENTLY idx_audit_logs_service_action
    ON audit_logs (service, action, created_at DESC);

CREATE INDEX CONCURRENTLY idx_audit_logs_resource
    ON audit_logs (resource_id)
    WHERE resource_id IS NOT NULL;

-- BRIN index on created_at for time-range scans (extremely
--           small, fast for sequential data)
CREATE INDEX CONCURRENTLY idx_audit_logs_brin_time
    ON audit_logs USING brin (created_at)
    WITH (pages_per_range = 128);

```

ssh_sessions

```
CREATE TABLE ssh_sessions (  
    id                UUID                NOT NULL DEFAULT  
        gen_random_uuid(),  
    user_id           UUID                NOT NULL,  
    host_id           UUID                NOT NULL,  
    proxy_pod         VARCHAR(64) NOT NULL,  
    status            VARCHAR(16) NOT NULL DEFAULT 'active',  
        -- active|closed|failed  
    bytes_sent        BIGINT              NOT NULL DEFAULT 0,  
    bytes_received    BIGINT              NOT NULL DEFAULT 0,  
    started_at        TIMESTAMPTZ NOT NULL DEFAULT now(),  
    ended_at          TIMESTAMPTZ,  
    CONSTRAINT ssh_sessions_pkey PRIMARY KEY (id, started_at)  
) PARTITION BY RANGE (started_at);  
  
-- Index 1: List active sessions per user  
CREATE INDEX CONCURRENTLY idx_ssh_sessions_user_active  
    ON ssh_sessions (user_id, started_at DESC)  
    WHERE status = 'active';  
  
-- Index 2: Sessions per host (for access analytics)  
CREATE INDEX CONCURRENTLY idx_ssh_sessions_host  
    ON ssh_sessions (host_id, started_at DESC);  
  
-- Index 3: Sessions per proxy pod (for drain detection)  
CREATE INDEX CONCURRENTLY idx_ssh_sessions_proxy  
    ON ssh_sessions (proxy_pod)  
    WHERE status = 'active';
```

hosts

```
-- Index 1: Primary list view  
CREATE INDEX CONCURRENTLY idx_hosts_user_label  
    ON hosts (user_id, label)  
    WHERE deleted_at IS NULL;  
  
-- Index 2: Group-based listing  
CREATE INDEX CONCURRENTLY idx_hosts_group  
    ON hosts (user_id, group_id, label)  
    WHERE deleted_at IS NULL;  
  
-- Index 3: Full-text search across hostname + label  
CREATE INDEX CONCURRENTLY idx_hosts_fts  
    ON hosts USING gin (
```

```

        to_tsvector('english', coalesce(hostname, '') || ' '
        || coalesce(label, ''))
    )
WHERE deleted_at IS NULL;

```

users

```

-- Index 1: Auth lookup (most critical – every login)
CREATE UNIQUE INDEX CONCURRENTLY idx_users_email
    ON users (lower(email))
    WHERE deleted_at IS NULL;

-- Index 2: SSO lookup by provider + subject
CREATE UNIQUE INDEX CONCURRENTLY idx_users_sso
    ON users (sso_provider, sso_subject)
    WHERE sso_provider IS NOT NULL;

-- Index 3: Team membership queries
CREATE INDEX CONCURRENTLY idx_users_team
    ON users (team_id)
    WHERE deleted_at IS NULL;

```

5.3 Partitioning Strategy

```

-- audit_logs: RANGE partition by month, retain 90 days
-- Automated partition management via pg_partman

SELECT partman.create_parent(
    p_parent_table := 'public.audit_logs',
    p_control      := 'created_at',
    p_type         := 'range',
    p_interval     := 'monthly',
    p_premake      := 3,           -- Pre-create 3 future
    partitions
    p_retention    := '90 days', -- Auto-drop partitions older
    than 90 days
    p_retention_keep_table := false
);

-- ssh_sessions: RANGE partition by month, retain 1 year
SELECT partman.create_parent(
    p_parent_table := 'public.ssh_sessions',
    p_control      := 'started_at',
    p_type         := 'range',
    p_interval     := 'monthly',
    p_premake      := 2,

```

```

        p_retention := '1 year',
        p_retention_keep_table := false
    );

-- Maintenance job (runs nightly at 02:00 UTC via scheduler-
    service)
SELECT partman.run_maintenance_proc();

```

5.4 PgBouncer Configuration

```

# /etc/pgbouncer/pgbouncer.ini
# PgBouncer 1.22+ recommended

[databases]
helixterm = host=postgres-primary port=5432 dbname=helixterm
           \
           pool_size=100 reserve_pool=20

helixterm_ro = host=postgres-replica-1,postgres-replica-2 \
               port=5432 dbname=helixterm pool_size=200

[pgbouncer]
# Listening
listen_addr = 0.0.0.0
listen_port = 5432
unix_socket_dir = /var/run/postgresql

# Authentication
auth_type = scram-sha-256
auth_file = /etc/pgbouncer/userlist.txt

# Pool mode – CRITICAL: transaction mode for microservices
# Session mode would defeat the purpose of pooling
pool_mode = transaction

# Limits
max_client_conn
    = 10000      # Total client connections across all
                pools
default_pool_size = 50      # Default per database/user pair
reserve_pool_size = 10      # Emergency reserve
reserve_pool_timeout = 5    # Seconds before using reserve

# Timeouts
connect_timeout = 15        # Seconds to wait for backend
                           connection
client_idle_timeout = 600   # Disconnect idle clients after
                           10 min

```

```

server_idle_timeout = 300      # Return idle server connections
                                after 5 min
server_lifetime = 3600        # Recycle server connections
                                after 1 hour
query_timeout = 60            # Kill queries exceeding 60
                                seconds
client_login_timeout = 10

# TLS
server_tls_sslmode = require
server_tls_ca_file = /etc/ssl/certs/ca-bundle.crt
client_tls_sslmode = require
client_tls_key_file = /etc/pgbouncer/server.key
client_tls_cert_file = /etc/pgbouncer/server.crt

# Logging
log_connections = 0
log_disconnections = 0
log_pooler_errors = 1
stats_period = 60

# Admin
admin_users = pgbouncer_admin
stats_users = pgbouncer_stats

```

5.5 Read Replica Routing Logic

```

// helixterm.io/services/vault-service/internal/db/router.go

package db

import (
    "context"
    "database/sql"
    "sync/atomic"

    "github.com/jackc/pgx/v5/pgxpool"
)

type PoolRouter struct {
    primary    *pgxpool.Pool
    replicas   []*pgxpool.Pool
    rrCounter  atomic.Uint64
}

type QueryMode int

const (
    ModeReadWrite QueryMode = iota

```

```

        ModeReadOnly
        ModeReadOnlyStrong // Requires reading from primary
                             (read-your-writes)
    )

    type ctxKey int
    const queryModeKey ctxKey = 0

    func WithQueryMode(ctx context.Context, mode QueryMode)
        context.Context {
        return context.WithValue(ctx, queryModeKey, mode)
    }

    func GetQueryMode(ctx context.Context) QueryMode {
        mode, _ := ctx.Value(queryModeKey).(QueryMode)
        return mode
    }

    func (r *PoolRouter) Pool(ctx context.Context) *pgxpool.Pool
    {
        mode := GetQueryMode(ctx)
        switch mode {
        case ModeReadWrite, ModeReadOnlyStrong:
            return r.primary
        case ModeReadOnly:
            if len(r.replicas) == 0 {
                return r.primary
            }
            // Round-robin across healthy replicas
            for attempts := 0; attempts < len(r.replicas);
            attempts++ {
                idx := r.rrCounter.Add(1) %
                uint64(len(r.replicas))
                replica := r.replicas[idx]
                if r.isHealthy(replica) {
                    return replica
                }
            }
            // All replicas unhealthy: fall back to primary
            return r.primary
        default:
            return r.primary
        }
    }

    func (r *PoolRouter) isHealthy(pool *pgxpool.Pool) bool {
        ctx, cancel := context.WithTimeout(context.Background(),
            100*time.Millisecond)
        defer cancel()
    }

```

```
    return pool.Ping(ctx) == nil
}
```

5.6 Slow Query Detection

```
-- Top 20 slowest queries by total execution time
SELECT
    round(total_exec_time::numeric, 2) AS total_exec_ms,
    calls,
    round(mean_exec_time::numeric, 2) AS mean_exec_ms,
    round(stddev_exec_time::numeric, 2) AS stddev_ms,
    round((100 * total_exec_time / sum(total_exec_time)
        OVER())::numeric, 2) AS pct_total,
    rows,
    round(rows::numeric / calls, 2) AS rows_per_call,
    shared_blks_hit,
    shared_blks_read,
    round(shared_blks_hit::numeric /
        nullif(shared_blks_hit + shared_blks_read, 0) * 100,
        2) AS cache_hit_pct,
    left(query, 120) AS query_snippet
FROM pg_stat_statements
WHERE calls > 100
ORDER BY total_exec_time DESC
LIMIT 20;

-- Alerting: export pg_stat_statements via postgres_exporter
-- Prometheus alert for newly slow queries
```

```
# prometheus/rules/postgres.yaml
groups:
  - name: postgres-slow-queries
    rules:
      - alert: SlowQueryDetected
        expr: |
          rate(pg_stat_statements_mean_exec_time_ms[5m]) >
          100
        for: 5m
        labels:
          severity: warning
        annotations:
          summary: "Slow PostgreSQL query detected (mean >
          100ms)"
```

5.7 VACUUM Strategy

```
-- Per-table autovacuum overrides for high-write tables
ALTER TABLE audit_logs SET (
    autovacuum_vacuum_scale_factor = 0.005,
    autovacuum_vacuum_threshold = 5000,
    autovacuum_vacuum_cost_delay = 1,
    autovacuum_vacuum_cost_limit = 2000,
    autovacuum_analyze_scale_factor = 0.01,
    toast.autovacuum_vacuum_scale_factor = 0.005
);

ALTER TABLE vault_items SET (
    autovacuum_vacuum_scale_factor = 0.01,
    autovacuum_vacuum_threshold = 1000,
    autovacuum_vacuum_cost_delay = 2,
    autovacuum_vacuum_cost_limit = 1000
);

ALTER TABLE ssh_sessions SET (
    autovacuum_vacuum_scale_factor = 0.02,
    autovacuum_vacuum_threshold = 2000,
    autovacuum_vacuum_cost_delay = 2
);

-- Manual VACUUM schedule for hot tables (run during off-peak
    hours 03:00-05:00 UTC)
-- Executed by scheduler-service weekly
VACUUM (ANALYZE, VERBOSE, PARALLEL 4) vault_items;
VACUUM (ANALYZE, VERBOSE, PARALLEL 8) ssh_sessions;

-- Freeze old tuples to prevent transaction ID wraparound
VACUUM (FREEZE, ANALYZE, VERBOSE) users;
```

5.8 EXPLAIN ANALYZE for Critical Queries

```
-- Query 1: List vault items for sync (most critical)
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT id, type, label, encrypted_data, encryption_iv,
       vector_clock, updated_at
FROM vault_items
WHERE user_id = '550e8400-e29b-41d4-a716-446655440000'
    AND updated_at > '2026-06-27 00:00:00+00'
    AND deleted_at IS NULL
ORDER BY updated_at DESC
LIMIT 100;
```

```

-- Expected plan:
-- Index Scan using idx_vault_items_sync on vault_items
    (cost=0.57..8.84 rows=23 width=285)
--   (actual time=0.045..0.123 rows=23 loops=1)
--   Index Cond: ((user_id = '550e8400... '::uuid) AND
    (updated_at > '2026-06-27...'))
--   Filter: (deleted_at IS NULL)
--   Buffers: shared hit=8
-- Planning Time: 0.215 ms
-- Execution Time: 0.158 ms ✓ (within 2 ms budget)

-- Query 2: SSH session list for user
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT s.id, s.host_id, s.status, s.started_at, h.hostname,
       h.label
FROM ssh_sessions s
JOIN hosts h ON h.id = s.host_id
WHERE s.user_id = '550e8400-e29b-41d4-a716-446655440000'
      AND s.status = 'active'
ORDER BY s.started_at DESC
LIMIT 20;

-- Expected: Index Scan on idx_ssh_sessions_user_active +
--           Nested Loop on hosts
-- Execution Time: ~0.8 ms ✓

-- Query 3: RBAC policy lookup (must be < 1 ms)
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT p.actions, p.resources, p.conditions
FROM rbac_policies p
JOIN role_assignments ra ON ra.policy_id = p.id
WHERE ra.user_id = $1
      AND ra.team_id = $2
      AND p.active = true;

-- Query 4: Audit log time-range query (reporting)
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT action, count(*), min(created_at), max(created_at)
FROM audit_logs
WHERE user_id = $1
      AND created_at BETWEEN '2026-06-01' AND '2026-06-28'
GROUP BY action
ORDER BY count(*) DESC;

-- Expected: Partition pruning to audit_logs_2026_06 only
-- Index Scan on idx_audit_logs_user_time
-- Execution Time: ~2 ms ✓

-- Query 5: Full-text vault search
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT id, label, type, folder_id

```

```

FROM vault_items
WHERE user_id = $1
      AND to_tsvector('english', label) @@
           plainto_tsquery('english', 'production server')
      AND deleted_at IS NULL
LIMIT 20;

-- Expected: Bitmap Index Scan on idx_vault_items_label_fts
-- Execution Time: ~3 ms (dependent on result count)

-- Query 6: Get host with SSH key (JOIN – no N+1)
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT h.id, h.hostname, h.port, h.username,
       sk.encrypted_private_key
FROM hosts h
LEFT JOIN ssh_keys sk ON sk.id = h.ssh_key_id
WHERE h.id = $1
      AND h.user_id = $2
      AND h.deleted_at IS NULL;

-- Query 7: Bulk audit log insert (from audit-service
           consumer)
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
INSERT INTO audit_logs (user_id, service, action,
                       resource_id, metadata)
SELECT
    unnest($1::uuid[]),
    unnest($2::text[]),
    unnest($3::text[]),
    unnest($4::text[]),
    unnest($5::jsonb[])
RETURNING id;

-- Expected: Multi-row insert directly to partition
-- Execution Time: ~5 ms for 1,000 rows ✓

-- Query 8: Certificate existence check (cache miss path)
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT id, cert_data, expires_at
FROM ssh_certificates
WHERE user_id = $1
      AND host_id = $2
      AND expires_at > now() + interval '5 minutes'
ORDER BY expires_at DESC
LIMIT 1;

-- Query 9: Team member listing with permissions
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT u.id, u.email, u.display_name, ra.role
FROM users u
JOIN role_assignments ra ON ra.user_id = u.id

```

```

WHERE ra.team_id = $1
  AND u.deleted_at IS NULL
ORDER BY u.display_name ASC;

-- Query 10: Session analytics rollup
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT
  date_trunc('hour', started_at) AS hour,
  count(*) AS session_count,
  sum(bytes_sent + bytes_received) AS total_bytes,
  avg(extract(epoch FROM (ended_at - started_at))) AS
    avg_duration_seconds
FROM ssh_sessions
WHERE started_at >= now() - interval '24 hours'
  AND status = 'closed'
GROUP BY 1
ORDER BY 1;


-- Expected: Parallel Seq Scan on recent partitions only
-- Execution Time: ~50 ms with partitioning ✓

```

6. Caching Strategy & Cache Architecture

6.1 Cache Architecture Overview

HelixTerminator uses a three-tier cache hierarchy:



For static assets + cacheable API resp
Latency: 1–5 ms (edge)
TTL: 1h–7d Invalidation: by tag

↓ All miss

Origin: PostgreSQL / Source of Truth
Latency: 1–10 ms

6.2 L1 In-Process Cache (ristretto)

```
// helixterm.io/services/vault-service/internal/cache/l1.go
// Uses digital.vasic.cache submodule

package cache

import (
    "time"
    "github.com/dgraph-io/ristretto"
    "digital.vasic.cache"
)

// L1Cache is a per-pod in-process cache backed by ristretto.
// Thread-safe, LFU eviction, supports TTL.
type L1Cache struct {
    inner    *ristretto.Cache
    metrics *cacheMetrics
}

func NewL1Cache(maxBytes int64) (*L1Cache, error) {
    c, err := ristretto.NewCache(&ristretto.Config{
        NumCounters: maxBytes / 100, // 10× expected items
        MaxCost:     maxBytes,
        BufferItems:  64,
        Cost: func(value interface{}) int64 {
            return int64(estimateSize(value))
        },
        Metrics: true,
        OnEvict: func(item *ristretto.Item) {
            l1Evictions.Inc()
        },
    })
    if err != nil {
        return nil, err
    }
    return &L1Cache{inner: c, metrics:
        newCacheMetrics("l1")}, nil
}
```

```

}

func (c *L1Cache) Get(key string) (interface{}, bool) {
    val, ok := c.inner.Get(key)
    if ok {
        c.metrics.hits.Inc()
    } else {
        c.metrics.misses.Inc()
    }
    return val, ok
}

func (c *L1Cache) SetWithTTL(key string, value interface{},
    ttl time.Duration) bool {
    return c.inner.SetWithTTL(key, value, 0, ttl)
}

// sync.Map for ultra-hot keys (JWKS, global config)
// 0(1) read with no lock contention after initial population
var globalHotCache sync.Map

func GetHot(key string) (interface{}, bool) {
    return globalHotCache.Load(key)
}

func SetHot(key string, value interface{}) {
    globalHotCache.Store(key, value)
}

```

6.3 L2 Redis Cluster Configuration

```

# kubernetes/redis/cluster-config.yaml
# 6-node Redis cluster: 3 primary + 3 replica
# 16,384 slots distributed across 3 primaries

apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: redis-cluster
  namespace: helixterm
spec:
  replicas: 6
  serviceName: redis-cluster
  template:
    spec:
      containers:
        - name: redis
          image: redis:7.4-alpine
          command:

```



```

lazyfree-lazy-eviction yes                # Async eviction
      (avoid blocking)
lazyfree-lazy-expire yes
lazyfree-lazy-server-del yes
replica-lazy-flush yes

# Slow log
slowlog-log-slower-than 10000            # Log commands > 10 ms
slowlog-max-len 256

# Client buffer limits
client-output-buffer-limit normal 0 0 0
client-output-buffer-limit replica 256mb 64mb 60
client-output-buffer-limit pubsub 32mb 8mb 60

```

Redis Key Design:

```

# Key naming conventions (namespace:entity:id:field)
sess:{token_hash}           → Session object (TTL: 24h)
vault:key:{user_id}         → Vault encryption key (TTL:
1h)
host:list:{user_id}         → Serialised host list (TTL:
5min)
rbac:policy:{user_id}:{team_id} → RBAC policy set (TTL: 5min)
cert:{user_id}:{host_id}    → SSH certificate (TTL: cert
expiry - 5min)
ratelimit:{user_id}:{endpoint} → Rate limit counter (TTL:
window duration)
ws:route:{session_id}       → WebSocket pod routing (TTL:
session lifetime)
singleflight:{key_hash}     → Cache stampede lock (TTL:
5s)

```

6.4 Cache Invalidation Strategies

Data Type	Strategy	Rationale
Session tokens	TTL + event-driven delete on logout	Logout must be immediate
Vault items	Write-through on every update	Consistency required for sync
Host list	Write-through + TTL (5 min)	Stale reads acceptable briefly

Data Type	Strategy	Rationale
RBAC policies	Event-driven (Kafka policy.changed topic)	Security-critical
SSH certificates	TTL (cert expiry - 5 min)	Cert rotation is time-based
JWKS keys	TTL (15 min) + event-driven on key rotation	Key rotation is rare

```
// helixterm.io/services/vault-service/internal/cache/
// invalidation.go

package cache

// WriteThrough writes to both cache and DB atomically.
// Cache write failure is non-fatal (DB is source of truth).
func WriteThrough(
    ctx context.Context,
    key string,
    value interface{},
    ttl time.Duration,
    dbWrite func() error,
) error {
    // Always write to DB first
    if err := dbWrite(); err != nil {
        return fmt.Errorf("db write: %w", err)
    }

    // Then update cache (non-critical path)
    if err := redisSet(ctx, key, value, ttl); err != nil {
        // Cache write failure is acceptable – DB is source
        // of truth
        cacheWriteErrors.WithLabelValues("write_through").Inc()
        slog.Warn("Cache write-through failed", "key", key,
            "error", err)
        // Invalidate rather than leave stale
        redisDel(ctx, key)
    }
    return nil
}
```

6.5 Cache Stampede Prevention

```
// helixterm.io/services/api-gateway/internal/cache/
// singleflight.go
// Uses digital.vasic.cache singleflight utilities

package cache

import (
    "context"
    "time"
    "golang.org/x/sync/singleflight"
    "digital.vasic.cache/probabilistic"
)

var sfGroup singleflight.Group

// GetOrFetch implements singleflight + probabilistic early
// expiry to prevent stampede.
func GetOrFetch(
    ctx context.Context,
    key string,
    ttl time.Duration,
    fetch func(ctx context.Context) (interface{}, error),
) (interface{}, error) {
    // L1 check with probabilistic early expiry
    // XFetch algorithm: recompute before TTL expires with
    // probability
    // proportional to remaining TTL (avoids hard expiry
    // cliff)
    if val, ok := probabilisticGet(key, ttl); ok {
        return val, nil
    }

    // Deduplicate concurrent misses via singleflight
    result, err, _ := sfGroup.Do(key, func() (interface{},
        error) {
        // Double-check after singleflight serialisation
        if val, ok := llCache.Get(key); ok {
            return val, nil
        }

        // Actual fetch (DB or upstream service)
        val, err := fetch(ctx)
        if err != nil {
            return nil, err
        }

        // Populate cache
        llCache.SetWithTTL(key, val, ttl)
    })
    return result, err
}
```

```

        redisSetAsync(key, val, ttl) // async to avoid
        blocking caller
        return val, nil
    })

    return result, err
}

// probabilisticGet implements XFetch cache stampede
// prevention.
// Returns (value, true) if cache hit, or (nil, false) for
// miss/probabilistic recompute.
func probabilisticGet(key string, ttl time.Duration)
    (interface{}, bool) {
    entry, ok := llCache.GetWithExpiry(key)
    if !ok {
        return nil, false
    }

    remaining := entry.Expiry.Sub(time.Now())
    if remaining <= 0 {
        return nil, false
    }

    // XFetch: recompute early with probability based on
    // remaining TTL
    // beta = 1.0 (standard), delta = fetch duration estimate
    delta := 10 * time.Millisecond // estimated fetch time
    beta := 1.0
    threshold := float64(delta) * beta *
        math.Log(rand.Float64())
    if float64(-remaining) <= threshold {
        // Probabilistically trigger early recompute
        return nil, false
    }

    return entry.Value, true
}

```

6.6 Full Cache Manager

```

// helixterm.io/services/vault-service/internal/cache/
// manager.go
// Orchestrates L1 + L2 with digital.vasic.cache

package cache

import (
    "context"

```

```

"encoding/json"
"time"

"github.com/redis/go-redis/v9"
"digital.vasic.cache"
)

// Manager is the unified cache facade for vault-service.
type Manager struct {
    l1 *L1Cache
    l2 *redis.ClusterClient
    sf singleflight.Group
}

func NewManager(l1 *L1Cache, l2 *redis.ClusterClient)
    *Manager {
    return &Manager{l1: l1, l2: l2}
}

// Get implements L1→L2→origin with stampede protection.
func (m *Manager) Get(ctx context.Context, key string, dest
    interface{}) (bool, error) {
    // L1 hit
    if val, ok := m.l1.Get(key); ok {
        return true, json.Unmarshal(val.([]byte), dest)
    }

    // L2 hit
    data, err := m.l2.Get(ctx, key).Bytes()
    if err == nil {
        m.l1.SetWithTTL(key, data, 30*time.Second) //
        Populate L1
        return true, json.Unmarshal(data, dest)
    }
    if err != redis.Nil {
        return false, fmt.Errorf("redis get: %w", err)
    }

    return false, nil // Cache miss
}

// Set populates both L1 and L2.
func (m *Manager) Set(ctx context.Context, key string, value
    interface{}, ttl time.Duration) error {
    data, err := json.Marshal(value)
    if err != nil {
        return err
    }

    m.l1.SetWithTTL(key, data, min(ttl, 30*time.Second))

```

```

        if err := m.l2.Set(ctx, key, data, ttl).Err(); err !=
            nil {
            cacheWriteErrors.WithLabelValues("l2_set").Inc()
            return fmt.Errorf("redis set: %w", err)
        }
        return nil
    }

    // Delete invalidates from both tiers.
    func (m *Manager) Delete(ctx context.Context, key string)
        error {
        m.l1.Delete(key)
        return m.l2.Del(ctx, key).Err()
    }

    // GetOrFetch is the preferred high-level accessor.
    func (m *Manager) GetOrFetch(
        ctx context.Context,
        key string,
        ttl time.Duration,
        dest interface{},
        fetch func() (interface{}, error),
    ) error {
        if ok, err := m.Get(ctx, key, dest); ok || err != nil {
            return err
        }

        // Singleflight to prevent stampede
        result, err, _ := m.sf.Do(key, func() (interface{},
            error) {
            // Re-check after singleflight serialisation
            if ok, _ := m.Get(ctx, key, dest); ok {
                return dest, nil
            }
            return fetch()
        })
        if err != nil {
            return err
        }

        return m.Set(ctx, key, result, ttl)
    }

```

7. SSH Proxy Performance Engineering

7.1 Goroutine Pool Sizing

```
// helixterm.io/services/ssh-proxy/internal/pool/sizing.go
// Uses digital.vasic.concurrency submodule

package pool

import (
    "runtime"
    "digital.vasic.concurrency/adaptivpool"
)

// PoolConfig returns optimal pool configuration based on
// available resources.
func PoolConfig() adaptivpool.Config {
    cpus := runtime.NumCPU()

    return adaptivpool.Config{
        // Goroutine pool for SSH session handlers
        // Each session uses 2 goroutines (read + write)
        // Budget: 100k sessions × 2 = 200k goroutines
        // Each goroutine stack: ~8 KB initial, grows as
        // needed
        // 200k × 8 KB = 1.6 GB minimum goroutine overhead
        MinWorkers: cpus * 4,
        MaxWorkers: 200_000, // Hard cap; shed load above
        this

        // Work queue depth: bounded to apply backpressure
        QueueDepth: 10_000,

        // Goroutine stack size hint (Go runtime handles
        // actual sizing)
        // Reduce initial allocation for leaf goroutines
        StackHint: 8 * 1024,

        // Dynamic scaling
        ScaleUpThreshold: 0.7, // Scale up when 70%
        utilised
        ScaleDownThreshold: 0.3, // Scale down when below 30%
        ScaleInterval: 10 * time.Second,
    }
}
```

7.2 io.Copy vs Custom Buffer Strategy

The choice of buffer size has significant impact on throughput for SSH data forwarding.

```
// helixterm.io/services/ssh-proxy/internal/forward/
// bench_test.go

package forward

import (
    "bytes"
    "io"
    "testing"
)

// BenchmarkForwardSmallBuf tests standard io.Copy with 32 KB
// buffer
func BenchmarkForwardSmallBuf(b *testing.B) {
    data := bytes.Repeat([]byte("x"), 1*1024*1024) // 1 MB
    b.SetBytes(int64(len(data)))
    b.ResetTimer()
    for i := 0; i < b.N; i++ {
        src := bytes.NewReader(data)
        dst := io.Discard
        buf := make([]byte, 32*1024)
        io.CopyBuffer(dst, src, buf)
    }
}
// Result: ~10 GB/s throughput, 0 allocs/op (with pool)

// BenchmarkForwardLargeBuf tests with 256 KB buffer (better
// for high-latency links)
func BenchmarkForwardLargeBuf(b *testing.B) {
    data := bytes.Repeat([]byte("x"), 1*1024*1024)
    b.SetBytes(int64(len(data)))
    b.ResetTimer()
    for i := 0; i < b.N; i++ {
        src := bytes.NewReader(data)
        dst := io.Discard
        buf := make([]byte, 256*1024)
        io.CopyBuffer(dst, src, buf)
    }
}
// Result: ~12 GB/s throughput (20% improvement for large
// transfers)

// Optimal strategy: adaptive buffer size based on session
// type
func adaptiveBuffer(sessionType SessionType) []byte {
```

```

switch sessionType {
case SessionTypeInteractive:
    // Interactive sessions: small buffer for low latency
    // Prefer 4 KB to match typical keystroke payloads
    buf := smallBufPool.Get().([]byte)
    return buf[:4*1024]
case SessionTypeSCP, SessionTypeRSync:
    // Bulk transfer: large buffer for throughput
    return largeBufPool.Get().([]byte)
default:
    return stdBufPool.Get().([]byte)
}
}

var (
    smallBufPool = sync.Pool{New: func() interface{} {
        return make([]byte, 4*1024) }}
    stdBufPool   = sync.Pool{New: func() interface{} {
        return make([]byte, 32*1024) }}
    largeBufPool = sync.Pool{New: func() interface{} {
        return make([]byte, 256*1024) }}
)

```

7.3 TCP Keep-Alive and SO_KEEPALIVE Tuning

```

// helixterm.io/services/ssh-proxy/internal/net/tcp.go

package net

import (
    "net"
    "syscall"
    "time"
)

// ConfigureSSHConn applies TCP optimisations for SSH proxy
// connections.
func ConfigureSSHConn(conn *net.TCPConn) error {
    // Enable TCP keep-alive to detect dead connections
    if err := conn.SetKeepAlive(true); err != nil {
        return fmt.Errorf("SetKeepAlive: %w", err)
    }
    if err := conn.SetKeepAlivePeriod(30 * time.Second);
    err != nil {
        return fmt.Errorf("SetKeepAlivePeriod: %w", err)
    }
}

rawConn, err := conn.SyscallConn()
if err != nil {

```

```

        return err
    }

    return rawConn.Control(func(fd uintptr) {
        // TCP_KEEPIIDLE: wait 30s before sending first
        // keepalive
        syscall.SetsockoptInt(int(fd), syscall.IPPROTO_TCP,
            syscall.TCP_KEEPIIDLE, 30)

        // TCP_KEEPIIDVL: send keepalive every 10s after
        // first
        syscall.SetsockoptInt(int(fd), syscall.IPPROTO_TCP,
            syscall.TCP_KEEPIIDVL, 10)

        // TCP_KEEPCNT: give up after 5 failed keepalives
        // (50s total)
        syscall.SetsockoptInt(int(fd), syscall.IPPROTO_TCP,
            syscall.TCP_KEEPCNT, 5)

        // TCP_NODELAY: disable Nagle algorithm for
        // interactive sessions
        // Critical for keystroke latency – without this,
        // small packets are buffered
        syscall.SetsockoptInt(int(fd), syscall.IPPROTO_TCP,
            syscall.TCP_NODELAY, 1)

        // SO_RCVBUF / SO_SNDBUF: tuned for SSH channel sizes
        syscall.SetsockoptInt(int(fd), syscall.SOL_SOCKET,
            syscall.SO_RCVBUF, 256*1024)
        syscall.SetsockoptInt(int(fd), syscall.SOL_SOCKET,
            syscall.SO_SNDBUF, 256*1024)
    })
}

```

7.4 Connection Draining for Zero-Downtime Restarts

```

// helixterm.io/services/ssh-proxy/internal/server/drain.go

package server

import (
    "context"
    "net/http"
    "time"
    "digital.vasic.recovery"
)

type Drainer struct {
    sessionRegistry *SessionRegistry
}

```

```

    maxDrainTime    time.Duration
}

// GracefulShutdown implements connection draining for zero-
// downtime pod restarts.
// K8s sends SIGTERM → Drainer starts → waits for sessions to
// complete.
func (d *Drainer) GracefulShutdown(ctx context.Context)
    error {
    // Phase 1: Mark pod as draining (removed from load
    // balancer)
    // Kubernetes readiness probe returns 503 from this point
    d.sessionRegistry.SetDraining(true)

    // Phase 2: Stop accepting new SSH connections
    d.stopListener()

    // Phase 3: Wait for existing sessions to complete (with
    // deadline)
    deadline := time.Now().Add(d.maxDrainTime)
    ticker := time.NewTicker(5 * time.Second)
    defer ticker.Stop()

    for {
        active := d.sessionRegistry.ActiveCount()
        if active == 0 {
            slog.Info("All sessions drained")
            return nil
        }

        if time.Now().After(deadline) {

            slog.Warn("Drain timeout reached, forcing session
            close",
                "remaining_sessions", active)
            // Phase 4: Force-close remaining sessions after
            // deadline
            d.sessionRegistry.ForceCloseAll()
            return ErrDrainTimeout
        }

        select {
        case <-ctx.Done():
            return ctx.Err()
        case <-ticker.C:
            slog.Info("Draining sessions", "remaining",
            active)
        }
    }
}

```

7.5 Benchmarks: Concurrent SSH Sessions

Concurrent Sessions	CPU (8-core pod)	Memory	p50 Latency	p99 Latency	Throughput
10,000	15%	2.1 GB	0.3 ms	1.2 ms	50 MB/s
50,000	48%	8.4 GB	0.5 ms	2.8 ms	200 MB/s
100,000	78%	15.2 GB	0.9 ms	5.1 ms	380 MB/s
150,000	95%	22.1 GB	2.1 ms	12 ms	450 MB/s

Session mix: 80% interactive (low byte rate), 20% SCP transfer (high byte rate)

8. Kafka Performance Tuning

8.1 Producer Configurations

```
// helixterm.io/services/audit-service/internal/kafka/
// producer_config.go

package kafka

import (
    "time"
    "github.com/IBM/sarama"
)

// DurabilityFirstProducerConfig for audit events.
// Prioritises zero data loss over throughput.
func DurabilityFirstProducerConfig() *sarama.Config {
    cfg := sarama.NewConfig()

    // Acknowledgement: wait for all ISR replicas
    cfg.Producer.RequiredAcks = sarama.WaitForAll

    // Retry: aggressive retry for durability
    cfg.Producer.Retry.Max = 10
    cfg.Producer.Retry.Backoff = 200 * time.Millisecond

    // Idempotent producer: prevents duplicates under retry
    cfg.Producer.Idempotent = true
    cfg.Net.MaxOpenRequests = 1 // Required for idempotent
    // producer
```

```

// Batching: small batches for lower latency
cfg.Producer.Flush.Frequency = 5 * time.Millisecond
cfg.Producer.Flush.Bytes = 64 * 1024 // 64 KB
cfg.Producer.Flush.Messages = 1000

// Compression: snappy (good compression ratio, low CPU)
cfg.Producer.Compression = sarama.CompressionSnappy

// Buffer
cfg.Producer.Return.Successes = true
cfg.Producer.Return.Errors = true
cfg.ChannelBufferSize = 1024

return cfg
}

// ThroughputFirstProducerConfig for terminal stream events.
// Prioritises throughput over durability (terminal data is
// ephemeral).
func ThroughputFirstProducerConfig() *sarama.Config {
    cfg := sarama.NewConfig()

    // Only wait for leader ack (replica lag acceptable)
    cfg.Producer.RequiredAcks = sarama.WaitForLocal

    // Larger batches for throughput
    cfg.Producer.Flush.Frequency = 50 * time.Millisecond //
    // Tolerate 50ms latency
    cfg.Producer.Flush.Bytes = 1 * 1024 *
    // 1024 // 1 MB batches
    cfg.Producer.Flush.Messages = 10000

    // LZ4 compression: fastest decompression (important for
    // consumer side)
    cfg.Producer.Compression = sarama.CompressionLZ4

    // Relaxed retry
    cfg.Producer.Retry.Max = 3
    cfg.Producer.Retry.Backoff = 100 * time.Millisecond

    cfg.ChannelBufferSize = 4096
    return cfg
}

```

8.2 Topic Configuration

```

# Create topics with production-ready configuration

# audit.events: durability-first

```

```

kafka-topics.sh --create \
  --bootstrap-server kafka:9092 \
  --topic audit.events \
  --partitions 96 \           # 96 partitions = 96 max
                             concurrent consumers
  --replication-factor 3 \
  --config retention.ms=604800000 \ # 7 day retention
  --config min.insync.replicas=2 \ # At least 2 replicas
                             must be in sync
  --config cleanup.policy=delete \
  --config segment.bytes=268435456 \ # 256 MB segments
  --config compression.type=snappy \
  --config max.message.bytes=1048576 # 1 MB max message

# ssh.events: high-throughput
kafka-topics.sh --create \
  --bootstrap-server kafka:9092 \
  --topic ssh.events \
  --partitions 192 \         # 2x audit partitions for
                             higher throughput
  --replication-factor 3 \
  --config retention.ms=86400000 \ # 24 hour retention
                             (ephemeral)
  --config min.insync.replicas=2 \
  --config compression.type=lz4 \
  --config max.message.bytes=1048576

# terminal.recordings: bulk transfer
kafka-topics.sh --create \
  --bootstrap-server kafka:9092 \
  --topic terminal.recordings \
  --partitions 48 \
  --replication-factor 3 \
  --config retention.ms=2592000000 \ # 30 day retention
  --config min.insync.replicas=2 \
  --config max.message.bytes=10485760 \ # 10 MB max (chunked
  recordings)
  --config compression.type=lz4 \
  --config segment.bytes=536870912 # 512 MB segments for
  bulk data

```

8.3 Consumer Group Sizing

Topic	Partitions	Consumer Group	Consumers	Throughput Target
audit.events	96	audit-service	96 (1/ partition)	500k events/s

Topic	Partitions	Consumer Group	Consumers	Throughput Target
ssh.events	192	session-service	96 (2 partitions/consumer)	1M events/s
ssh.events	192	analytics-service	48	500k events/s
terminal.recordings	48	recording-service	24	50 GB/h
vault.sync.completed	24	notification-service	12	100k events/s

8.4 Consumer Lag Monitoring

```
# prometheus/rules/kafka-lag.yaml
groups:
- name: kafka-consumer-lag
  rules:
    - record: kafka:consumer_lag:sum_by_group
      expr: sum by (consumer_group, topic)
        (kafka_consumer_group_lag)

    - alert: KafkaAuditConsumerLagWarning
      expr: kafka_consumer_group_lag{consumer_group="audit-
        service", topic="audit.events"} > 5000
      for: 3m
      labels:
        severity: warning
        team: platform
      annotations:
        summary: "audit-service Kafka lag {{ $value }} on
          audit.events"
        runbook: "https://runbooks.helixterm.io/kafka/
          consumer-lag"

    - alert: KafkaSSHConsumerLagCritical
      expr: kafka_consumer_group_lag{topic="ssh.events"} >
        50000
      for: 5m
      labels:
        severity: critical
```

8.5 Kafka Streams for Session Analytics

```
// helixterm.io/services/metrics-aggregator/internal/streams/
// session_analytics.go

package streams

// SessionAnalytics uses Kafka Streams to compute real-time
// session metrics.
// Topology: ssh.events → windowed aggregation →
// session.metrics topic

func BuildSessionAnalyticsTopology(builder
    *kstreams.StreamsBuilder) {
    // Input: SSH connection events
    stream := builder.Stream("ssh.events",
        kstreams.Consumed.With(
            kstreams.Serdes.String(),
            kstreams.Serdes.Struct(SSHEvent{})),
    ),
)

    // 1-minute tumbling window aggregation
    windowedStream := stream.
        Filter(func(key string, event *SSHEvent) bool {
            return event.Type == "CONNECTED" || event.Type
            == "DISCONNECTED"
        }).
        GroupByKey().
        WindowedBy(kstreams.TimeWindows.OfSizeWithNoGrace(1
            * time.Minute))

    // Count active sessions per minute
    windowedStream.
        Count(kstreams.Materialized.As("session-counts")).
        ToStream().
        To("session.metrics")
}
```

9. API Gateway & Rate Limiting Performance

9.1 Rate Limiter Algorithm Comparison

Algorithm	Burst Handling	Memory	Distributed?	Smoothness	Best For
Token Bucket	Excellent	O(1)	Via Redis	Medium	API endpoints
Sliding Window	Good	O(window/bucket)	Via Redis	Good	User rate limits
Fixed Window	Poor (boundary spike)	O(1)	Via Redis	Poor	Avoid
Leaky Bucket	Poor (strict)	O(1)	Via Redis	Excellent	Upstream call shaping
GCRA (Token Bucket variant)	Excellent	O(1)	Via Redis	Excellent	Recommended

Recommendation: Use GCRA (Generic Cell Rate Algorithm) — it provides smooth rate limiting without burst artifacts, implements efficiently with a single Redis value, and handles distributed deployment without coordination overhead.

9.2 Distributed Rate Limiter Implementation

```
// helixterm.io/services/api-gateway/internal/ratelimit/
gcrago
// Uses digital.vasic.ratelimiter submodule

package ratelimit

import (
    "context"
    "fmt"
    "time"

    "github.com/redis/go-redis/v9"
    "digital.vasic.ratelimiter"
)

// GCRARateLimiter implements the Generic Cell Rate Algorithm
using Redis.
```

```

// Single atomic Lua script ensures correctness under
// concurrent access.
type GCRARateLimiter struct {
    redis    *redis.ClusterClient
    rate     float64           // requests per second
    burst    int              // max burst above rate
    emission time.Duration    // time between emissions = 1/rate
}

// Lua script for GCRA: atomic check-and-increment.
// Avoids race conditions without distributed locking.
var gcraScript = redis.NewScript(`
local key = KEYS[1]
local now = tonumber(ARGV[1])
local emission_interval = tonumber(ARGV[2])
local burst_offset = tonumber(ARGV[3])
local ttl = tonumber(ARGV[4])

local tat = redis.call('GET', key)
if tat == false then
    tat = now
else
    tat = tonumber(tat)
end

-- TAT = Theoretical Arrival Time
-- If TAT is in the past, reset to now (burst recovery)
local new_tat = math.max(tat, now) + emission_interval
local allow_at = new_tat - burst_offset

if allow_at <= now then
    -- Request allowed: update TAT
    redis.call('SET', key, new_tat, 'PX', ttl)
    local remaining = math.floor((burst_offset - (new_tat -
        now)) / emission_interval)
    return {1, remaining, math.ceil(allow_at - now)}
else
    -- Request denied: return retry-after
    return {0, 0, math.ceil(allow_at - now)}
end
`)

type AllowResult struct {
    Allowed    bool
    Remaining  int
    RetryAfter time.Duration
}

func (r *GCRARateLimiter) Allow(ctx context.Context, key
    string) (*AllowResult, error) {
    now := time.Now().UnixNano()

```

```

    emissionNs := r.emission.Nanoseconds()
    burstOffsetNs := int64(r.burst) * emissionNs
    ttlMs := (r.emission * time.Duration(r.burst+1) *
        2).Milliseconds()

    result, err := gcraScript.Run(ctx, r.redis,
        []string{fmt.Sprintf("ratelimit:%s", key)},
        now, emissionNs, burstOffsetNs, ttlMs,
    ).Slice()

    if err != nil {
        // On Redis failure: fail open (allow request)
        rateLimitRedisErrors.Inc()
        return &AllowResult{Allowed: true}, nil
    }

    allowed := result[0].(int64) == 1
    remaining := int(result[1].(int64))
    retryAfterNs := result[2].(int64)

    rateLimitDecisions.WithLabelValues(
        boolStr(allowed),
    ).Inc()

    return &AllowResult{
        Allowed:    allowed,
        Remaining:  remaining,
        RetryAfter: time.Duration(retryAfterNs) *
            time.Nanosecond,
    }, nil
}

```

9.3 Gin Routing Performance: Trie vs Regexp

```

// helixterm.io/services/api-gateway/internal/router/
// bench_test.go

package router

import (
    "net/http"
    "net/http/httptest"
    "testing"
    "github.com/gin-gonic/gin"
)

// BenchmarkGinRouterTrie tests httprouter (trie-based)
// routing - Gin's default
func BenchmarkGinRouterTrie(b *testing.B) {

```

```

gin.SetMode(gin.ReleaseMode)
r := gin.New()
registerAllRoutes(r) // Register 200 routes

req := httptest.NewRequest("GET", "/api/v1/hosts/host-id-123/sessions", nil)
w := httptest.NewRecorder()

b.ResetTimer()
b.RunParallel(func(pb *testing.PB) {
    for pb.Next() {
        r.ServeHTTP(w, req)
    }
})
}
// Result: 195 ns/op, 0 allocs/op ← Trie routing

// BenchmarkRegexpRouter tests regexp-based routing (DO NOT USE)
func BenchmarkRegexpRouter(b *testing.B) {
    // Shows why Gin's trie routing is 50× faster than regexp
    // regexp routing result: ~9,800 ns/op, 15 allocs/op
}

```

Gin best practices for HelixTerminator:

```

// helixterm.io/services/api-gateway/internal/router/setup.go

package router

import (
    "github.com/gin-gonic/gin"
)

func SetupRouter() *gin.Engine {
    gin.SetMode(gin.ReleaseMode)

    // Use gin.New() not gin.Default() – don't include
    // default Logger/Recovery
    // We use our own structured logger and recovery
    // middleware
    r := gin.New()

    // Middleware cost (measured per request, in order):
    // 1. RequestID: ~50 ns
    // 2. Recovery: ~20 ns (digital.vasic.recovery)
    // 3. OTel: ~150 ns
    // 4. RateLimit: ~500 µs (Redis round-trip)
    // 5. Auth JWT: ~300 µs (cache hit)
}

```

```

// 6. RBAC coarse: ~200 µs (cache hit)
r.Use(
    RequestIDMiddleware(),
    RecoveryMiddleware(),
    OtelMiddleware(),
    RateLimitMiddleware(),
    JWTMiddleware(),
    CoarseRBACMiddleware(),
)

// Group routes by version to avoid route conflict
// overhead
v1 := r.Group("/api/v1")
{
    v1.GET("/hosts", handler.ListHosts)
    v1.GET("/hosts/:id", handler.GetHost)
    v1.POST("/hosts", handler.CreateHost)
    // ... 200 more routes
}

return r
}

```

9.4 HTTP/2 Multiplexing Configuration

```

// helixterm.io/services/api-gateway/cmd/main.go

package main

import (
    "net/http"
    "golang.org/x/net/http2"
    "golang.org/x/net/http2/h2c"
)

func newHTTP2Server(handler http.Handler) *http.Server {
    h2s := &http2.Server{
        MaxConcurrentStreams: 1000, // per-
            connection stream limit
        MaxReadFrameSize: 1 << 20, // 1 MB
        PermitProhibitedCipherSuites: false,
        IdleTimeout: 120 * time.Second,
    }

    srv := &http.Server{
        Addr: ":8443",
        Handler: h2c.NewHandler(handler,
            h2s), // Also accept H2C (cleartext) for internal
        TLSConfig: &tls.Config{

```

```

        MinVersion:          tls.VersionTLS13,
        PreferServerCipherSuites: true,
        // Session tickets for 0-RTT resumption
        SessionTicketsDisabled: false,
    },
    // Connection-level timeouts
    ReadTimeout:      30 * time.Second,
    WriteTimeout:     30 * time.Second,
    IdleTimeout:      120 * time.Second,
    ReadHeaderTimeout: 5 * time.Second,
    MaxHeaderBytes:   1 << 20, // 1 MB
}

http2.ConfigureServer(srv, h2s)
return srv
}

```

10. Flutter Client Performance

10.1 Widget Rebuild Minimisation

```

// lib/features/ssh/presentation/widgets/terminal_widget.dart

/// TerminalWidget uses RepaintBoundary and const
/// constructors aggressively.
/// Only the terminal canvas is rebuilt on new output;
/// toolbar and sidebar
/// are isolated via RepaintBoundary.
class TerminalWidget extends StatelessWidget {
  const TerminalWidget({
    super.key,
    required this.sessionId,
  });

  final String sessionId;

  @override
  Widget build(BuildContext context) {
    return Column(
      children: [
        // Toolbar: const → never rebuilds
        const TerminalToolbar(),

        // Terminal area: RepaintBoundary isolates repaints
        Expanded(
          child: RepaintBoundary(
            child: BlocBuilder<TerminalBloc, TerminalState>(

```

10.2 Terminal Canvas Performance: CustomPainter vs RichText

```

    @override
    State<TerminalCanvas> createState() =>
        _TerminalCanvasState();
}

class _TerminalCanvasState extends State<TerminalCanvas> {
    // Pre-allocate TextPainter cache to avoid re-creation
    final _charPainters = <String, TextPainter>{};
    final _paint = Paint();

    @override
    Widget build(BuildContext context) {
        return CustomPaint(
            painter: _TerminalPainter(
                buffer: widget.buffer,
                cursorPosition: widget.cursorPosition,
                charPainters: _charPainters,
            ),
            // Size: fills available space
            size: Size.infinite,
        );
    }
}

class _TerminalPainter extends CustomPainter {
    const _TerminalPainter({
        required this.buffer,
        required this.cursorPosition,
        required this.charPainters,
    });

    final TerminalBuffer buffer;
    final CursorPosition cursorPosition;
    final Map<String, TextPainter> charPainters;

    static const _cellWidth = 8.0;
    static const _cellHeight = 16.0;
    static const _fontStyle = TextStyle(
        fontFamily: 'JetBrains Mono',
        fontSize: 14.0,
        height: 1.143, // 16/14
    );

    @override
    void paint(Canvas canvas, Size size) {
        // Dirty rectangle optimisation: only repaint changed
        // rows
        final dirtyRows = buffer.getDirtyRows();

        for (final rowIdx in dirtyRows) {

```

```

        final row = buffer.getRow(rowIdx);
        final y = rowIdx * _cellHeight;

        for (int col = 0; col < row.length; col++) {
            final cell = row[col];
            final x = col * _cellWidth;

            // Background fill (only if non-default)
            if (cell.bgColor != TerminalColor.defaultBg) {
                canvas.drawRect(
                    Rect.fromLTWH(x, y, _cellWidth, _cellHeight),
                    Paint()..color = cell.bgColor.toColor(),
                );
            }

            // Character rendering via cached TextPainter
            if (cell.char != ' ') {
                final painter = _getOrCreatePainter(cell);
                painter.paint(canvas, Offset(x, y));
            }
        }
    }

    // Draw cursor
    _drawCursor(canvas);
    buffer.clearDirtyFlags();
}

TextPainter _getOrCreatePainter(TerminalCell cell) {
    final cacheKey = '${cell.char}:${cell.fgColor.index}:${cell.attrs.index}';
    return charPainters.putIfAbsent(cacheKey, () {
        final painter = TextPainter(
            text: TextSpan(
                text: cell.char,
                style: _fontStyle.copyWith(
                    color: cell.fgColor.toColor(),
                    fontWeight: cell.attrs.bold ? FontWeight.bold :
FontWeight.normal,
                ),
            ),
            textDirection: TextDirection.ltr,
        );
        painter.layout();
        return painter;
    });
}

@override
bool shouldRepaint(_TerminalPainter oldDelegate) {
    return oldDelegate.buffer.version != buffer.version ||

```

```

        oldDelegate.cursorPosition != cursorPosition;
    }
}

```

10.3 BLoC Stream Optimisation

```

// lib/features/ssh/bloc/terminal_bloc.dart

import 'package:flutter_bloc/flutter_bloc.dart';
import 'package:rxdart/rxdart.dart';

class TerminalBloc extends Bloc<TerminalEvent,
    TerminalState> {
  TerminalBloc({required this.sshChannel}) :
    super(TerminalState.initial()) {
    on<TerminalOutputReceived>(_onOutputReceived,
      // CRITICAL: use sequential transformer to process
      // events in order
      transformer: sequential(),
    );

    on<TerminalInputSent>(_onInputSent);
    on<TerminalResized>(_onResized);

    _subscribeToChannel();
  }

  final SSHChannel sshChannel;
  StreamSubscription<Uint8List>? _channelSubscription;

  void _subscribeToChannel() {
    _channelSubscription = sshChannel.outputStream
      // Throttle at 60 fps (16.67 ms) for interactive
      // sessions
      // This prevents 10k+ events/second from overwhelming
      // the UI
      .throttleTime(
        const Duration(milliseconds: 16),
        trailing: true, // Always include the latest value
        leading: false,
      )
      // Buffer rapid output: accumulate bytes within 16ms
      // window
      .bufferTime(const Duration(milliseconds: 16))
      .where((batch) => batch.isNotEmpty)
      // Flatten batched Uint8Lists into one
      .map((batch) => Uint8List.fromList(batch.expand((b)
        => b).toList()))
  }
}

```

```

        // distinct: don't emit if identical to last (edge
        case)
        .distinct()
        .listen(
            (data) => add(TerminalOutputReceived(data)),
            onError: (err) => add(TerminalError(err)),
        );
    }

    Future<void> _onOutputReceived(
        TerminalOutputReceived event,
        Emitter<TerminalState> emit,
    ) async {
        // VT100 parsing happens in isolate to avoid blocking UI
        thread
        final parsed = await _parseVT100Isolate(event.data);
        emit(state.copyWith(
            outputBuffer: state.outputBuffer.apply(parsed),
        ));
    }

    @override
    Future<void> close() {
        _channelSubscription?.cancel();
        return super.close();
    }
}

```

10.4 Dart Isolates for SSH Crypto

```

// lib/features/ssh/infrastructure/crypto_isolate.dart

import 'dart:isolate';
import 'dart:typed_data';

/// CryptoIsolate offloads SSH key generation and signing to
/// a background
/// isolate to avoid blocking the main UI thread.
///
/// Benchmark (main thread vs isolate):
/// - Key generation (Ed25519) on main thread: 45 ms → UI
///   jank
/// - Key generation on isolate: 0 ms UI impact (runs
///   concurrently)
class CryptoIsolate {
    static Isolate? _isolate;
    static SendPort? _sendPort;

    static Future<void> spawn() async {

```

```

    final receivePort = ReceivePort();
    _isolate = await Isolate.spawn(_cryptoEntryPoint,
        receivePort.sendPort);
    _sendPort = await receivePort.first;
}

static Future<Uint8List> signData(Uint8List
    data, Uint8List privateKey) async {
    final responsePort = ReceivePort();
    _sendPort!.send({
        'op': 'sign',
        'data': data,
        'key': privateKey,
        'replyTo': responsePort.sendPort,
    });
    return await responsePort.first;
}

static Future<KeyPair> generateKeyPair() async {
    final responsePort = ReceivePort();
    _sendPort!.send({
        'op': 'generate_keypair',
        'replyTo': responsePort.sendPort,
    });
    return await responsePort.first;
}
}

/// Entry point for the crypto isolate (runs in separate Dart
    thread).
void _cryptoEntryPoint(SendPort mainSendPort) {
    final receivePort = ReceivePort();
    mainSendPort.send(receivePort.sendPort);

    receivePort.listen((message) async {
        final replyTo = message['replyTo'] as SendPort;
        switch (message['op'] as String) {
            case 'sign':
                final sig = Ed25519().sign(message['data'],
                    message['key']);
                replyTo.send(sig);
            case 'generate_keypair':
                final pair = Ed25519().generateKeyPair();
                replyTo.send(pair);
        }
    });
}

```

10.5 App Startup Optimisation

```
// lib/main.dart

/// Deferred imports reduce initial bundle size and TTI (Time
    to Interactive).
/// Features not needed on startup are loaded lazily.

import 'package:helixterm/features/auth/auth_feature.dart';
// Deferred imports for heavy features:
import 'package:helixterm/features/vault/vault_feature.dart'
    deferred as vault;
import 'package:helixterm/features/terminal/
    terminal_feature.dart' deferred as terminal;
import 'package:helixterm/features/settings/
    settings_feature.dart' deferred as settings;

Future<void> main() async {
  WidgetsFlutterBinding.ensureInitialized();

  // Phase 1: Critical initialisation only (target: < 200 ms)
  await Future.wait([
    _initSecureStorage(),
    _initHive(),           // Local DB for offline vault
    _initSentryLite(),    // Crash reporting (lightweight)
  ]);

  runApp(const HelixTermApp());

  // Phase 2: Background preload (after first frame)
  WidgetsBinding.instance.addPostFrameCallback((_) async {
    await Future.wait([
      vault.loadLibrary(),
      terminal.loadLibrary(),
      CryptoIsolate.spawn(), // Warm up crypto isolate
    ]);
  });
}
```

10.6 Flutter Web: CanvasKit vs HTML Renderer

Metric	HTML Renderer	CanvasKit Renderer
Bundle size	+0 KB	+2.7 MB (CanvasKit WASM)
	~1.5 s	~3.5 s (WASM parse)

Metric	HTML Renderer	CanvasKit Renderer
TTI (Time to Interactive)		
Terminal render at 80×24	~12 ms/frame (CSS limits)	~0.8 ms/frame
Text rendering quality	Browser-native	Pixel-perfect (custom rasteriser)
Font loading	Browser cache	Embedded in CanvasKit
Recommendation	Login, settings pages	Terminal view, vault UI

Selective renderer configuration:

```
<!-- web/index.html: force CanvasKit for performance-critical app -->
<script>
  window.flutterConfiguration = {
    renderer: "canvaskit",
    // Preload CanvasKit WASM in background while showing splash
    canvaskitBaseUrl: "https://cdn.helixterm.io/canvaskit/",
  };
</script>
```

11. Observability for Performance

11.1 Prometheus Metric Definitions

```
// helixterm.io/services/api-gateway/internal/metrics/
// metrics.go
// Uses digital.vasic.observability submodule

package metrics

import (
  "github.com/prometheus/client_golang/prometheus"
  "github.com/prometheus/client_golang/prometheus/promauto"
)

var (
  // REQUEST METRICS
```

```

HTTPRequestDuration = promauto.NewHistogramVec(
    prometheus.HistogramOpts{
        Namespace: "helixterm",
        Subsystem: "api_gateway",
        Name:      "http_request_duration_seconds",
        Help:
            "HTTP request latency by method, path, and status
            code",
        Buckets:
            []float64{.0005, .001, .0025, .005, .01, .025, .05, .1, .25, .5,
            1},
    },
    []string{"method", "path_pattern", "status_code"},
)

HTTPRequestTotal = promauto.NewCounterVec(
    prometheus.CounterOpts{
        Namespace: "helixterm",
        Subsystem: "api_gateway",
        Name:      "http_requests_total",
        Help:      "Total HTTP requests by method, path
        pattern, and status code",
    },
    []string{"method", "path_pattern", "status_code"},
)

HTTPRequestInFlight = promauto.NewGauge(
    prometheus.GaugeOpts{
        Namespace: "helixterm",
        Subsystem: "api_gateway",
        Name:      "http_requests_in_flight",
        Help:      "Number of HTTP requests currently
        being processed",
    },
)

// ----- SSH METRICS -----

SSHConnectionsActive = promauto.NewGaugeVec(
    prometheus.GaugeOpts{
        Namespace: "helixterm",
        Subsystem: "ssh_proxy",
        Name:      "connections_active",
        Help:      "Number of active SSH connections by
        proxy pod",
    },
    []string{"pod", "protocol_version"},
)

SSHConnectionEstablishDuration =
    promauto.NewHistogramVec(

```

```

    prometheus.HistogramOpts{
        Namespace: "helixterm",
        Subsystem: "ssh_proxy",
        Name:
            "connection_establish_duration_seconds",
        Help:      "Time to establish an SSH connection,
from client request to first prompt",
        Buckets:   []float64{.05, .1, .2, .3, .5, 1, 2,
5},
    },
    []string{"auth_method", "cert_cache_status"},
)

SSHDataBytesForwarded = promauto.NewCounterVec(
    prometheus.CounterOpts{
        Namespace: "helixterm",
        Subsystem: "ssh_proxy",
        Name:      "data_bytes_forwarded_total",
        Help:      "Total bytes forwarded through SSH
proxy by direction",
    },
    []string{"direction"}, // "inbound" or "outbound"
)

SSHConnectionErrors = promauto.NewCounterVec(
    prometheus.CounterOpts{
        Namespace: "helixterm",
        Subsystem: "ssh_proxy",
        Name:      "connection_errors_total",
        Help:      "Total SSH connection errors by error
type",
    },
    []string{"error_type"}, // "auth_failed",
    "host_unreachable", "timeout", "cert_expired"
)

// ----- VAULT METRICS -----

VaultSyncDuration = promauto.NewHistogramVec(
    prometheus.HistogramOpts{
        Namespace: "helixterm",
        Subsystem: "vault_service",
        Name:      "sync_duration_seconds",
        Help:      "Duration of vault sync operations",
        Buckets:   []float64{.005, .01, .025, .05, .1, .25, .5, 1, 2.5,
5},
    },
    []string{"item_count_bucket", "cache_status"},
)

VaultItemCRDTMergeDuration = promauto.NewHistogram(

```

```

    prometheus.HistogramOpts{
        Namespace: "helixterm",
        Subsystem: "vault_service",
        Name:      "crdt_merge_duration_seconds",
        Help:      "Duration of CRDT merge operations on
vault items",
        Buckets:
        []float64{.0001, .0005, .001, .005, .01, .05, .1},
    },
)

// ----- CACHE METRICS -----

CacheOperationsTotal = promauto.NewCounterVec(
    prometheus.CounterOpts{
        Namespace: "helixterm",
        Subsystem: "cache",
        Name:      "operations_total",
        Help:      "Total cache operations by tier,
operation type, and result",
    },
    []string{"tier", "operation", "result"}, // tier: l1|
l2|l3; result: hit|miss|error
)

CacheMemoryBytes = promauto.NewGaugeVec(
    prometheus.GaugeOpts{
        Namespace: "helixterm",
        Subsystem: "cache",
        Name:      "memory_bytes",
        Help:      "Memory used by cache tier",
    },
    []string{"tier"},
)

// ----- RATE LIMITING METRICS -----

RateLimitDecisionsTotal = promauto.NewCounterVec(
    prometheus.CounterOpts{
        Namespace: "helixterm",
        Subsystem: "rate_limiter",
        Name:      "decisions_total",
        Help:      "Total rate limit decisions by
endpoint and result",
    },
    []string{"endpoint", "decision"}, // decision:
allowed|denied
)

// ----- CIRCUIT BREAKER METRICS -----

```

```

    CircuitBreakerState = promauto.NewGaugeVec(
        prometheus.GaugeOpts{
            Namespace: "helixterm",
            Subsystem: "circuit_breaker",
            Name:      "state",
            Help:      "Current circuit breaker state
(0=closed, 1=half-open, 2=open)",
        },
        []string{"name"},
    )
)

```

11.2 OpenTelemetry Trace Sampling Strategy

```

// helixterm.io/services/api-gateway/internal/otel/
// sampling.go

package otel

import (
    "go.opentelemetry.io/otel/sdk/trace"
    "go.opentelemetry.io/otel/attribute"
)

// TailBasedSampler implements tail-based sampling:
// - 100% of error traces
// - 100% of traces exceeding p95 latency
// - 1% of successful fast traces (statistical baseline)
// - 10% of traces for important endpoints
type TailBasedSampler struct {
    p95Threshold time.Duration
}

func (s *TailBasedSampler) ShouldSample(p
    trace.SamplingParameters) trace.SamplingResult {
    // Always sample errors
    for _, event := range p.ParentContext.
        (*spanContext).events {
        if event.Name == "exception" {
            return trace.SamplingResult{Decision:
                trace.RecordAndSample}
        }
    }

    // Always sample slow requests
    elapsed := time.Since(p.ParentContext.
        (*spanContext).startTime)
    if elapsed > s.p95Threshold {
        return trace.SamplingResult{

```

```

        Decision: trace.RecordAndSample,
        Attributes: []attribute.KeyValue{
            attribute.Bool("sampled.slow_request", true),
        },
    }
}

// Sample important endpoints at 10%
if isImportantEndpoint(p.Name) && rand.Float64() < 0.10 {
    return trace.SamplingResult{Decision:
        trace.RecordAndSample}
}

// Baseline: 1% sample rate for healthy traffic
if rand.Float64() < 0.01 {
    return trace.SamplingResult{Decision:
        trace.RecordAndSample}
}

return trace.SamplingResult{Decision: trace.Drop}
}

func NewTracerProvider(endpoint string)
(*trace.TracerProvider, error) {
    exporter, err := otlptracehttp.New(context.Background(),
        otlptracehttp.WithEndpoint(endpoint),
        otlptracehttp.WithCompression(otlptracehttp.GzipCompression),
    )
    if err != nil {
        return nil, err
    }

    return trace.NewTracerProvider(
        trace.WithBatcher(exporter,
            trace.WithMaxExportBatchSize(512),
            trace.WithBatchTimeout(5*time.Second),
        ),
        trace.WithSampler(&TailBasedSampler{
            p95Threshold: 25 * time.Millisecond,
        }),
        trace.WithResource(resource.NewWithAttributes(
            semconv.SchemaURL,
            semconv.ServiceName("api-gateway"),
            semconv.ServiceVersion(version.String()),
        )),
    ), nil
}

```

11.3 Grafana Dashboard JSON (Key Panels)

The following panels form the core HelixTerminator Platform Overview dashboard. Import via Grafana UI → Dashboards → Import → Paste JSON.

```
{
  "title": "HelixTerminator Platform Overview",
  "uid": "helixterm-platform-overview",
  "panels": [
    {
      "id": 5,
      "title": "PostgreSQL – Slow Queries & Pool Utilisation",
      "type": "timeseries",
      "gridPos": {"h": 8, "w": 24, "x": 0, "y": 16},
      "targets": [
        {
          "expr":
            "sum(helixterm_db_pool_size{state='in_use'}) /
            sum(helixterm_db_pool_size) * 100",
          "legendFormat": "Pool Utilisation %"
        },
        {
          "expr":
            "sum(rate(helixterm_db_query_duration_seconds_count{quantile='0.99'}
            [5m]))",
          "legendFormat": "Slow Queries/s (p99 > 100ms)"
        }
      ]
    }
  ]
}
```

11.4 SLO Burn Rate Alerting – Multi-Window Strategy

Multi-window alerting detects SLO burns at both fast (1 h) and slow (6 h / 24 h) rates, eliminating both false positives and delayed pages.

```
# helm/charts/helixterm-alerts/templates/slo-alerts.yaml
# Strategy: Google SRE Workbook multi-window burn-rate alerts
# Error budget = 1 - SLO target (e.g. 99.9% SLO → 0.1% error
# budget per 30 days)

groups:
  - name: helixterm.slo.api_gateway
    rules:
      # — TICKET: Fast burn (1h window)
```

```

- alert: APIGatewayErrorBudgetFastBurn
  expr: |
    (
      sum(rate(helixterm_api_gateway_http_requests_total{status_code=~"5.."}
        [1h]))
      /
      sum(rate(helixterm_api_gateway_http_requests_total[1h]))
    ) > (14.4 * 0.001)
  for: 5m
  labels:
    severity: critical
    slo: api_gateway_availability
    window: 1h
  annotations:
    summary: "API Gateway burning error budget 14.4×
      fast (1h window)"
    description: |
      Error rate {{ $value | humanizePercentage }} over
      1h window.
      At this rate the 30-day error budget will be
      exhausted in ~2 hours.
      Runbook: https://runbooks.helixterm.io/api-
      gateway/fast-burn

# — PAGE: Slow burn (6h window)

- alert: APIGatewayErrorBudgetSlowBurn
  expr: |
    (
      sum(rate(helixterm_api_gateway_http_requests_total{status_code=~"5.."}
        [6h]))
      /
      sum(rate(helixterm_api_gateway_http_requests_total[6h]))
    ) > (6 * 0.001)
  for: 30m
  labels:
    severity: warning
    slo: api_gateway_availability
    window: 6h
  annotations:
    summary: "API Gateway burning error budget 6× rate
      (6h window)"
    description: |
      Sustained elevated error rate over 6h.
      Runbook: https://runbooks.helixterm.io/api-
      gateway/slow-burn

```

```

# — TICKET: 24h burn rate watch
- alert: APIGatewayErrorBudgetDailyBurn
  expr: |
    (
      sum(rate(helixterm_api_gateway_http_requests_total{status_code=~"5
[24h]))
      /
      sum(rate(helixterm_api_gateway_http_requests_total[24h]))
    ) > (3 * 0.001)
  for: 2h
  labels:
    severity: warning
    slo: api_gateway_availability
    window: 24h
  annotations:
    summary: "API Gateway error rate elevated over 24h
    - budget at risk"

- name: helixterm.slo.ssh_proxy
  rules:
    - alert: SSHProxyConnectionLatencySLOBreach
      expr: |
        histogram_quantile(0.99,
          sum(rate(helixterm_ssh_proxy_connection_establish_duration_seconds_b
            by (le)
          ) > 0.500
        for: 5m
      labels:
        severity: critical
        slo: ssh_connection_latency_p99
      annotations:
        summary: "SSH connection p99 latency {{ $value |
        humanizeDuration }} - SLO breach (>500ms)"

    - alert: SSHProxyAvailabilityFastBurn
      expr: |
        (
          sum(rate(helixterm_ssh_proxy_connections_total{result="failed"}
[1h]))
          /
          sum(rate(helixterm_ssh_proxy_connections_total[1h]))
        ) > (14.4 * 0.005)
      for: 5m
      labels:
        severity: critical
        slo: ssh_proxy_availability

```

```

- name: helixterm.slo.vault_sync
  rules:
    - alert: VaultSyncLatencySLOBreach
      expr: |
        histogram_quantile(0.95,
          sum(rate(helixterm_vault_sync_duration_seconds_bucket[5m]))
            by (le)
          ) > 2.0
      for: 10m
      labels:
        severity: warning
        slo: vault_sync_p95_latency
      annotations:
        summary: "Vault sync p95 latency {{ $value |
          humanizeDuration }} - SLO breach (>2s)"

- name: helixterm.slo.database
  rules:
    - alert: DatabaseQueryLatencySLOBreach
      expr: |
        histogram_quantile(0.99,
          sum(rate(helixterm_db_query_duration_seconds_bucket[5m]))
            by (le, service, query_type)
          ) > 0.100
      for: 5m
      labels:
        severity: warning
        slo: db_query_latency_p99
      annotations:
        summary: "DB query p99 >100ms on
          {{ $labels.service }}/{{ $labels.query_type }}"

    - alert: DatabaseConnectionPoolExhaustion
      expr: |
        (
          sum by (service)
            (helixterm_db_pool_size{state="in_use"})
          /
          sum by (service) (helixterm_db_pool_size)
        ) > 0.90
      for: 2m
      labels:
        severity: critical
      annotations:
        summary: "DB pool >90% exhausted on
          {{ $labels.service }}"

- name: helixterm.slo.kafka
  rules:

```

```

- alert: KafkaConsumerLagCritical
  expr: |
    sum by (consumer_group, topic)
    (kafka_consumer_group_lag) > 50000
  for: 5m
  labels:
    severity: critical
    slo: kafka_consumer_lag
  annotations:
    summary: "Kafka consumer lag {{ $value }} for
    {{ $labels.consumer_group }}/{{ $labels.topic }}"

- alert: KafkaConsumerLagGrowing
  expr: |
    rate(kafka_consumer_group_lag[10m]) > 500
  for: 5m
  labels:
    severity: warning
  annotations:
    summary: "Kafka consumer lag growing at
    {{ $value }} msg/s for {{ $labels.consumer_group }}"

```

11.5 Continuous Profiling with Pyroscope

```

// helixterm.io/services/internal/profiling/profiling.go
// Integrates Pyroscope (or Grafana Parca) for always-on Go
// profiling.
// Module: digital.vasic.observability

package profiling

import (
    "context"
    "os"
    "runtime"
    "time"

    "github.com/grafana/pyroscope-go"
    "go.uber.org/zap"
)

// Config holds profiling configuration.
type Config struct {
    ServerAddress string
    ApplicationName string
    // Tags are appended to every profile sample.
    Tags map[string]string
    // SampleRate is the number of CPU samples per second
    // (default 100).

```

```

    CPUSampleRate uint32
    // EnableMutex enables mutex contention profiling.
    EnableMutex bool
    // EnableBlock enables goroutine block profiling.
    EnableBlock bool
    // EnableGoroutine enables goroutine count profiling.
    EnableGoroutine bool
}

// Start initialises Pyroscope continuous profiling.
// Call once at service startup, before accepting traffic.
func Start(cfg Config, log *zap.Logger) (func(), error) {
    if cfg.CPUSampleRate == 0 {
        cfg.CPUSampleRate = 100
    }

    profileTypes := []pyroscope.ProfileType{
        pyroscope.ProfileCPU,
        pyroscope.ProfileAllocObjects,
        pyroscope.ProfileAllocSpace,
        pyroscope.ProfileInuseObjects,
        pyroscope.ProfileInuseSpace,
    }

    if cfg.EnableMutex {
        runtime.SetMutexProfileFraction(5)
        profileTypes = append(profileTypes,
            pyroscope.ProfileMutexCount,
            pyroscope.ProfileMutexDuration)
    }

    if cfg.EnableBlock {
        runtime.SetBlockProfileRate(5000)
        profileTypes = append(profileTypes,
            pyroscope.ProfileBlockCount,
            pyroscope.ProfileBlockDuration)
    }

    if cfg.EnableGoroutine {
        profileTypes = append(profileTypes,
            pyroscope.ProfileGoroutines)
    }

    tags := map[string]string{
        "pod":      os.Getenv("POD_NAME"),
        "node":     os.Getenv("NODE_NAME"),
        "environment": os.Getenv("ENVIRONMENT"),
        "version":  os.Getenv("APP_VERSION"),
    }

    for k, v := range cfg.Tags {
        tags[k] = v
    }
}

```

```

    profiler, err := pyroscope.Start(pyroscope.Config{
        ApplicationName: cfg.ApplicationName,
        ServerAddress:   cfg.ServerAddress,
        Logger:          nil, // use Pyroscope's built-in
        Tags:            tags,
        ProfileTypes:    profileTypes,
        SampleRate:      cfg.CPUSampleRate,
    })
    if err != nil {
        return nil, fmt.Errorf("pyroscope start: %w", err)
    }

    log.Info("continuous profiling started",
        zap.String("server", cfg.ServerAddress),
        zap.String("app", cfg.ApplicationName),
        zap.Uint32("cpu_sample_rate", cfg.CPUSampleRate),
    )

    stop := func() {
        ctx, cancel :=
            context.WithTimeout(context.Background(),
                5*time.Second)
        defer cancel()
        if err := profiler.Stop(); err != nil {
            log.Error("pyroscope stop error", zap.Error(err))
        }
        _ = ctx
    }
    return stop, nil
}

```

```

# helm/charts/helixterm-pyroscope/values.yaml
pyroscope:
  image:
    repository: grafana/pyroscope
    tag: "1.5.0"
  replicaCount: 2
  resources:
    requests:
      cpu: "500m"
      memory: "1Gi"
    limits:
      cpu: "2"
      memory: "4Gi"
  storage:
    backend: s3
    s3:
      bucket: helixterm-pyroscope-profiles
      region: us-east-1

```

```

retention: 30d
ingestion:
  # Accept pprof, jfr, and OpenTelemetry profiles
  protocols:
    - pprof
    - otlp
compactor:
  enabled: true
  blockRetention: 720h # 30 days

```

11.6 OTel Tail-Sampling Collector Configuration

```

// helixterm.io/services/internal/tracing/sampler.go

package tracing

import (
    "go.opentelemetry.io/otel/sdk/trace"
)

// AdaptiveSampler implements tail-based sampling logic at
// the collector.
// At the SDK level we use a parent-based sampler with 10%
// head sampling
// for normal traffic, 100% for errors and high-latency
// spans.
//
// For tail-based decisions, deploy the OpenTelemetry
// Collector with the
// tailsamplingprocessor (see collector config below).
func NewSampler(targetRPS float64) trace.Sampler {
    return trace.ParentBased(
        // Root span sampler: adaptive rate targeting ~500
        // traces/s
        newAdaptiveRateSampler(targetRPS),
        // If parent was sampled, always sample child
        trace.WithRemoteParentSampled(trace.AlwaysSample()),
        // If parent was not sampled, never sample child
        // (except errors)
        trace.WithRemoteParentNotSampled(newErrorOnlySampler()),
    )
}

// newAdaptiveRateSampler samples at a rate to hit targetRPS.
// At 100k req/s with target 500 traces/s → 0.5% sampling
// rate.
func newAdaptiveRateSampler(targetRPS float64) trace.Sampler {
    {

```

```

    // Use TraceIDRatioBased – deterministic, no coordination
    // overhead.
    // Ratio is adjusted at startup based on expected load.
    return trace.TraceIDRatioBased(targetRPS / 100_000)
}

type errorOnlySampler struct{}

func newErrorOnlySampler() trace.Sampler { return
    &errorOnlySampler{} }

func (s *errorOnlySampler) ShouldSample(p
    trace.SamplingParameters) trace.SamplingResult {
    // Sample 100% of error spans regardless of parent
    // decision
    for _, attr := range p.Attributes {
        if attr.Key == "error" && attr.Value.AsBool() {
            return trace.SamplingResult{Decision:
                trace.RecordAndSample}
        }
    }
    return trace.SamplingResult{Decision: trace.Drop}
}

func (s *errorOnlySampler) Description() string { return
    "ErrorOnlySampler" }

```

```

# k8s/otel-collector/tail-sampling-config.yaml
# OpenTelemetry Collector tail-sampling processor
# configuration.
# Deployed as a stateful DaemonSet so all spans for a trace
# land on the same collector.

processors:
  tail_sampling:
    decision_wait: 10s          # wait up to 10s for all
    spans before deciding
    num_traces: 50000          # max traces in memory
    expected_new_traces_per_sec: 5000
    policies:
      # Always sample errors
      - name: errors-policy
        type: status_code
        status_code: {status_codes: [ERROR]}

      # Always sample slow traces (p99 target 500ms → alert
      # on >1s)
      - name: slow-traces-policy
        type: latency
        latency: {threshold_ms: 1000}

```

```
# Always sample SSH connection establishment (critical
path)
- name: ssh-connect-policy
  type: string_attribute
  string_attribute:
    key: "helixterm.operation"
    values: ["ssh.connect", "vault.sync", "auth.mfa"]

# Probabilistic policy: 1% of remaining healthy traces
- name: probabilistic-policy
  type: probabilistic
  probabilistic: {sampling_percentage: 1}
```

12. Danger Zones – Security-Performance Tradeoffs

12.1 AES-256-GCM Encryption Overhead

AES-256-GCM is hardware-accelerated on modern x86-64 CPUs via AES-NI instructions. At 10,000 req/s with an average payload of 4 KB:

Metric	Value
Throughput (AES-NI, single core)	~4 GB/s
Latency per 4 KB block	~1 μ s
CPU cost at 10k req/s $\times$ 4 KB	~40 MB/s – negligible
CPU cost at 100k req/s $\times$ 64 KB	~6.4 GB/s – saturates 2 cores

```
// helixterm.io/services/vault/internal/crypto/
gcm_bench_test.go

package crypto_test

import (
    "crypto/aes"
    "crypto/cipher"
    "crypto/rand"
    "testing"
)

func BenchmarkAES256GCMEncrypt4KB(b *testing.B) {
```

```

    key := make([]byte, 32)
    rand.Read(key)
    plaintext := make([]byte, 4096)
    rand.Read(plaintext)
    nonce := make([]byte, 12)
    rand.Read(nonce)

    block, _ := aes.NewCipher(key)
    gcm, _ := cipher.NewGCM(block)
    dst := make([]byte, 0, len(plaintext)+gcm.Overhead())

    b.SetBytes(int64(len(plaintext)))
    b.ResetTimer()
    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            dst = gcm.Seal(dst[:0], nonce, plaintext, nil)
        }
    })
}

func BenchmarkAES256GCMDecrypt4KB(b *testing.B) {
    key := make([]byte, 32)
    rand.Read(key)
    plaintext := make([]byte, 4096)
    rand.Read(plaintext)
    nonce := make([]byte, 12)
    rand.Read(nonce)

    block, _ := aes.NewCipher(key)
    gcm, _ := cipher.NewGCM(block)
    ciphertext := gcm.Seal(nil, nonce, plaintext, nil)
    dst := make([]byte, 0, len(plaintext))

    b.SetBytes(int64(len(plaintext)))
    b.ResetTimer()
    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            dst, _ = gcm.Open(dst[:0], nonce, ciphertext,
                nil)
        }
    })
}

// Results on c5.2xlarge (Intel Xeon Platinum, AES-NI
// enabled):
// BenchmarkAES256GCMEncrypt4KB-8      1500000      780 ns/op
//          5247 MB/s    0 allocs/op
// BenchmarkAES256GCMDecrypt4KB-8      1500000      760 ns/op
//          5390 MB/s    0 allocs/op
//
// WITHOUT AES-NI (software fallback):

```

```
// BenchmarkAES256GCMEncrypt4KB-8      80000    14800 ns/op
//      276 MB/s
//
// ⚠ DANGER ZONE: Verify AES-NI is available on all node
//      types.
// Enforce via node affinity if needed:
// nodeAffinity:
//   requiredDuringSchedulingIgnoredDuringExecution:
//     nodeSelectorTerms:
//       - matchExpressions:
//         - key: feature.node.kubernetes.io/cpu-
//           cpuid.AESNI
//       operator: In
//       values: ["true"]
```

12.2 Argon2id KDF Cost — Brute-Force Protection vs Latency

Argon2id is the OWASP-recommended KDF for password hashing. The parameters are a direct tradeoff between user latency and brute-force resistance.

```
// helixterm.io/services/auth/internal/crypto/
//      argon2_bench_test.go

package crypto_test

import (
    "crypto/rand"
    "testing"

    "golang.org/x/crypto/argon2"
)

// BenchmarkArgon2idOWASP tests OWASP minimum (m=19456, t=2,
//      p=1)
func BenchmarkArgon2idOWASP(b *testing.B) {
    password := []byte("correct horse battery staple")
    salt := make([]byte, 16)
    rand.Read(salt)
    b.ResetTimer()
    for i := 0; i < b.N; i++ {
        _ = argon2.IDKey(password, salt, 2, 19456, 1, 32)
    }
}

// BenchmarkArgon2idHelixDefault tests HelixTerminator
//      default (m=65536, t=3, p=4)
func BenchmarkArgon2idHelixDefault(b *testing.B) {
    password := []byte("correct horse battery staple")
    salt := make([]byte, 16)
```

```

    rand.Read(salt)
    b.ResetTimer()
    for i := 0; i < b.N; i++ {
        _ = argon2.IDKey(password, salt, 3, 65536, 4, 32)
    }
}

// Results on c5.2xlarge:
// BenchmarkArgon2id0WASP-8          100    10.2 ms/op    (→ 98
//      hashes/s per core)
// BenchmarkArgon2idHelixDefault-8    30    34.5 ms/op    (→ 29
//      hashes/s per core)
//
// At 29 hashes/s per 4-core auth pod and 10k concurrent
//      login attempts:
// → Need 10000/29 ≈ 345 cores just for hashing.
//
// MITIGATION: Rate-limit login to max 50 req/s per pod →
//      50/29 ≈ 2 cores.
// Use a dedicated "hash worker" goroutine pool with
//      semaphore to prevent
// Argon2 from saturating all CPU cores.

```

```

// helixterm.io/services/auth/internal/crypto/argon2_pool.go

package crypto

import (
    "context"
    "crypto/rand"

    "golang.org/x/crypto/argon2"
)

// Argon2Params holds the KDF parameters.
type Argon2Params struct {
    Time      uint32 // iterations (t)
    Memory    uint32 // KB (m)
    Threads   uint8  // parallelism (p)
    KeyLen    uint32 // output key length
}

// DefaultArgon2Params are HelixTerminator production
//      defaults.
// Tuned to ~34ms on c5.2xlarge while providing strong brute-
//      force resistance.
var DefaultArgon2Params = Argon2Params{
    Time:      3,
    Memory:    65536,
}

```

```

    Threads: 4,
    KeyLen: 32,
}

// HashPool serialises Argon2 operations to a bounded worker
// pool,
// preventing CPU saturation from concurrent hash requests.
type HashPool struct {
    sem    chan struct{}
    params Argon2Params
}

// NewHashPool creates a pool allowing at most maxConcurrent
// simultaneous hashes.
// Recommended: maxConcurrent = runtime.NumCPU() / 2
func NewHashPool(maxConcurrent int, params Argon2Params)
    *HashPool {
    return &HashPool{
        sem:    make(chan struct{}, maxConcurrent),
        params: params,
    }
}

// Hash derives a key from password and salt using Argon2id.
// Blocks until a worker slot is available or ctx is
// cancelled.
func (p *HashPool) Hash(ctx context.Context, password, salt
    []byte) ([]byte, error) {
    select {
    case p.sem <- struct{}{}:
        defer func() { <-p.sem }()
    case <-ctx.Done():
        return nil, ctx.Err()
    }
    return argon2.IDKey(password, salt, p.params.Time,
        p.params.Memory, p.params.Threads, p.params.KeyLen),
        nil
}

// GenerateSalt creates a cryptographically random salt.
func GenerateSalt() ([]byte, error) {
    salt := make([]byte, 16)
    _, err := rand.Read(salt)
    return salt, err
}

```

Parameter Set	Time (ms)	Memory (MB)	Brute-force Rate (attacker GPU)	Recommended Use
OWASP minimum	~10 ms	19 MB	1 hash/hour per RTX 4090	Acceptable minimum
HelixTerminator default	~34 ms	64 MB	~0.3 hash/hour per RTX 4090	Production
High-security mode	~100 ms	256 MB	~0.09 hash/hour per RTX 4090	Admin account recovery

12.3 mTLS Handshake Overhead in Istio Service Mesh

```
// helixterm.io/services/internal/perf/mtls_bench_test.go

package perf_test

// mTLS ECDSA-P256 handshake benchmark results (measured with
// netperf + wireshark):
//
// ECDSA-P256 full TLS 1.3 handshake: ~0.8–1.2 ms (first
// connection)
// TLS 1.3 session resumption (0-RTT): ~0.05 ms
//
// At 10,000 new connections/s across all services:
// Total mTLS overhead = 10,000 × 1.0ms = 10 CPU-seconds/s
// = ~2.5 cores
//
// With connection reuse (keep-alive), most are resumed:
// 90% resumed × 0.05ms + 10% new × 1.0ms = ~0.14ms avg →
// ~1.4 CPU-seconds/s
//
// DANGER ZONE: Services that create a new HTTP client per
// request will incur
// full handshake cost every time. ALWAYS reuse http.Client /
// grpc.ClientConn.

// Mitigation: configure Istio DestinationRule for connection
// pooling
```

```
# k8s/istio/destination-rules/vault-service.yaml
apiVersion: networking.istio.io/v1beta1
```

```

kind: DestinationRule
metadata:
  name: vault-service-dr
  namespace: helixterm-prod
spec:
  host: vault-service.helixterm-prod.svc.cluster.local
  trafficPolicy:
    connectionPool:
      tcp:
        maxConnections: 1000
        connectTimeout: 30ms
        tcpKeepalive:
          time: 7200s
          interval: 75s
          probes: 9
    http:
      http1MaxPendingRequests: 200
      http2MaxRequests: 1000
      maxRequestsPerConnection: 0      # unlimited – reuse
      connection
      maxRetries: 3
      idleTimeout: 90s
      h2UpgradePolicy: UPGRADE        # force HTTP/2
      multiplexing
    loadBalancer:
      consistentHash:
        httpHeaderName: x-session-
        id      # sticky routing for session affinity
    outlierDetection:
      consecutiveGatewayErrors: 5
      interval: 10s
      baseEjectionTime: 30s
      maxEjectionPercent: 50

```

12.4 Audit Logging Write Amplification

```

// helixterm.io/services/audit/internal/writer/
//   async_writer.go
//
// DANGER ZONE: Synchronous audit writes block SSH session
//   establishment.
// Write amplification occurs when compliance mode requires:
// 1. Insert to audit_logs (PostgreSQL)
// 2. Publish to Kafka audit.events topic
// 3. Write to WORM S3 bucket (compliance archive)
// 4. Update SIEM via syslog
//
// Each write path adds 2–15ms. Total synchronous cost:
//   20–60ms.

```

```

// Solution: async buffered writer with guaranteed delivery.

package writer

import (
    "context"
    "encoding/json"
    "time"

    "github.com/IBM/sarama"
    "go.uber.org/zap"
)

// AuditEvent is the canonical audit log entry.
type AuditEvent struct {
    EventID      string    `json:"event_id"`
    Timestamp    time.Time `json:"timestamp"`
    UserID       string    `json:"user_id"`
    Action      string    `json:"action"`
    Resource     string    `json:"resource"`
    Result       string    `json:"result"`
    IPAddress    string    `json:"ip_address"`
    SessionID    string    `json:"session_id,omitempty"`
    Metadata     json.RawMessage `json:"metadata,omitempty"`
}

// AsyncAuditWriter buffers audit events and writes them
// asynchronously,
// decoupling the hot path from I/O latency.
type AsyncAuditWriter struct {
    producer sarama.AsyncProducer
    topic     string
    buffer    chan *AuditEvent
    errorsCh  <-chan *sarama.ProducerError
    log       *zap.Logger
    // drainTimeout is the maximum time to wait for buffer
    // drain on shutdown.
    drainTimeout time.Duration
}

func NewAsyncAuditWriter(producer sarama.AsyncProducer,
    topic string, bufferSize int, log *zap.Logger)
    *AsyncAuditWriter {
    w := &AsyncAuditWriter{
        producer:    producer,
        topic:       topic,
        buffer:      make(chan *AuditEvent, bufferSize),
        errorsCh:    producer.Errors(),
        log:         log,
        drainTimeout: 10 * time.Second,
    }
}

```

```

    go w.drainErrors()
    return w
}

// Write enqueues an audit event. This returns immediately (<
// 1µs) and never
// blocks the caller unless the buffer is full (> bufferSize
// events pending).
func (w *AsyncAuditWriter) Write(ctx context.Context, event
    *AuditEvent) error {
    data, err := json.Marshal(event)
    if err != nil {
        return fmt.Errorf("marshal audit event: %w", err)
    }

    msg := &sarama.ProducerMessage{
        Topic:      w.topic,
        Key:        sarama.StringEncoder(event.UserID),
        Value:      sarama.ByteEncoder(data),
        Timestamp:  event.Timestamp,
        Headers: []sarama.RecordHeader{
            {Key: []byte("event_type"), Value:
                []byte(event.Action)},
        },
    }

    select {
    case w.producer.Input() <- msg:
        return nil
    case <-ctx.Done():
        return ctx.Err()
    case <-time.After(5 * time.Millisecond):
        // Buffer full – this should never happen in normal
        // operation.
        // Increment a counter to alert on write pressure.
        auditWriteDropped.Inc()
        return ErrAuditBufferFull
    }
}

func (w *AsyncAuditWriter) drainErrors() {
    for err := range w.errorsCh {
        w.log.Error("audit write failed",
            zap.String("topic", err.Msg.Topic),
            zap.Error(err.Err),
        )
        auditWriteErrors.Inc()
    }
}

```

Write amplification risk table:

Compliance Mode	Write Destinations	Total Added Latency (sync)	Total Added Latency (async)
Standard	Kafka only	2 ms	< 0.01 ms
Enhanced	Kafka + PostgreSQL	8 ms	< 0.01 ms
Compliance (SOC 2 / ISO 27001)	Kafka + PostgreSQL + S3 WORM	45 ms	< 0.01 ms
Regulatory (FedRAMP)	All above + SIEM syslog	60 ms	< 0.01 ms

12.5 RBAC Policy Evaluation Cost (OPA)

```
// helixterm.io/services/authz/internal/opa/client.go
// digital.vasic.authz wraps OPA evaluation with caching and
// circuit breaking.

package opa

import (
    "context"
    "time"

    "github.com/open-policy-agent/opa/rego"
    lru "github.com/hashicorp/golang-lru/v2"
    "go.uber.org/zap"
)

// OPA policy evaluation costs:
//   Simple RBAC rule: ~0.05ms
//   Complex hierarchical RBAC: ~0.5ms
//   Full attribute-based policy (100 conditions): ~2ms
//
// At 100,000 req/s with uncached evaluation: 200 CPU-
// seconds/s = 50 cores.
// With 99% cache hit rate: 2 CPU-seconds/s = 0.5 cores. ✓
//
// Cache key: sha256(subject + action + resource +
// environment_hash)

type CachedOPAClient struct {
    query      rego.PreparedEvalQuery
```

```

    cache      *lru.Cache[string, *DecisionResult]
    cacheTTL   time.Duration
    log        *zap.Logger
    // Prometheus metrics
    cacheHits   prometheus.Counter
    cacheMisses prometheus.Counter
    evalDuration prometheus.Histogram
}

type DecisionResult struct {
    Allow      bool
    Reason     string
    expiresAt  time.Time
}

// Eval evaluates the OPA policy for a given input, with
// caching.
// Cache hit path: ~200ns. Cache miss path: ~500µs.
func (c *CachedOPAClient) Eval(ctx context.Context, input
    map[string]any) (*DecisionResult, error) {
    cacheKey := inputHash(input)

    if cached, ok := c.cache.Get(cacheKey); ok &&
        time.Now().Before(cached.expiresAt) {
        c.cacheHits.Inc()
        return cached, nil
    }

    c.cacheMisses.Inc()
    start := time.Now()

    results, err := c.query.Eval(ctx, rego.EvalInput(input))
    c.evalDuration.Observe(time.Since(start).Seconds())
    if err != nil {
        return nil, fmt.Errorf("opa eval: %w", err)
    }

    decision := &DecisionResult{
        Allow:      results.Allowed(),
        expiresAt:  time.Now().Add(c.cacheTTL),
    }
    c.cache.Add(cacheKey, decision)
    return decision, nil
}

```

OPA Benchmark Results (production policy complexity):

Scenario	Latency	CPU/1000req
Cache TTL		

Simple role check	0.05 ms	0.05 cores
60s		
Hierarchical RBAC	0.50 ms	0.50 cores
30s		
Attribute-based (full)	2.00 ms	2.00 cores
10s		
Cached (any policy, hit)	0.0002ms	0.0002 cores
—		

12.6 HSM Signing Latency for SSH Certificates

HSM Operation Latency Profiles (measured on AWS CloudHSM, Thales Luna Network HSM):

Operation (PKCS#11)	CloudHSM	Thales Luna	Software
RSA-2048 sign	12 ms	8 ms	0.5
ms			
ECDSA-P256 sign	3 ms	2 ms	0.15
ms			
ECDSA-P384 sign	4 ms	3 ms	0.25
ms			
AES-256 encrypt (1KB) (AES-NI)	0.5 ms	0.3 ms	0.001 ms
Key generation (P256)	20 ms	15 ms	2
ms			

Impact at 1,000 new SSH connections/s requiring certificate signing:

ECDSA-P256 via HSM: $1000 \times 3\text{ms} = 3 \text{ CPU-seconds/s} + \text{HSM queue depth}$.

HSM typically supports 1,000–10,000 sign ops/s depending on key type.

MITIGATION STRATEGY:

1. Cache signed SSH certificates for their validity period (default: 60s).

Cache key: `user_id + host_id + role_hash`

Cache hit rate: ~95% for recurring users → reduces HSM calls to 50/s.

2. Pre-sign a pool of short-lived certificates during off-peak hours.

3. Use separate HSM partitions per environment to prevent cross-env contention.

```
// helixterm.io/services/auth/internal/hsm/cert_cache.go

package hsm

import (
    "context"
    "crypto/rand"
    "time"

    lru "github.com/hashicorp/golang-lru/v2"
    "golang.org/x/crypto/ssh"
)

// CertCache caches short-lived SSH certificates to reduce
// HSM load.
type CertCache struct {
    cache *lru.Cache[string, *cachedCert]
    signer CertSigner // backed by HSM
    validity time.Duration
}

type cachedCert struct {
    cert *ssh.Certificate
    expiresAt time.Time
}

// CertSigner is implemented by the HSM client.
type CertSigner interface {
    SignCertificate(ctx context.Context, req *CertRequest)
        (*ssh.Certificate, error)
}

// GetOrSign returns a cached certificate or requests a new
// one from the HSM.
// Cache hit path: ~200ns. Cache miss path: HSM latency
// (~3ms).
func (c *CertCache) GetOrSign(ctx context.Context, req
    *CertRequest) (*ssh.Certificate, error) {
    key := req.cacheKey()

    if cached, ok := c.cache.Get(key); ok {
        // Return if at least 10s of validity remains (avoid
        // using a cert that
        // might expire during the SSH handshake).
        if time.Until(cached.expiresAt) > 10*time.Second {
            hsmCertCacheHits.Inc()
        }
    }
}
```

```

        return cached.cert, nil
    }
}

hsmCertCacheMisses.Inc()
cert, err := c.signer.SignCertificate(ctx, req)
if err != nil {
    return nil, err
}

c.cache.Add(key, &cachedCert{
    cert:      cert,
    expiresAt: time.Unix(int64(cert.ValidBefore), 0),
})
return cert, nil
}

```

12.7 Full-Text Search: pg_trgm vs Elasticsearch

```

-- PostgreSQL pg_trgm approach
-- vault_items has ~10M rows for a large enterprise customer.

-- Index creation
CREATE INDEX CONCURRENTLY vault_items_name_trgm_idx
ON vault_items
USING GIN (name gin_trgm_ops);

CREATE INDEX CONCURRENTLY vault_items_tags_trgm_idx
ON vault_items
USING GIN (tags gin_trgm_ops);

-- Query
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
SELECT id, name, type, tags
FROM vault_items
WHERE
    user_id = $1
    AND (
        name ILIKE '%' || $2 || '%'
        OR tags ILIKE '%' || $2 || '%'
    )
ORDER BY
    CASE WHEN name ILIKE $2 || '%' THEN 0 ELSE 1 END, --
    prefix match first
    similarity(name, $2) DESC
LIMIT 20;

/*
Results (10M rows, GIN trigram index):

```

```
Planning time: 2.1 ms
Execution time: 12.8 ms for 20 results
Index usage: YES (Bitmap Heap Scan via GIN)
Rows examined: ~2,400 (trigram filter is very selective)
*/
```

Dimension	pg_trgm	Elasticsearch
Latency (1M rows)	5–15 ms	1–5 ms
Latency (100M rows)	30–80 ms	2–8 ms
Operational complexity	Low (same DB)	High (separate cluster)
Infrastructure cost	Included in PostgreSQL	300–3,000/mo additional
Relevance ranking	Basic (trigram similarity)	Advanced (BM25, custom)
Faceted search	Difficult	Native
Real-time index updates	Immediate	~1s propagation delay
Break-even point	< 50M rows, < 1k searches/s	> 50M rows OR > 1k searches/s

Decision: Use `pg_trgm` for vaults up to 50M items; add Elasticsearch for enterprise accounts exceeding that threshold or requiring advanced relevance ranking.

13. Load Testing Results & Benchmarks

13.1 k6 Test Scenarios

```
// tests/k6/scenarios/api_gateway_load.js
// Run: k6 run --out prometheus=ns=helixterm tests/k6/
//       scenarios/api_gateway_load.js

import http from 'k6/http';
import ws from 'k6/ws';
import { check, sleep } from 'k6';
import { Rate, Trend, Counter } from 'k6/metrics';
```

```

// Custom metrics
const sshConnectLatency = new Trend('helix_ssh_connect_ms',
  true);
const vaultSyncLatency = new Trend('helix_vault_sync_ms',
  true);
const errorRate = new Rate('helix_error_rate');
const sshConnectionsOpen = new
  Counter('helix_ssh_connections_opened');

export const options = {
  scenarios: {
    // Scenario 1: API Gateway baseline load
    api_load: {
      executor: 'ramping-vus',
      startVUs: 0,
      stages: [
        { duration: '2m', target: 100 }, // warm-up
        { duration: '5m', target: 5000 }, // ramp to load
        { duration: '10m', target: 5000 }, // sustained
        load
        { duration: '2m', target: 10000 }, // spike
        { duration: '5m', target: 10000 }, // sustained
        spike
        { duration: '2m', target: 0 }, // ramp-down
      ],
      gracefulRampDown: '30s',
      exec: 'apiGatewayScenario',
    },

    // Scenario 2: SSH connection establishment
    ssh_connections: {
      executor: 'constant-arrival-rate',
      rate:
        500, // 500 new SSH
        connections/s
      timeUnit: '1s',
      duration: '10m',
      preAllocatedVUs: 200,
      maxVUs: 1000,
      exec: 'sshConnectionScenario',
    },

    // Scenario 3: Vault sync storm
    vault_sync: {
      executor: 'ramping-arrival-rate',
      startRate: 10,
      timeUnit: '1s',
      stages: [
        { duration: '5m', target: 1000 }, // ramp to 1000
        vault syncs/s
        { duration: '5m', target: 1000 }, // sustain
      ]
    }
  }
};

```

```

        { duration: '2m', target: 5000 },    // 5x spike
        (e.g. all users reconnect after outage)
        { duration: '5m', target: 5000 },    // sustain spike
        { duration: '2m', target: 0 },       // recovery
    ],
    preAllocatedVUs: 500,
    maxVUs: 2000,
    exec: 'vaultSyncScenario',
  },
},

thresholds: {
  // API Gateway SLOs
  'http_req_duration{scenario:api_load}':
    ['p(99)<500', 'p(95)<200'],
  'http_req_failed{scenario:api_load}':
    ['rate<0.001'], // < 0.1% errors

  // SSH connection SLOs
  'helix_ssh_connect_ms':
    ['p(99)<500', 'p(95)<300'],
  'http_req_failed{scenario:ssh_connections}':
    ['rate<0.005'], // < 0.5% failures

  // Vault sync SLOs
  'helix_vault_sync_ms':
    ['p(99)<2000', 'p(95)<1000'],
  'http_req_failed{scenario:vault_sync}':
    ['rate<0.01'], // < 1% under spike
},
};

export function apiGatewayScenario() {
  const BASE = 'https://api.helixterm.io';
  const token = __ENV.TEST_JWT_TOKEN;

  // Typical read-heavy API workload: 80% reads, 15% writes,
  // 5% deletes
  const r = Math.random();

  let res;
  if (r < 0.80) {
    // GET /api/v1/hosts – most common operation
    res = http.get(`${BASE}/api/v1/hosts?
      page=1&page_size=50`, {
      headers: { Authorization: `Bearer ${token}` },
      timeout: '5s',
    });
    check(res, {
      'hosts list 200': (r) => r.status === 200,
      'hosts list <200ms': (r) => r.timings.duration < 200,
    });
  }
}

```

```

    } else if (r < 0.95) {
      // POST /api/v1/hosts/:id/connect – initiate SSH
      connection
      const hostId = `host-${Math.floor(Math.random() * 10000)}
      `;
      const startMs = Date.now();
      res = http.post(`${BASE}/api/v1/hosts/${hostId}/
        connect`, JSON.stringify({
          username: 'ubuntu',
          auth_method: 'certificate',
        })), {
        headers: {
          'Content-Type': 'application/json',
          Authorization: `Bearer ${token}`,
        },
        timeout: '10s',
      });
      sshConnectLatency.add(Date.now() - startMs);
      sshConnectionsOpen.add(1);
      check(res, {
        'connect 200 or 201': (r) => r.status === 200 ||
          r.status === 201,
      });
    } else {
      // DELETE /api/v1/sessions/:id – disconnect session
      const sessionId = `session-${Math.floor(Math.random() *
        1000)}`;
      res = http.del(`${BASE}/api/v1/sessions/${sessionId}`,
        null, {
          headers: { Authorization: `Bearer ${token}` },
          timeout: '5s',
        });
      check(res, { 'delete 204': (r) => r.status === 204 });
    }

    errorRate.add(res.status >= 400);
    sleep(Math.random() * 0.5 + 0.1); // 100–600ms think time
  }

export function sshConnectionScenario() {
  const BASE = 'https://api.helixterm.io';
  const token = __ENV.TEST_JWT_TOKEN;

  const startMs = Date.now();
  const hostId = `loadtest-host-${Math.floor(Math.random()
    * 1000)}`;

  const res = http.post(`${BASE}/api/v1/hosts/${hostId}/
    connect`, JSON.stringify({
      username: 'ubuntu',

```

```

    auth_method: 'certificate',
    terminal: { rows: 24, cols: 80 },
  )), {
    headers: {
      'Content-Type': 'application/json',
      Authorization: `Bearer ${token}`,
    },
    timeout: '10s',
  });

sshConnectLatency.add(Date.now() - startMs);
errorRate.add(res.status >= 400);
check(res, {
  'ssh connect 201': (r) => r.status === 201,
  'ssh connect <500ms': (r) => r.timings.duration < 500,
});
}

export function vaultSyncScenario() {
  const BASE = 'https://api.helixterm.io';
  const token = __ENV.TEST_JWT_TOKEN;

  const startMs = Date.now();
  const res = http.post(`${BASE}/api/v1/vault/sync`,
    JSON.stringify({
      device_id: __ENV.DEVICE_ID || 'loadtest-
        device-001',
      last_sync_ts: new Date(Date.now() -
        60000).toISOString(),
      client_version: '2.0.0',
    )), {
    headers: {
      'Content-Type': 'application/json',
      Authorization: `Bearer ${token}`,
    },
    timeout: '30s',
  });

  vaultSyncLatency.add(Date.now() - startMs);
  errorRate.add(res.status >= 400);
  check(res, {
    'vault sync 200': (r) => r.status === 200,
    'vault sync <2s': (r) => r.timings.duration < 2000,
  });
}

```

13.2 Go Benchmark Table — Critical Functions

```
// helixterm.io/services/benchmarks/
critical_path_bench_test.go
```

```
package benchmarks_test
```

```
// Benchmark results collected on c5.4xlarge (16 vCPU, 32 GB
RAM), Go 1.25
```

```
// Command: go test -bench=. -benchmem -benchtime=10s
-count=5 ./...
```

```
/*
```

Function	allocs/op	B/op	MB/s	ns/op
BenchmarkJWTValidateRS256-16	12	1536	—	42000
BenchmarkJWTValidateEd25519-16	8	896	—	3100
BenchmarkJWTValidateEd25519_Cached-16	2	128	—	280
BenchmarkRateLimiterTokenBucket-16	0	0	—	320
BenchmarkRateLimiterSlidingWindowRedis-16	14	832	—	18000
BenchmarkAES256GCMEncrypt4KB-16	0	0	5247	780
BenchmarkAES256GCMDecrypt4KB-16	0	0	5390	760
BenchmarkSSHHandshakeECDSA-16	240	32768	—	820000
BenchmarkSSHHandshakeEd25519-16	180	24576	—	180000
BenchmarkTerminalDataForward1KB-16	0	0	2380	420
BenchmarkTerminalDataForward64KB-16	0	0	8000	8200
BenchmarkCRDTMerge1000ps-16	42	4096	—	15000
BenchmarkCRDTMerge100000ps-16	3800	327680	—	1200000
BenchmarkRedisGet_Local-16	4	256	—	210
BenchmarkRedisGet_Remote-16	6	320	—	18000
BenchmarkPgQueryVaultItems_WithIndex-16	28	3584	—	85000
BenchmarkPgQueryVaultItems_NoIndex-16	28	3584	—	9500000

BenchmarkKafkaProducerAsync_Audit-16	12	1024	—	4200
BenchmarkKafkaProducerSync_Terminal-16	14	1280	—	12000
BenchmarkOPAPolicyEval_Simple-16	38	6144	—	52000
BenchmarkOPAPolicyEval_Cached-16	2	128	—	190
BenchmarkGinRouter_1000Routes-16	4	384	—	1800
BenchmarkTLSHandshake_mTLS-16	320	49152	—	950000
BenchmarkTLSResumption_SessionTicket-16	80	8192	—	48000
BenchmarkArgon2id_OWASPMIN-16	10200000	10200	19922944	—
BenchmarkArgon2id_HelixDefault-16	34500000	34500	67108864	—
BenchmarkSingleflightCacheGet-16	1	128	—	310

Key observations:

- Ed25519 JWT validation is 13.5× faster than RS256; always use Ed25519 for new tokens.
- OPA policy evaluation with caching is 273× faster than uncached; always cache decisions.
- AES-256-GCM at ~5 GB/s makes encryption cost negligible for payloads < 1 MB.
- pg_trgm without index is 111× slower – GIN index is mandatory for vault search.
- CRDT merge scales $O(n \log n)$; for 10,000 ops it takes 1.2ms – acceptable for sync.

*/

13.3 Flamegraph Analysis Methodology

```
#!/usr/bin/env bash
# scripts/perf/capture_flamegraph.sh
# Capture a CPU flamegraph for a running service pod.
# Usage: ./capture_flamegraph.sh <pod_name> <namespace>
#         <duration_seconds>

set -euo pipefail

POD_NAME="${1:?pod name required}"
NAMESPACE="${2:-helixterm-prod}"
DURATION="${3:-30}"
OUTPUT_DIR="./flamegraphs/${date +%Y%m%d_%H%M%S}"
mkdir -p "$OUTPUT_DIR"
```

```

echo "==> Capturing pprof CPU profile from $POD_NAME for $
      {DURATION}s..."

# Port-forward to the pprof endpoint (exposed on port 6060,
      not in production ingress)
kubectl port-forward -n "$NAMESPACE" "pod/$POD_NAME"
      6060:6060 &
PF_PID=$!
sleep 2

# Capture CPU profile
go tool pprof -proto -seconds "$DURATION" \
      "http://localhost:6060/debug/pprof/profile?seconds=$
      {DURATION}" \
      > "$OUTPUT_DIR/cpu.pb.gz"

# Also capture heap, goroutine, and mutex profiles
go tool pprof -proto "http://localhost:6060/debug/pprof/
      heap" > "$OUTPUT_DIR/heap.pb.gz"
go tool pprof -proto "http://localhost:6060/debug/pprof/
      goroutine"> "$OUTPUT_DIR/goroutine.pb.gz"
go tool pprof -proto "http://localhost:6060/debug/pprof/
      mutex" > "$OUTPUT_DIR/mutex.pb.gz"

kill "$PF_PID" 2>/dev/null || true

# Generate SVG flamegraphs
echo "==> Generating SVG flamegraphs..."
go tool pprof -svg "$OUTPUT_DIR/cpu.pb.gz" >
      "$OUTPUT_DIR/cpu_flame.svg"
go tool pprof -svg "$OUTPUT_DIR/heap.pb.gz" >
      "$OUTPUT_DIR/heap_flame.svg"
go tool pprof -svg "$OUTPUT_DIR/goroutine.pb.gz">
      "$OUTPUT_DIR/goroutine_flame.svg"

# Upload to Pyroscope for long-term storage
echo "==> Uploading to Pyroscope..."
pyroscope adhoc \
      --url "http://pyroscope.helixterm-
      monitoring.svc.cluster.local:4040" \
      --app-name "helixterm.${POD_NAME%%-*}.adhoc" \
      "$OUTPUT_DIR/cpu.pb.gz"

echo "==> Done. Flamegraphs saved to $OUTPUT_DIR"
echo "      CPU:      $OUTPUT_DIR/cpu_flame.svg"
echo "      Heap:     $OUTPUT_DIR/heap_flame.svg"
echo "      Goroutine:$OUTPUT_DIR/goroutine_flame.svg"

```

Flamegraph interpretation guide for HelixTerminator:

Pattern to Look For	Root Cause	Action
Wide <code>runtime.mallocgc</code> bar	Excessive allocations	Pool objects; reduce slice growth
Wide <code>runtime.gcBgMarkWorker</code>	GC pressure (high allocation rate)	Reduce alloc rate, tune <code>GOGC</code>
Deep <code>sync.(*Mutex).Lock</code> stack	Lock contention	Use <code>sync.RWMutex</code> , sharding, or lock-free structures
Wide <code>tls.(*Conn).Handshake</code>	Too many new TLS connections	Enable connection reuse / keep-alive
Wide <code>json.Marshal</code> / <code>json.Unmarshal</code>	Slow serialization	Switch to <code>encoding/json/v2</code> or <code>jsoniter</code>
Wide <code>crypto/sha256</code>	Hash-heavy path	Pre-compute hashes; cache digests
Wide <code>net/http.(*Transport).roundTrip</code>	External HTTP bottleneck	Check downstream service latency

13.4 Capacity Planning Formula

Capacity planning model for a single microservice:

Given:

P = peak requests per second (target)
 L = average processing latency per request (seconds)
 C = number of concurrent goroutines per pod (= $P \times L$, Little's Law)
 G = max goroutines per pod (default: 10,000 for Go)
 M = memory per goroutine stack (2KB initial, up to 8KB typical)
 CPU = average CPU fraction per request

Pods required (CPU bound):

$pods_cpu = \text{ceil}((P \times CPU) / (0.7 \times vCPU_per_pod))$
 0.7 = 70% target CPU utilisation (30% headroom for spikes)

Pods required (memory bound):

$pods_mem = \text{ceil}((C \times M) / (0.8 \times RAM_per_pod))$

```

Pods required (goroutine bound):
  pods_goroutine = ceil(C / G)

Take the maximum of the three.

HPA formula:
  min_replicas = ceil(max(pods_cpu, pods_mem, pods_goroutine)
× 1.0) # base
  max_replicas = ceil(max(pods_cpu, pods_mem, pods_goroutine)
× 3.0) # 3× safety factor

Example – API Gateway at 100,000 req/s peak:
  L   = 5ms (0.005s)
  CPU = 0.001 cores/req (1ms CPU per req on a 1GHz-
equivalent thread)
  C   = 100,000 × 0.005 = 500 concurrent goroutines/pod at 1
pod
  M   = 8KB

  pods_cpu      = ceil((100,000 × 0.001) / (0.7 × 4)) =
ceil(100/2.8) = 36
  pods_mem      = ceil((500×36 × 8192) / (0.8 × 4GB)) =
negligible
  pods_goroutine = ceil(500×36 / 10000) = 2 (irrelevant)

→ min_replicas = 36, max_replicas = 108

```

13.5 Vertical vs. Horizontal Scale Break-Even

Service	Vertical Scale Gains	Horizontal Scale Gains	Break-Even
API Gateway	Diminishing above 4 vCPU (Gin scales linearly)	Linear – add pods	Horizontal always
SSH Proxy	NUMA benefits up to 16 vCPU	Linear to 100k sessions/pod	Horizontal at >8 vCPU
Vault Service	PostgreSQL connection limits improve	Limited by DB connections	Vertical first to 8 vCPU, then horizontal
PostgreSQL primary	Up to 32 vCPU, 256 GB RAM	Read replicas for reads	Vertical to 32 vCPU; read replicas for reads
Redis			

Service	Vertical Scale Gains	Horizontal Scale Gains	Break-Even
	Up to 16 vCPU cluster node	Redis Cluster for hot keys	Horizontal (cluster) for > 100k ops/s
Kafka broker	Disk I/O bound: NVMe SSD needed	Add brokers for partition leadership	Horizontal after NVMe

14. Performance Improvement Roadmap

14.1 Quick Wins (< 1 Week Each)

Each item is achievable in ≤ 5 engineering-days and has a concrete, measurable expected gain.

#	Improvement	Service	Expected Gain	Effort	Priority
1	Switch JWT signing from RS256 to Ed25519	auth, api-gateway	-93% JWT validation CPU	1 day	P0
2	Enable <code>GOGC=200</code> on API Gateway (reduce GC frequency at cost of 2x heap)	api-gateway	-30% GC pause time	0.5 day	P0
3	Add <code>singleflight</code> to cache stampede paths in vault-service	vault-service	-80% DB storm on cold start	1 day	P0
4	Set <code>http.Client</code> transport <code>MaxIdleConnsPerHost=100</code> on all outbound callers	all services	-60% mTLS handshake overhead	1 day	P0
5	Add GIN trgm index on <code>vault_items.name</code> (missing, causes sequential scans)	vault DB	-95% vault search latency	0.5 day	P0
6	Enable <code>tcp_nodelay = on</code> in PgBouncer (reduces Nagle algorithm buffering)	pgbouncer	-2ms median DB query latency	0.5 day	P1
7		PostgreSQL			P1

#	Improvement	Service	Expected Gain	Effort	Priority
	Tune <code>autovacuum_vacuum_cost_delay=2ms</code> on <code>audit_logs</code> (default 20ms)		-70% vacuum lag on hot table	0.5 day	
8	Add <code>idle_timeout: 90s</code> to all Istio DestinationRules (currently unlimited)	all services	-20% Envoy memory usage	1 day	P1
9	Enable <code>batch.size=1000000</code> , <code>linger.ms=5</code> on Kafka audit producer	audit-service	+40× Kafka producer throughput	0.5 day	P1
10	Add <code>compression.type=lz4</code> to terminal Kafka topics	terminal-service	-60% Kafka broker disk I/O	0.5 day	P1
11	Cache OPA policy decisions in-process (TTL=30s) — currently uncached	authz-service	-99% OPA CPU overhead	2 days	P1
12	Enable Redis <code>OBJECT ENCODING</code> checking — migrate large string keys to hash	session-service	-40% Redis memory	1 day	P1
13	Add <code>RepaintBoundary</code> around Flutter terminal canvas	Flutter client	-80% unnecessary widget repaints	1 day	P1
14	Use <code>const</code> constructors for all static Flutter widgets	Flutter client	-30% Flutter framework CPU	1 day	P2
15	Add <code>LIMIT 1</code> to existence-check queries (currently <code>SELECT *</code>)	vault, host, auth	-50% DB I/O on common checks	1 day	P2
16	Enable <code>prepared_statements=true</code> in PgBouncer <code>pgbouncer.ini</code>	pgbouncer	-15% DB query overhead	0.5 day	P2
17	Add read replica routing for <code>GET /api/v1/hosts</code> and <code>GET /api/v1/vault</code>	api-gateway	-40% primary DB write pressure	2 days	P2
18				1 day	P2

#	Improvement	Service	Expected Gain	Effort	Priority
	Enable HTTP/2 server push for Flutter Web initial asset load	api-gateway	-200ms Flutter Web startup		
19	Tune <code>SO_SNDBUF</code> and <code>SO_RCVBUF</code> to 256KB on SSH proxy TCP sockets	ssh-proxy	+25% terminal throughput	0.5 day	P2
20	Add <code>runtime/pprof</code> CPU profile endpoint gated behind internal IP	all services	Enables continuous profiling	1 day	P2

14.2 Medium-Term Improvements (1–4 Weeks Each)

#	Improvement	Architecture Change Required	Expected Gain	Duration
1	Implement circuit breakers for all service-to-service calls using <code>digital.vasic.recovery</code>	Add breaker middleware to all Gin routes	Prevents cascade failures; recovery in <30s	2 weeks
2	Connection pool per-service with PgBouncer transaction mode	Deploy PgBouncer sidecar per service namespace	-70% DB connection overhead	1 week
3	Partition <code>audit_logs</code> by month with <code>pg_partman</code>	Schema migration + <code>partman</code> install	-90% VACUUM time on hot table	2 weeks
4	Add Redis Cluster (6-node, 3 shards) to replace single-node Redis	Infrastructure change + client library update	3× Redis throughput; eliminates hot key risk	3 weeks
5	Implement tail-based Kafka consumer with batch DB writes (1,000 row batches)	Refactor audit consumer to batch mode	+100× audit write throughput	2 weeks
6	Add Kafka Streams session analytics (real-time session duration, bytes transferred)	New Kafka Streams application	Enables real-time billing and	3 weeks

#	Improvement	Architecture Change Required	Expected Gain	Duration
			anomaly detection	
7	Implement bulkhead pattern for SSH proxy goroutine pools (separate pools per customer tier)	Refactor goroutine pool in ssh-proxy	Prevents noisy-neighbor SSH storms	2 weeks
8	Add L1 in-process cache (ristretto) to vault-service for hot vault items	Integrate <code>digital.vasic.cache</code> L1 tier	-80% L2 Redis traffic for popular vaults	2 weeks
9	Implement gRPC streaming for terminal data instead of WebSocket	Replace WS handler with gRPC bidirectional stream	+30% terminal throughput; native flow control	4 weeks
10	Deploy Pyroscope continuous profiling for all 25 services	Helm chart + service annotation rollout	Baseline profiling data; enables regression detection	1 week

14.3 Long-Term Improvements (1+ Months)

14.3.1 WASM for Flutter Web Terminal

The Flutter Web terminal currently uses the HTML renderer which goes through the DOM for text rendering. Switching to CanvasKit (WASM) provides a 3–5× rendering improvement for the terminal canvas.

Beyond CanvasKit, implementing the SSH crypto stack directly in WebAssembly (compiled from Rust or C) eliminates the Dart crypto overhead on the main isolate:

Current Flutter Web terminal pipeline:

SSH bytes → Dart decode (main isolate) → Dart decrypt (main isolate) → Canvas HTML

Target Flutter Web terminal pipeline:

SSH bytes → WASM crypto worker (off main thread) →

CanvasKit painter (GPU)

Expected improvement:

- Terminal frame rate: 30 fps → 60 fps on mid-range devices
- Crypto overhead: 2ms → 0.1ms (20× improvement)
- Main thread CPU: -70% (crypto moved to WASM worker)

```
// Target architecture: Dart isolate offloads crypto to WASM
// worker
// lib/features/ssh/presentation/wasm/crypto_bridge.dart

import 'dart:isolate';
import 'dart:typed_data';

/// WasmCryptoBridge offloads SSH crypto to a dedicated WASM
/// worker,
/// keeping the main isolate free for UI rendering.
class WasmCryptoBridge {
  late final SendPort _workerPort;
  final _pending = <int, Completer<Uint8List>>{};
  int _nextId = 0;

  Future<void> init() async {
    final receivePort = ReceivePort();
    await Isolate.spawnUri(
      Uri.parse('assets/wasm/ssh_crypto.wasm.dart'),
      [receivePort.sendPort],
      null,
    );
    _workerPort = await receivePort.first as SendPort;
    receivePort.listen(_handleResponse);
  }

  Future<Uint8List> decrypt(Uint8List ciphertext, Uint8List
    key, Uint8List nonce) {
    final id = _nextId++;
    final completer = Completer<Uint8List>();
    _pending[id] = completer;
    _workerPort.send({'id': id, 'op': 'decrypt', 'data':
      ciphertext, 'key': key, 'nonce': nonce});
    return completer.future;
  }

  void _handleResponse(dynamic msg) {
    final id = (msg as Map)['id'] as int;
    _pending.remove(id)?.complete(msg['result'] as
      Uint8List);
  }
}
```

```
}  
}
```

14.3.2 eBPF for Network Observability

eBPF programs attached to the SSH proxy provide sub-microsecond network metrics without any changes to the application code:

```
// bpf/ssh_latency_tracker.bpf.c  
// Tracks TCP connection latency and retransmits for SSH  
// proxy pods.  
// Compiled with: clang -O2 -target bpf -c  
//                ssh_latency_tracker.bpf.c  
  
#include <linux/bpf.h>  
#include <linux/if_ether.h>  
#include <linux/ip.h>  
#include <linux/tcp.h>  
#include <bpf/bpf_helpers.h>  
#include <bpf/bpf_endian.h>  
  
struct conn_key {  
    __u32 src_ip;  
    __u32 dst_ip;  
    __u16 src_port;  
    __u16 dst_port;  
};  
  
struct conn_metrics {  
    __u64 connect_ts_ns;           // timestamp of SYN  
    __u64 established_ts_ns;      // timestamp of first data  
    packet  
    __u32 retransmits;  
    __u32 bytes_sent;  
    __u32 bytes_recv;  
};  
  
// Map: connection key → metrics  
struct {  
    __uint(type, BPF_MAP_TYPE_HASH);  
    __type(key, struct conn_key);  
    __type(value, struct conn_metrics);  
    __uint(max_entries, 1000000);  
} conn_metrics_map SEC(".maps");  
  
// Histogram map: connection latency in microseconds  
// (0-10,000µs, 100µs buckets)  
struct {  
    __uint(type, BPF_MAP_TYPE_ARRAY);
```

```

    __type(key, __u32);
    __type(value, __u64);
    __uint(max_entries, 100);
} latency_histogram SEC(".maps");

SEC("tp/tcp/tcp_probe")
int trace_tcp_probe(struct trace_event_raw_tcp_probe *ctx) {
    // Record per-connection metrics on every TCP event
    struct conn_key key = {
        .src_ip    = ctx->saddr,
        .dst_ip    = ctx->daddr,
        .src_port  = ctx->sport,
        .dst_port  = ctx->dport,
    };

    struct conn_metrics *m =
        bpf_map_lookup_elem(&conn_metrics_map, &key);
    if (!m) {
        struct conn_metrics new_m = {};
        new_m.connect_ts_ns = bpf_ktime_get_ns();
        bpf_map_update_elem(&conn_metrics_map, &key, &new_m,
            BPF_ANY);
        return 0;
    }

    if (ctx->data_len > 0 && m->established_ts_ns == 0) {
        m->established_ts_ns = bpf_ktime_get_ns();
        __u64 latency_us = (m->established_ts_ns -
            m->connect_ts_ns) / 1000;
        __u32 bucket = latency_us / 100;
        if (bucket >= 100) bucket = 99;
        __u64 *count =
            bpf_map_lookup_elem(&latency_histogram, &bucket);
        if (count) __sync_fetch_and_add(count, 1);
    }

    m->bytes_sent += ctx->data_len;
    return 0;
}

char LICENSE[] SEC("license") = "GPL";

```

```

# k8s/daemonset/ebpf-observer.yaml
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: helixterm-ebpf-observer
  namespace: helixterm-monitoring
spec:

```

```
selector:
  matchLabels:
    app: helixterm-ebpf-observer
template:
  metadata:
    labels:
      app: helixterm-ebpf-observer
  spec:
    hostPID: true
    hostNetwork: true
    tolerations:
      - operator: Exists      # run on all nodes including
        control plane
    containers:
      - name: ebpf-observer
        image: helixterm/ebpf-observer:1.0.0
        securityContext:
          privileged: true
          capabilities:
            add: [CAP_BPF, CAP_NET_ADMIN, CAP_PERFMON]
        volumeMounts:
          - name: bpf
            mountPath: /sys/fs/bpf
          - name: proc
            mountPath: /host/proc
            readOnly: true
        env:
          - name: PROMETHEUS_PORT
            value: "9090"
          - name: FILTER_POD_LABELS
            value: "app=helixterm-ssh-proxy"
        resources:
          requests:
            cpu: "100m"
            memory: "128Mi"
          limits:
            cpu: "500m"
            memory: "512Mi"
    volumes:
      - name: bpf
        hostPath:
          path: /sys/fs/bpf
          type: DirectoryOrCreate
      - name: proc
        hostPath:
          path: /proc
```

14.3.3 Custom Kafka Client in Go

The Confluent/IBM Sarama Kafka clients have non-trivial per-message overhead from CGO (Confluent) and reflection-heavy encoding (Sarama). A purpose-built Go Kafka client for the terminal stream topic (highest throughput path) can reduce per-message overhead from $\sim 4\mu\text{s}$ to $\sim 0.5\mu\text{s}$:

Target architecture for terminal stream Kafka path:

Current: SSH proxy $\rightarrow$ Sarama client (reflect encoding) $\rightarrow$ Kafka
 $\rightarrow$ Sarama consumer $\rightarrow$ WS
Per-message overhead: $\sim 4.2\mu\text{s}$, ~ 2 allocs

Target: SSH proxy $\rightarrow$ custom client (zero-copy, binary protocol) $\rightarrow$ Kafka $\rightarrow$ custom consumer $\rightarrow$ WS
Per-message overhead: $\sim 0.4\mu\text{s}$, 0 allocs (buffer pool)

Key optimisations in custom client:

1. Pre-encode Kafka protocol headers into reusable byte buffers (no reflection).
2. Use `io.WriterTo` interface for zero-copy sends from terminal buffer.
3. Maintain a per-partition sorted ring buffer for out-of-order message assembly.
4. Implement produce pipelining (send next batch before ACK for previous).
5. Use `io_uring` on Linux 5.1+ for async I/O with zero system call overhead.

Expected throughput improvement: 5–10 $\times$ for small messages (< 256 bytes) typical of terminal keystrokes. For larger messages (paste operations, `ls` output), improvement is smaller ($\sim 2\times$) as serialization is a smaller fraction of total cost.

Engineering estimate: 6–8 weeks for a production-ready custom client covering the terminal stream use case. Recommend starting with the consumer side (easier) before the producer.

Appendix A: Service SLO Reference Table

Service	Availability SLO	p50 Latency	p95 Latency	p99 Latency	Error Rate SLO
api-gateway	99.99%	< 10ms	< 100ms	< 500ms	< 0.01%
auth-service	99.99%	< 20ms	< 150ms	< 500ms	< 0.01%
ssh-proxy	99.95%	< 5ms	< 50ms	< 200ms	< 0.05%
vault-service	99.99%	< 15ms	< 100ms	< 500ms	< 0.01%
sync-service	99.9%	< 200ms	< 1s	< 2s	< 0.1%
audit-service	99.9%	< 5ms (async)	< 50ms	< 200ms	< 0.1%
notification-service	99.5%	< 500ms	< 2s	< 5s	< 0.5%
billing-service	99.9%	< 100ms	< 500ms	< 2s	< 0.1%
session-service	99.95%	< 10ms	< 50ms	< 200ms	< 0.05%
host-service	99.99%	< 15ms	< 100ms	< 300ms	< 0.01%
user-service	99.99%	< 15ms	< 100ms	< 300ms	< 0.01%
config-service	99.9%	< 20ms	< 100ms	< 500ms	< 0.1%
tag-service	99.9%	< 10ms	< 50ms	< 200ms	< 0.1%
search-service	99.5%	< 50ms	< 500ms	< 2s	< 0.5%
analytics-service	99.5%	< 100ms	< 1s	< 5s	< 0.5%
keygen-service	99.99%	< 50ms	< 200ms	< 1s	< 0.01%
cert-service	99.99%	< 100ms	< 500ms	< 2s	< 0.01%
	99.5%	< 200ms	< 2s	< 10s	< 0.5%

Service	Availability SLO	p50 Latency	p95 Latency	p99 Latency	Error Rate SLO
webhook-service					
export-service	99.0%	< 500ms	< 5s	< 30s	< 1.0%
import-service	99.0%	< 500ms	< 5s	< 30s	< 1.0%
team-service	99.99%	< 15ms	< 100ms	< 300ms	< 0.01%
org-service	99.99%	< 15ms	< 100ms	< 300ms	< 0.01%
rbac-service	99.99%	< 5ms	< 20ms	< 100ms	< 0.01%
policy-service	99.9%	< 50ms	< 200ms	< 1s	< 0.1%
health-service	99.99%	< 2ms	< 10ms	< 50ms	< 0.01%

Appendix B: Production Kubernetes Resource Profiles

```
# Reference resource requests/limits per service (baseline,
# tune with HPA)
# Values derived from load testing at 10k req/s sustained
# load.

services:
  api-gateway:
    replicas: { min: 10, max: 100 }
    resources:
      requests: { cpu: "500m", memory: "512Mi" }
      limits: { cpu: "2", memory: "2Gi" }

  ssh-proxy:
    replicas: { min: 20, max: 200 }
    resources:
      requests: { cpu: "1", memory: "1Gi" }
      limits: { cpu: "4", memory: "8Gi" }
    # SSH proxy is memory-bound: each session uses ~80KB
    # goroutine stack + buffers
```

```
vault-service:
  replicas: { min: 5, max: 50 }
  resources:
    requests: { cpu: "500m", memory: "512Mi" }
    limits: { cpu: "2", memory: "2Gi" }

auth-service:
  replicas: { min: 5, max: 50 }
  resources:
    requests: { cpu: "1", memory: "512Mi" }
    limits: { cpu: "4", memory: "2Gi" }
  # CPU request higher due to Argon2id hash pool

audit-service:
  replicas: { min: 3, max: 30 }
  resources:
    requests: { cpu: "200m", memory: "256Mi" }
    limits: { cpu: "1", memory: "1Gi" }
  # I/O bound – minimal CPU needed; generous memory for
    Kafka batch buffers

sync-service:
  replicas: { min: 5, max: 30 }
  resources:
    requests: { cpu: "500m", memory: "1Gi" }
    limits: { cpu: "2", memory: "4Gi" }
  # Memory higher due to CRDT operation buffers
```

Document version: 1.0.0 | Last updated: 2026-06-28 | Owner: Platform Performance Engineering
Review cycle: Quarterly or after any incident with performance root cause
Source of truth: this document + [helixterm.io/services//internal/metrics/packages](https://helixterm.io/services/*/internal/metrics/packages)*